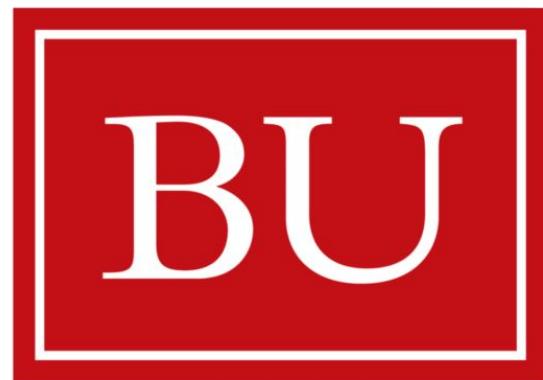


# ORQ: Complex Analytics on Private Data with Strong Security Guarantees

Eli Baum



**joint work with** Sam Buxbaum, Nitin Mathai, Muhammad Faisal, Vasia Kalavri, Mayank Varia, and John Liagouris  
SOSP 2025 – 16 October 2025

# Motivation: Collaborative Analytics on Sensitive Data

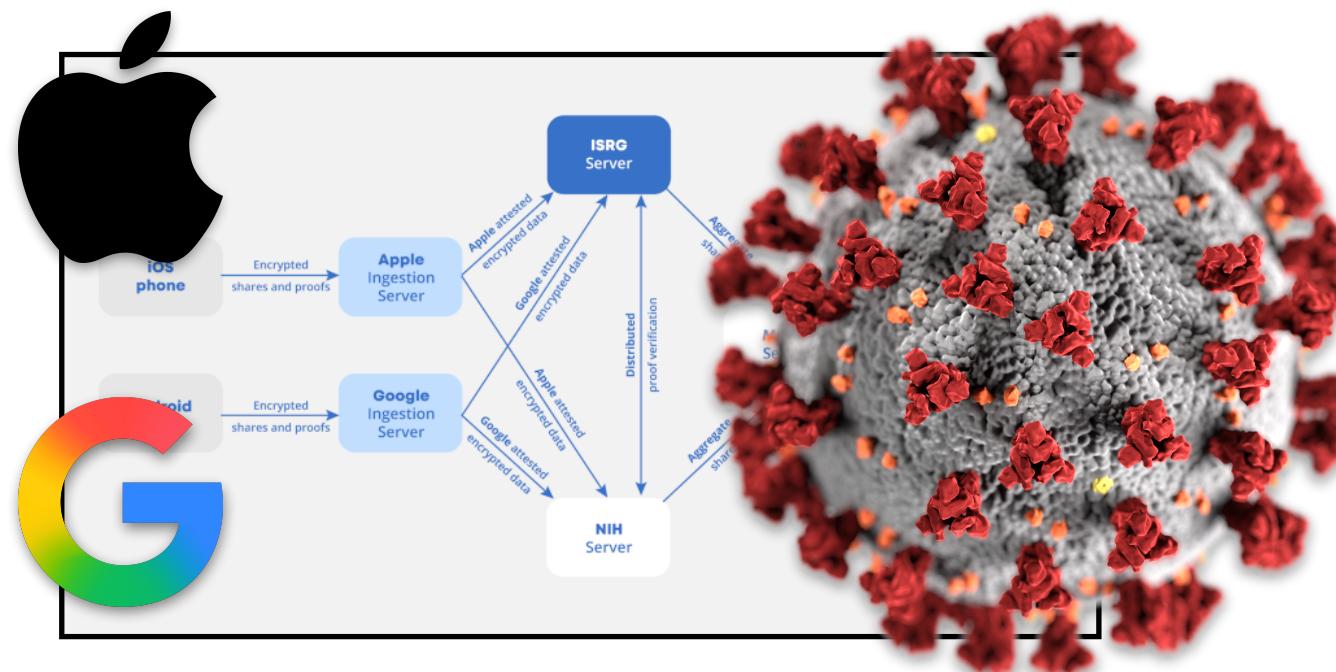
Common interest

Private data

---

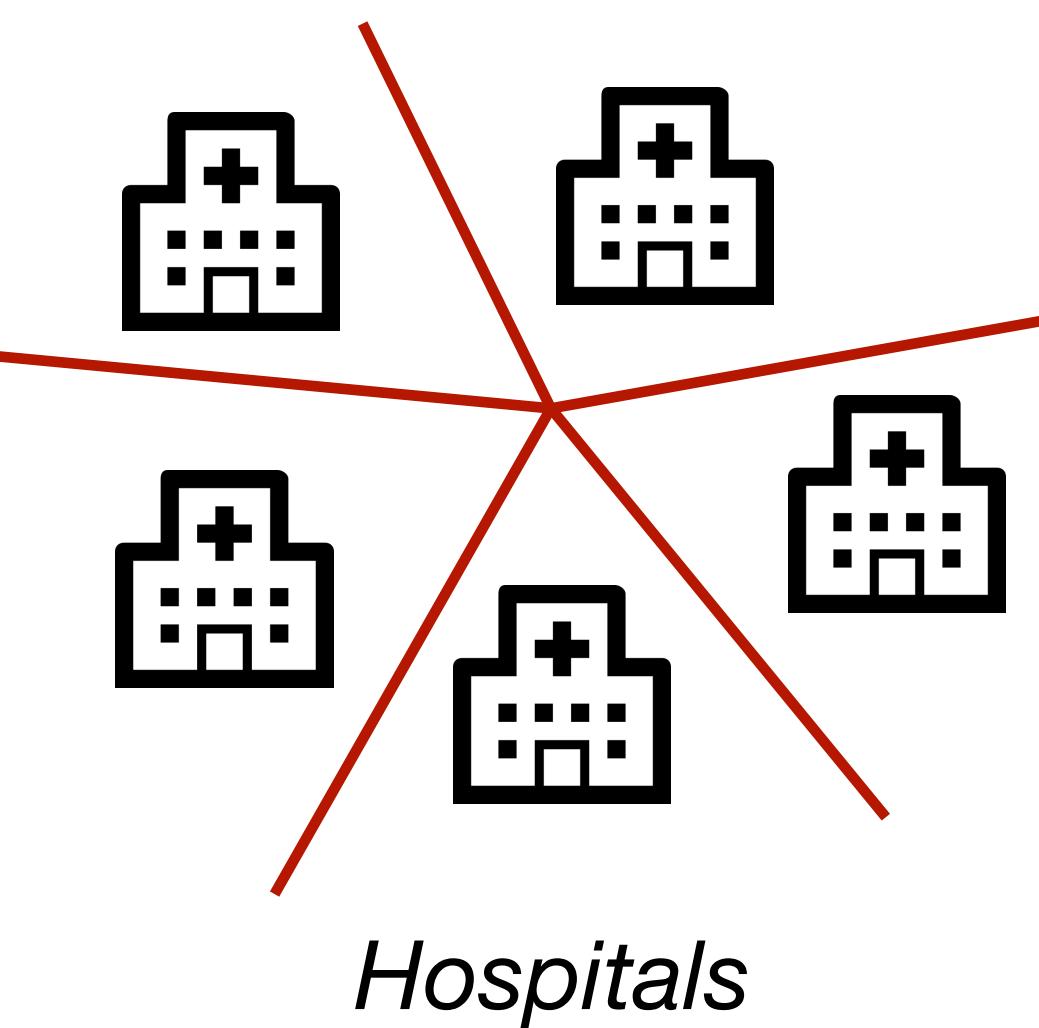
# Motivation: Collaborative Analytics on Sensitive Data

Common interest



Disease Surveillance

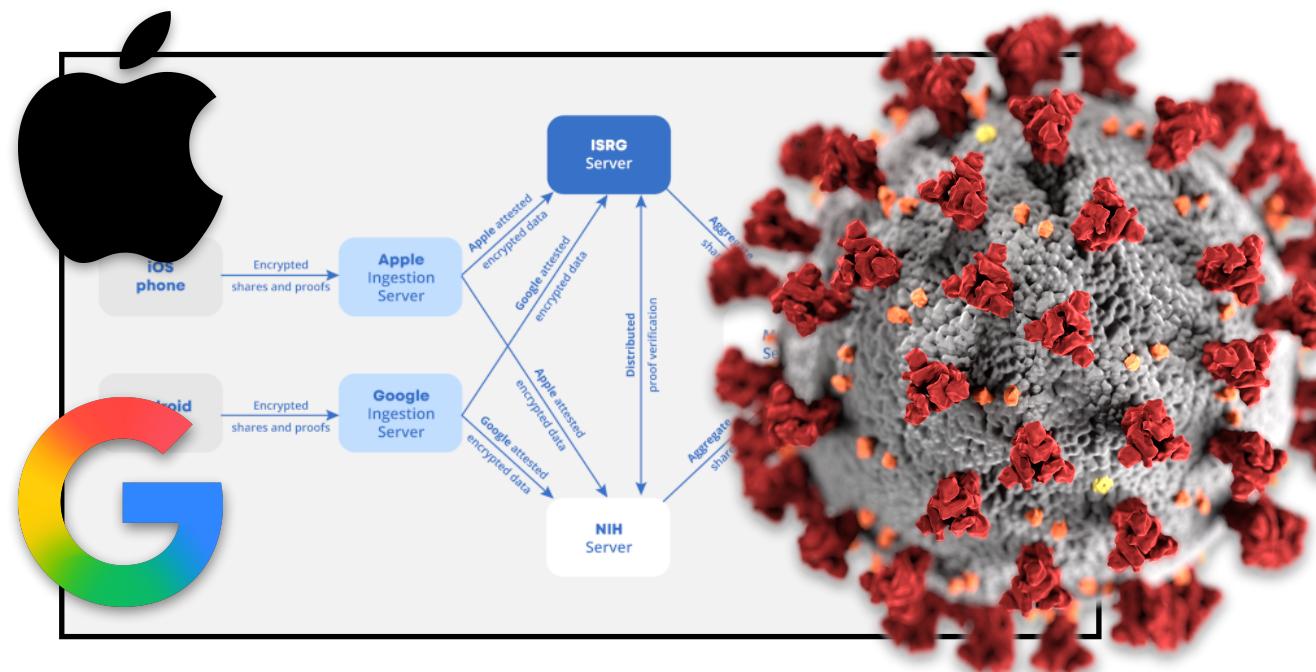
Private data



Hospitals

# Motivation: Collaborative Analytics on Sensitive Data

Common interest

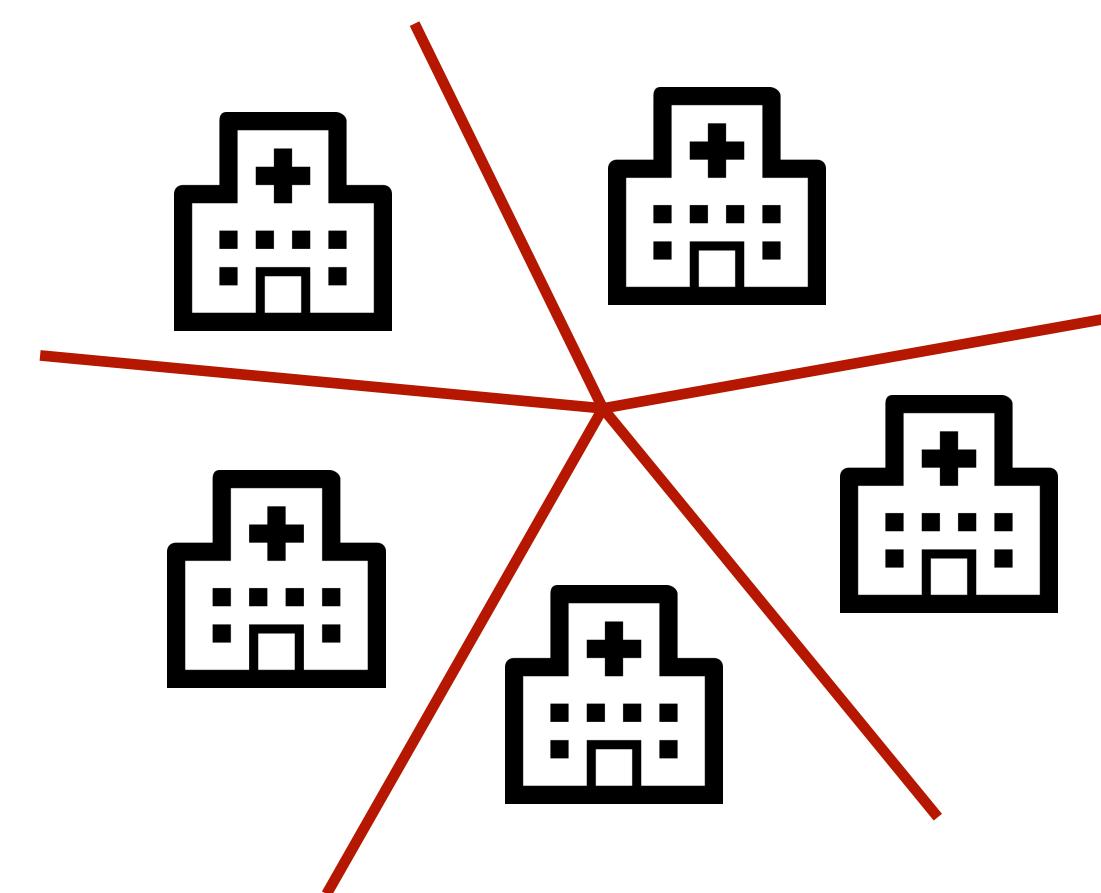


Disease Surveillance

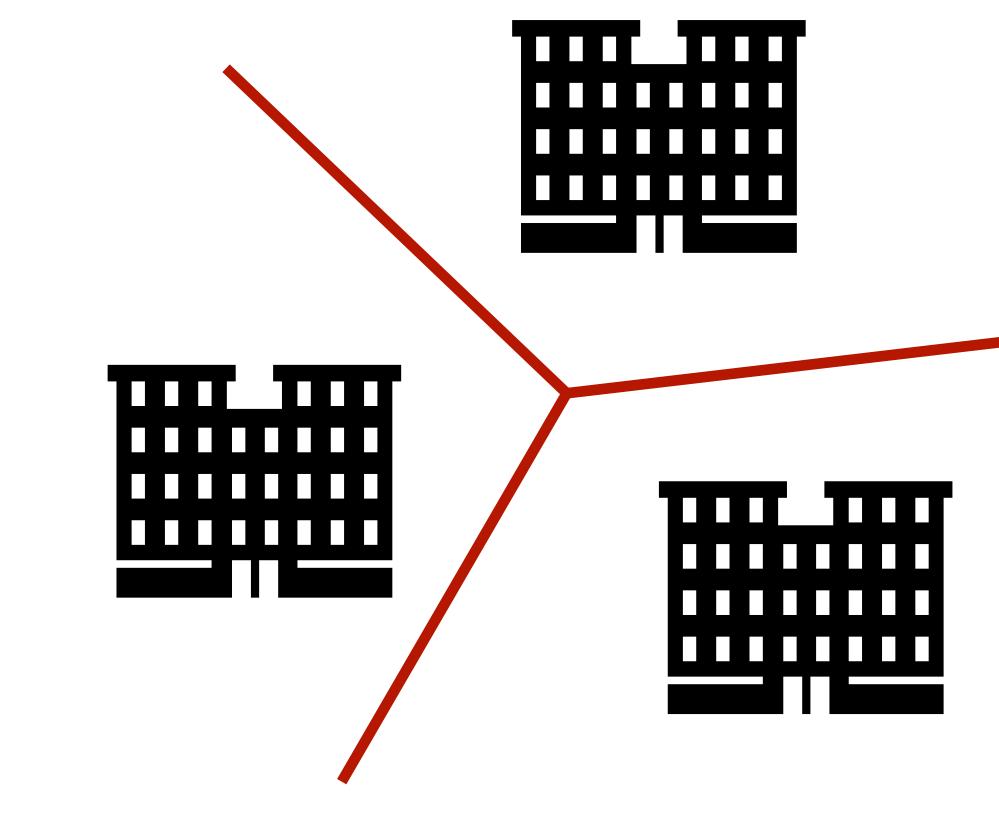


Gender Wage Gap

Private data



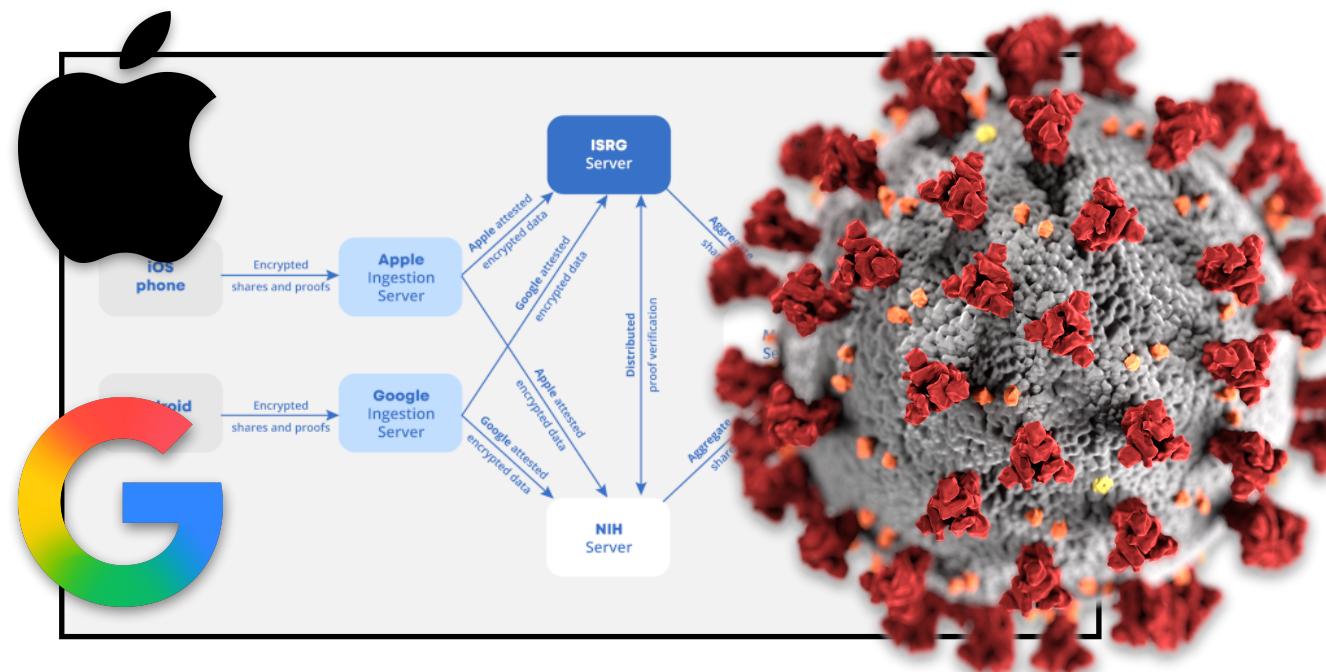
Hospitals



Employers

# Motivation: Collaborative Analytics on Sensitive Data

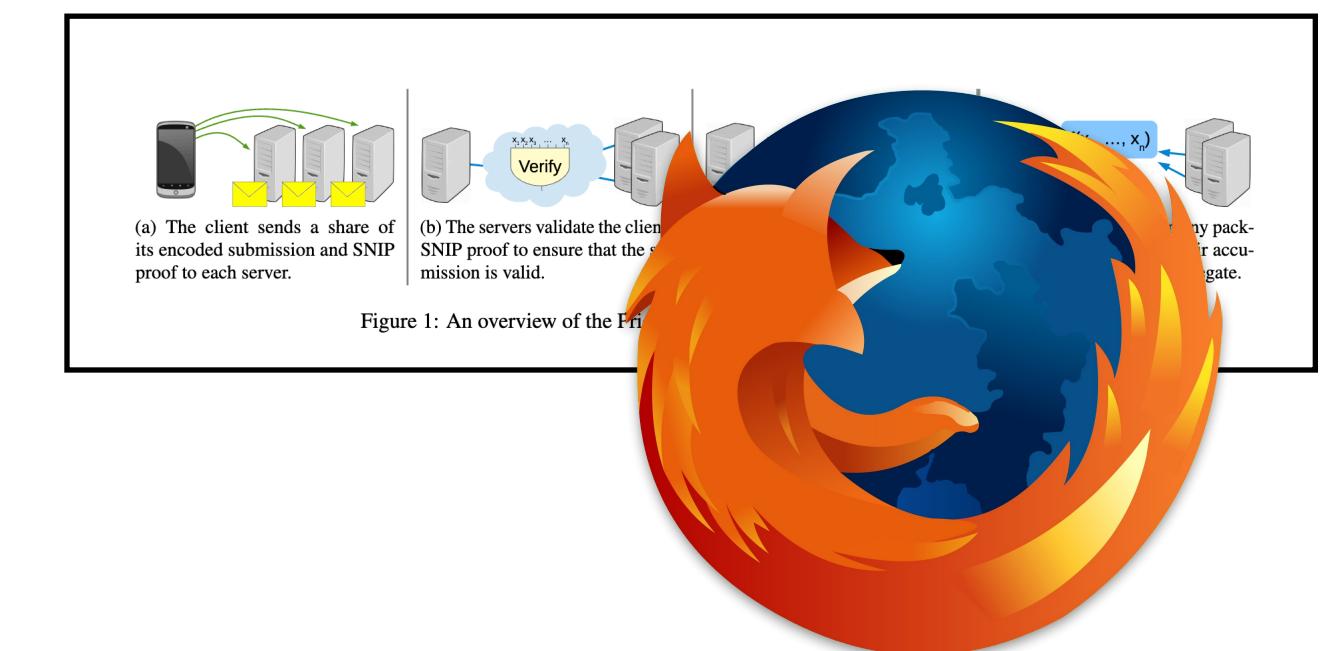
Common interest



Disease Surveillance

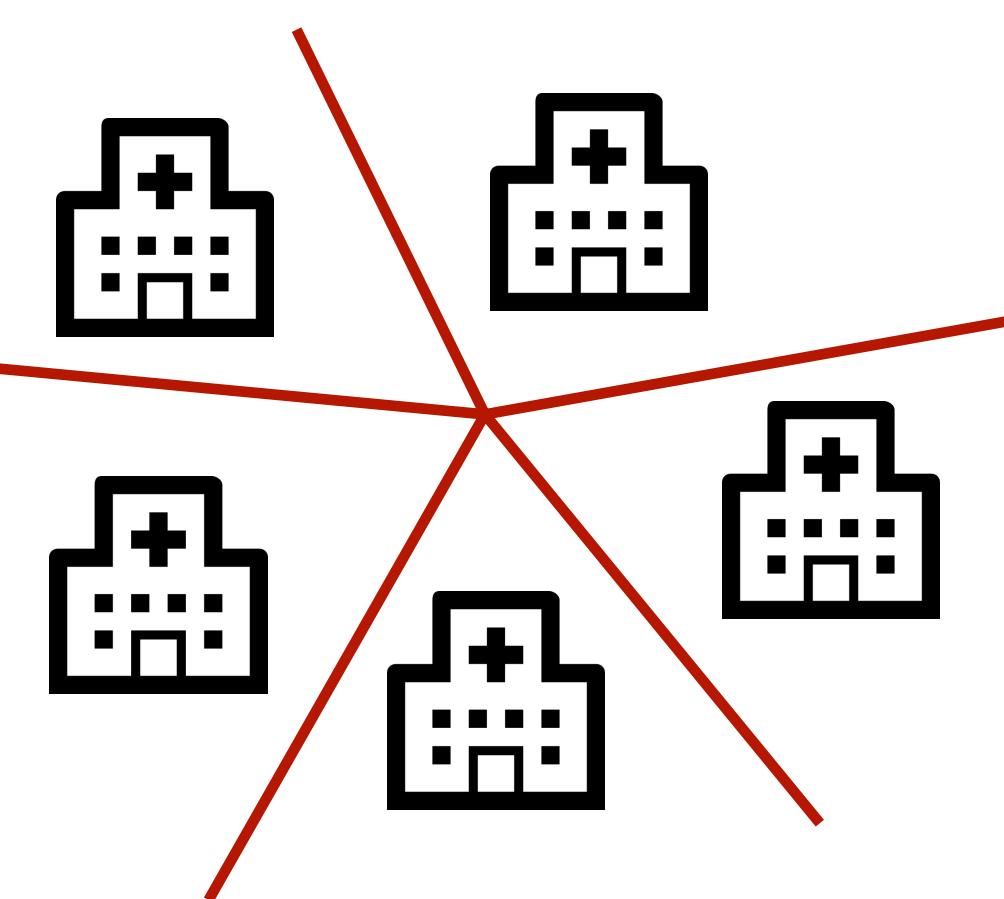


Gender Wage Gap

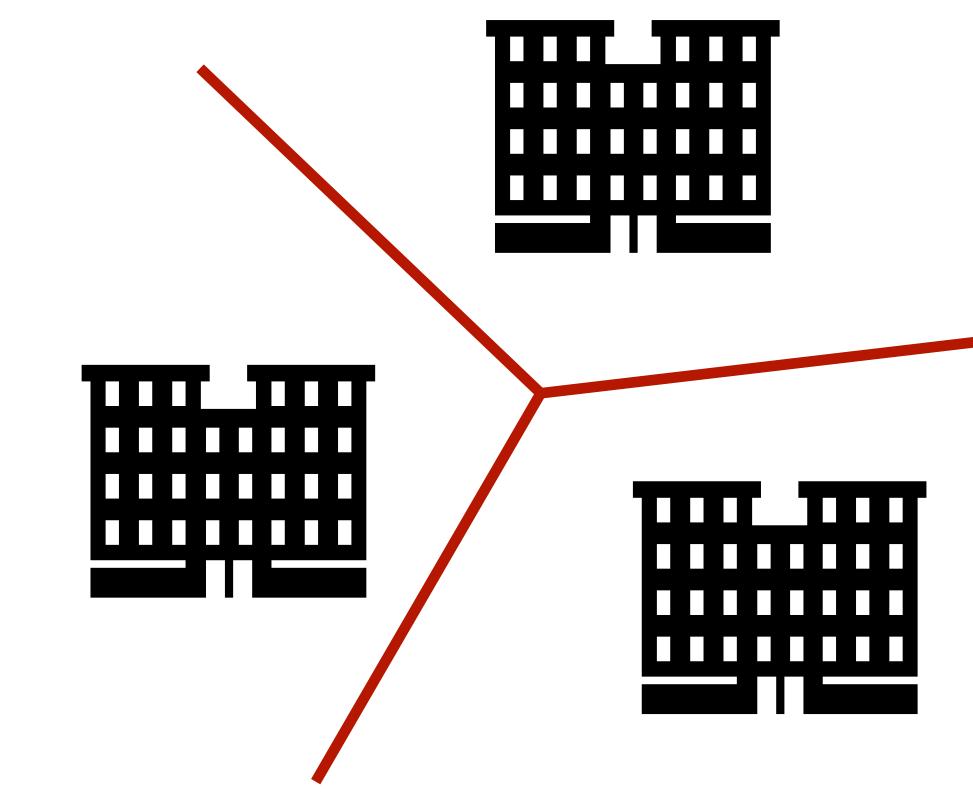


Browser Telemetry

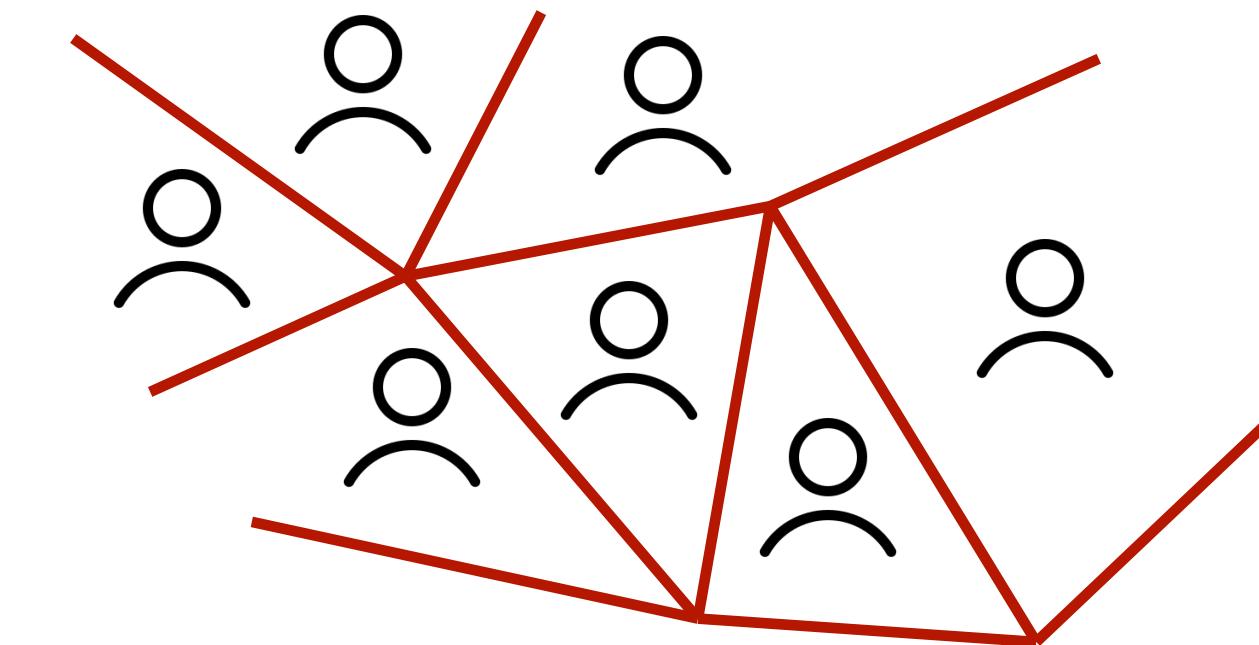
Private data



Hospitals



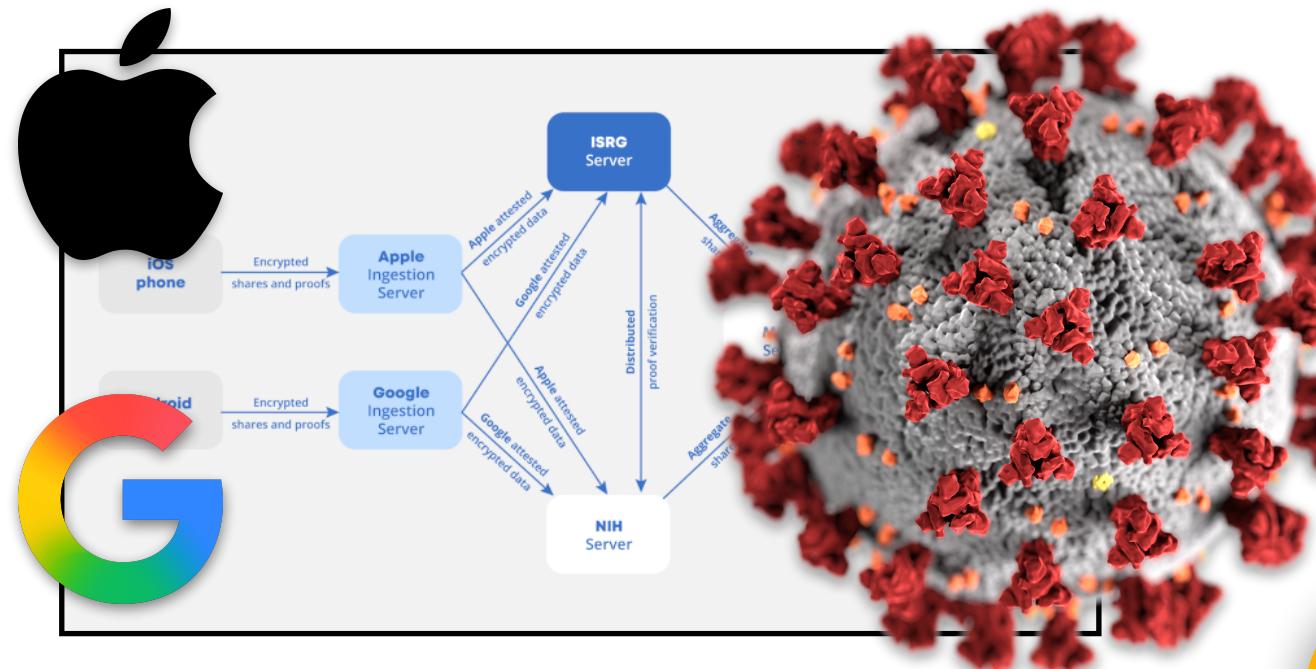
Employers



End Users

# Motivation: Collaborative Analytics on Sensitive Data

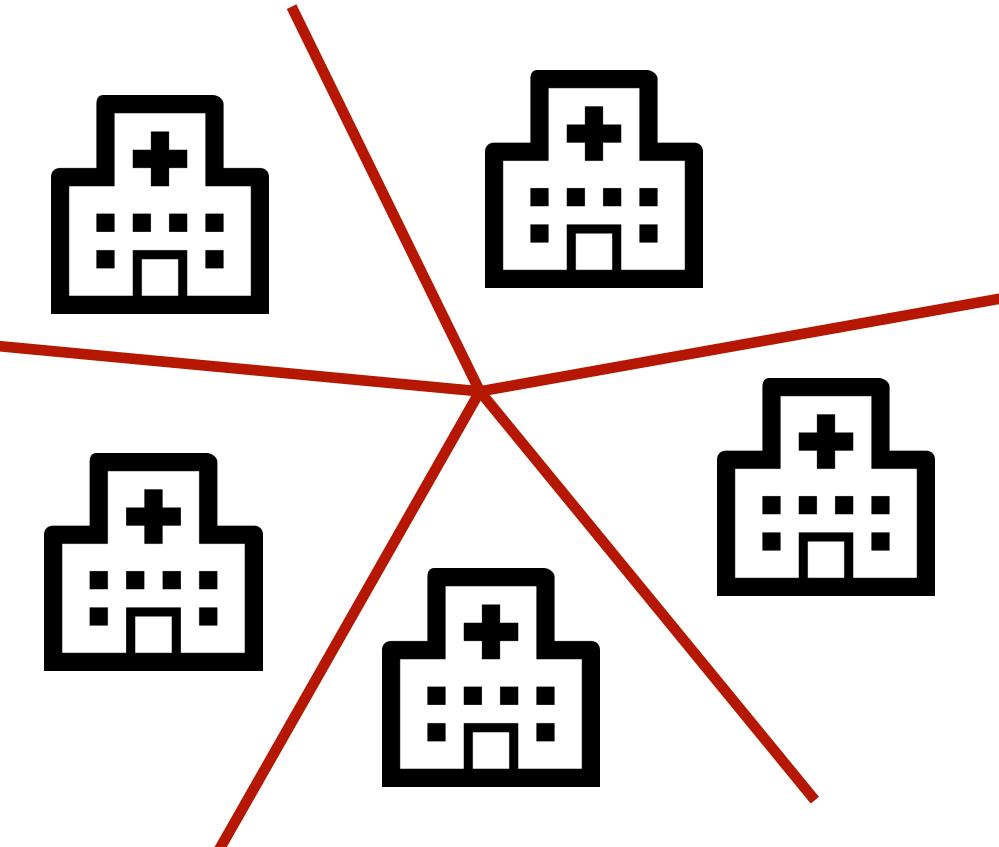
Common interest



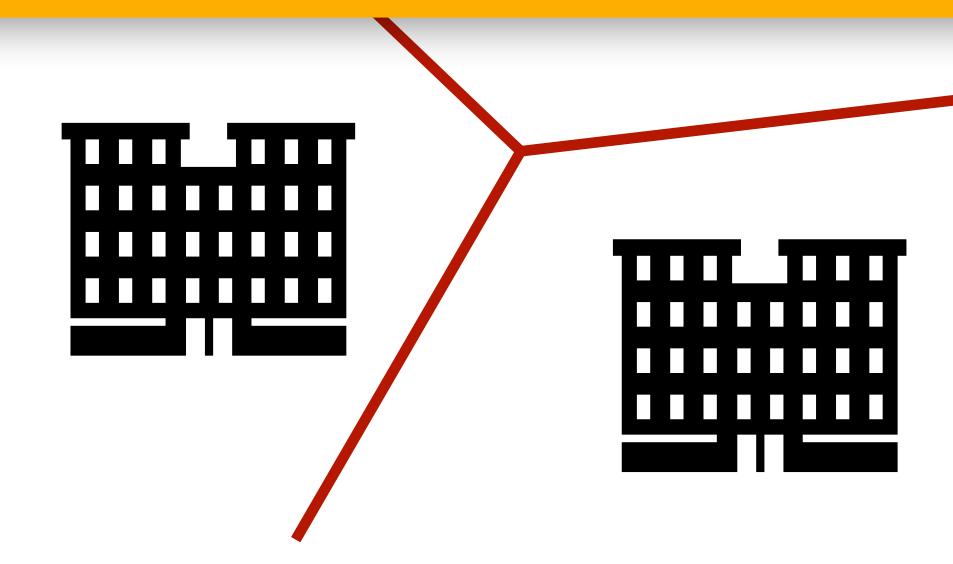
Disease Surveillance



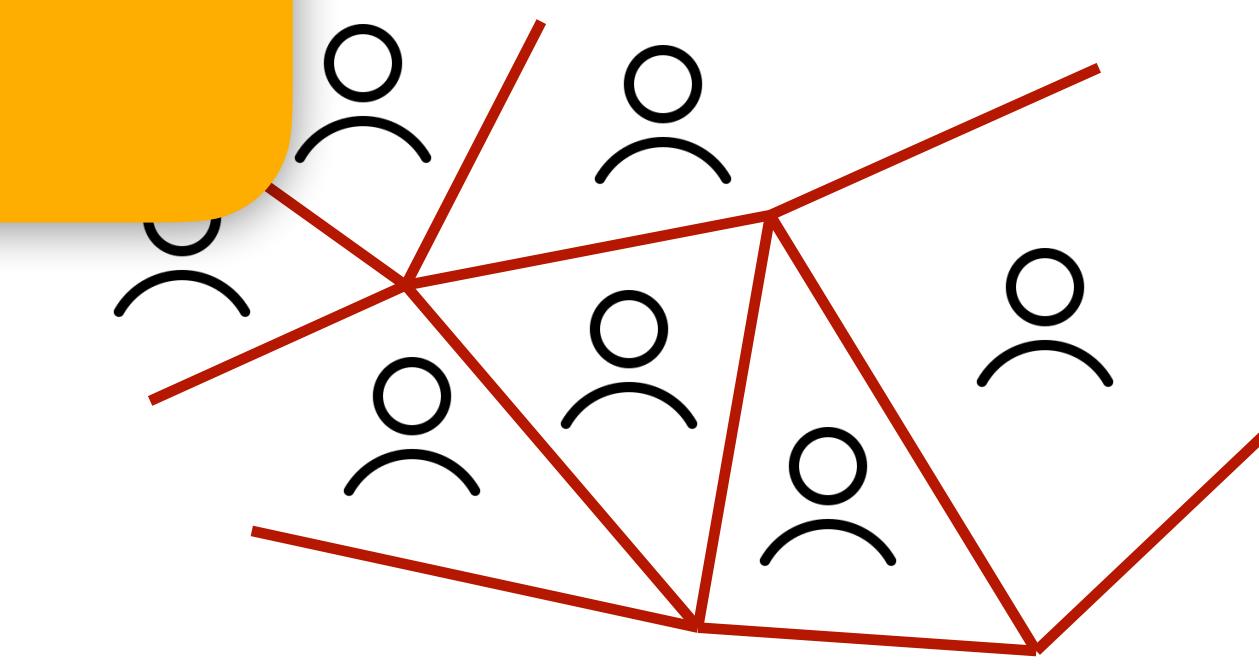
Private data



Hospitals

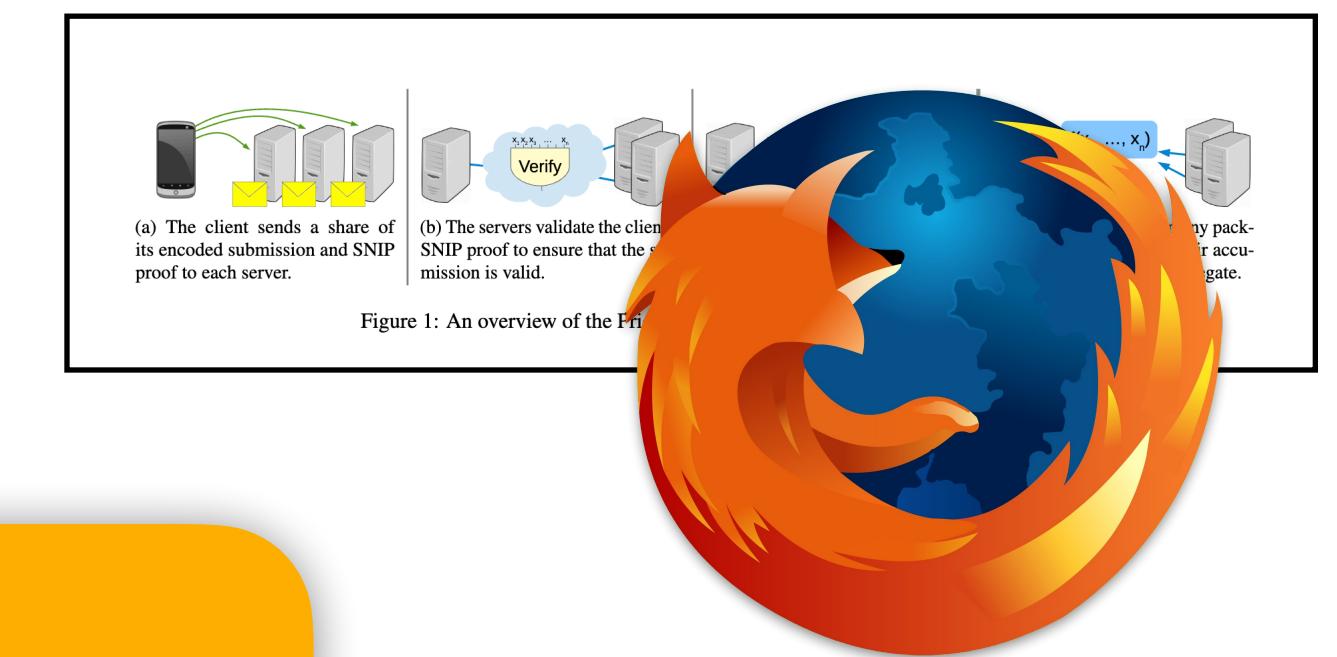


Employers



End Users

**Collaborative analytics requires  
securely joining data from  
multiple sources!**



Browser Telemetry

# Outline

- Background on Oblivious Computation
- Our Approach
- Making the system work
- Results

# Plaintext Analytics

Customer	AcctBalance
🤡	\$1000
😈	\$2000
🤡	\$500
🐸	\$60
🥶	\$8000
🥶	\$700
🐸	-\$200

filter(Customer ≠ 🐸)



Customer	AcctBalance
🤡	\$1000
😈	\$2000
🤡	\$500
🥶	\$8000
🥶	\$700

# Plaintext Analytics

# Adversary View

The figure displays a 7x2 grid of 14 grayscale images. The left column is labeled "Customer" and the right column is labeled "AcctBalance". Each image shows a highly pixelated and noisy version of a handwritten digit, likely a 0 or 1, with significant visual artifacts.

```
filter(Customer != null)
```



The figure displays a 5x2 grid of images. The left column is labeled "Customer" and the right column is labeled "AcctBalance". Each row contains two images, one from each column. The images are highly pixelated and show various patterns of black, white, and gray squares.

# Table sizes leak information!

# Oblivious Analytics

# Complex oblivious operators must preserve **worst-case** sizes

The figure displays a 7x2 grid of 14 28x28 pixel grayscale images. The left column is labeled "Customer" and the right column is labeled "AcctBalance". Each image is heavily noisy, consisting of a mix of black, white, and various shades of gray pixels. There is no discernible pattern or text within the images.

```
filter(Customer != 🐸)
```



# Oblivious Joins with Cartesian Product

Customer	Order
Id	CustomerId
🤡	🐸
😈	🤡
🤡	🥶
🐸	🥶
🥶	🤡
🥶	🤡
🐸	😈

# Oblivious Joins with Cartesian Product



# Oblivious Joins with Cartesian Product

Cust \ Ord	🐸	🤡	🥶	🥶	🤡	🤡	😈

# Oblivious Joins with Cartesian Product

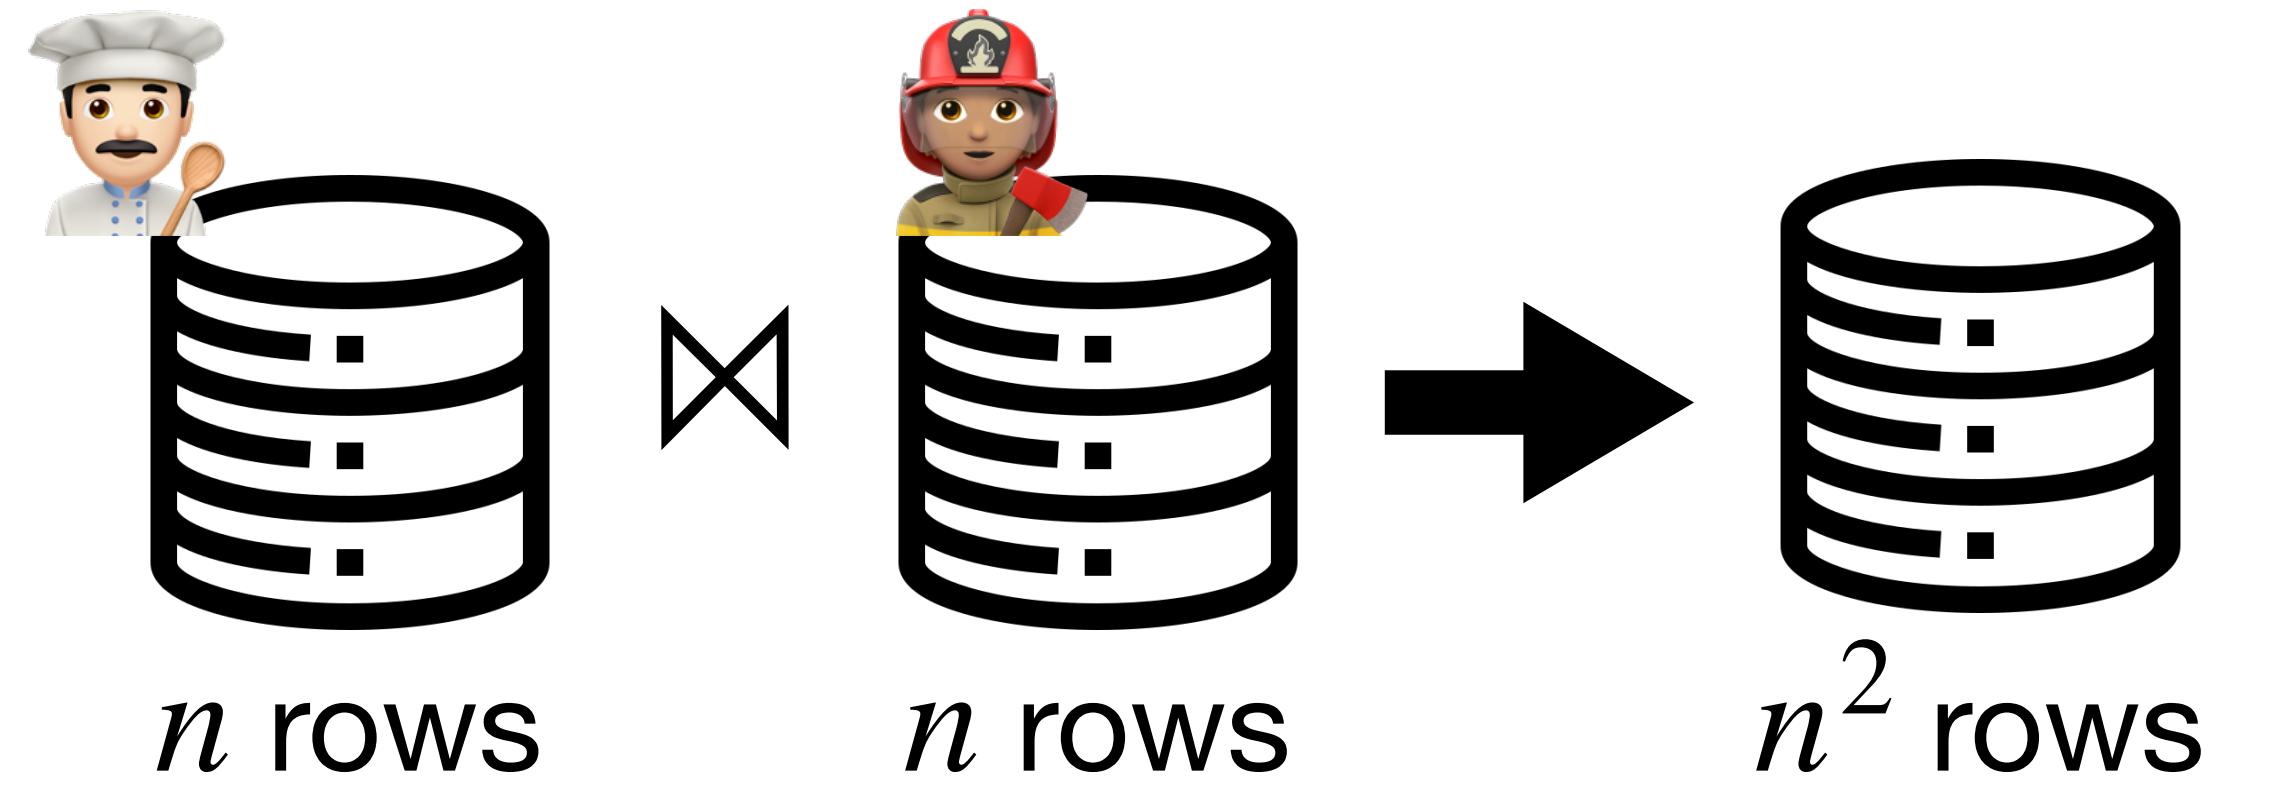
Cust \ Ord	🐸	🤡	🥶	🥶	🤡	🤡	😈
🤡	✗	✓	✗	✗	✓	✓	✗
😈	✗	✗	✗	✗	✗	✗	✓
🤡	✗	✓	✗	✗	✓	✓	✗
🐸	✓	✗	✗	✗	✗	✗	✗
🥶	✗	✗	✓	✓	✗	✗	✗
🥶	✗	✗	✓	✓	✗	✗	✗
🐸	✓	✗	✗	✗	✗	✗	✗

# Oblivious Joins with Cartesian Product

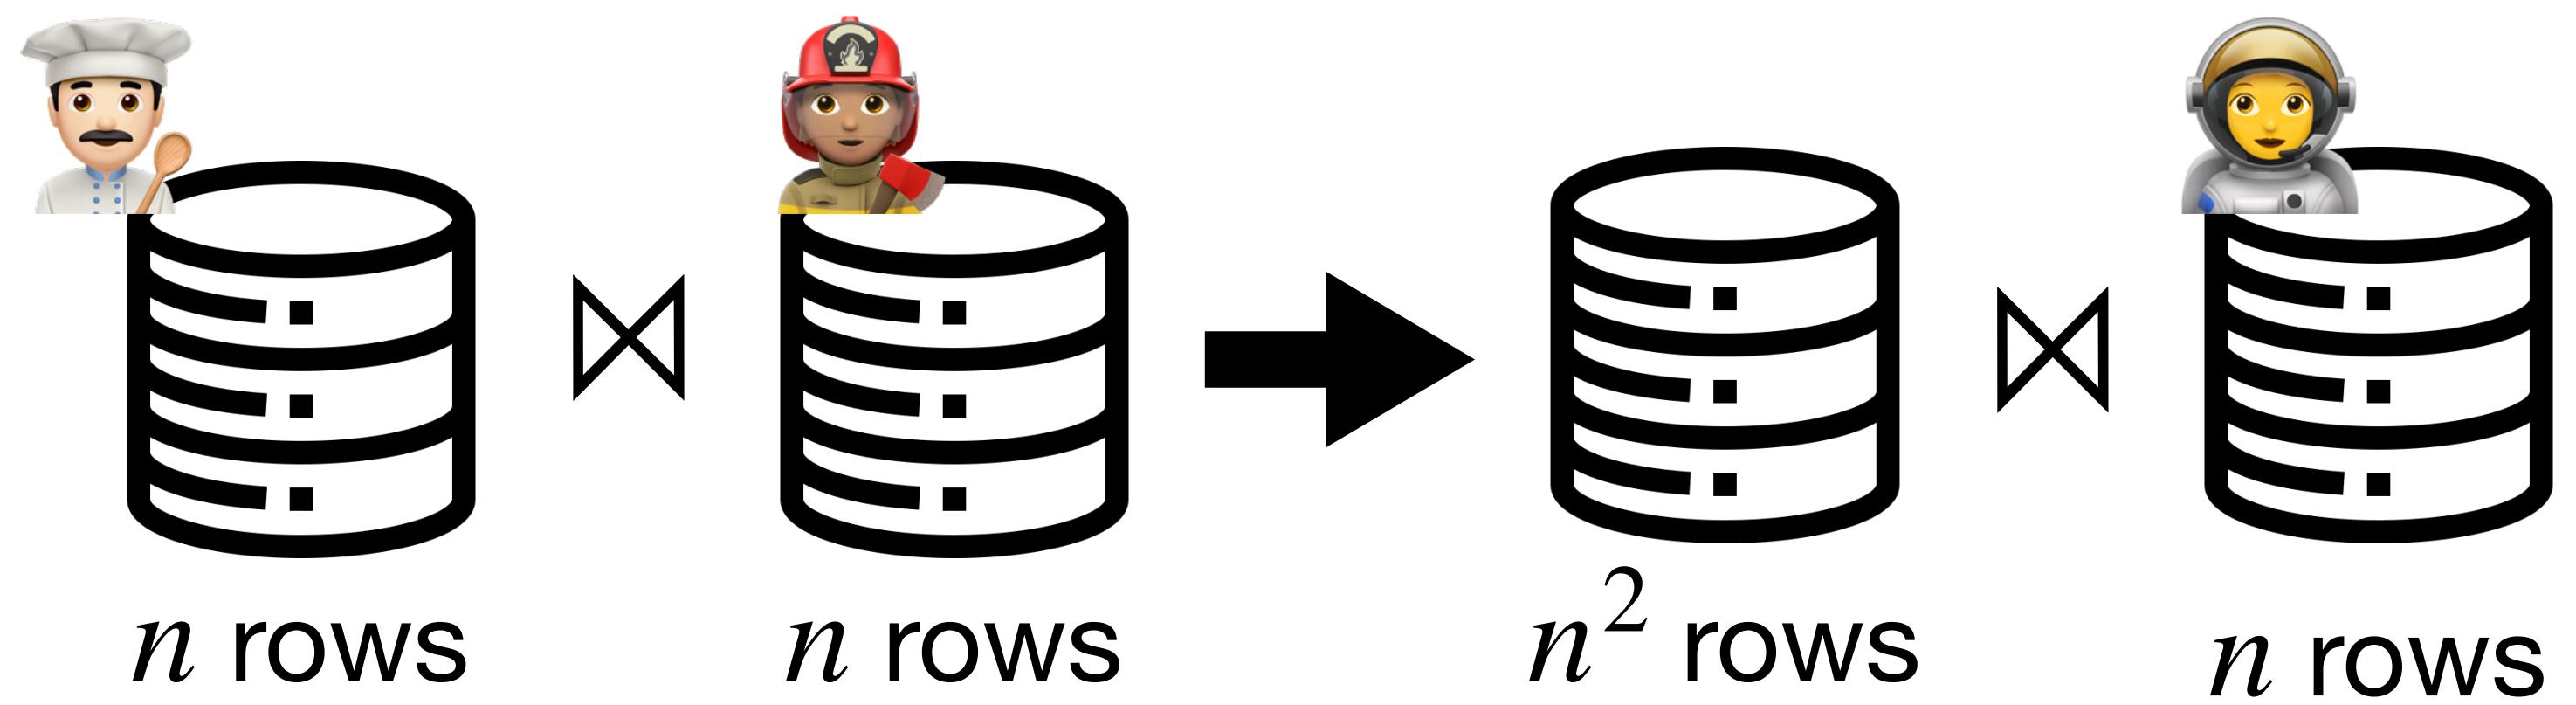
Cust \ Ord	🐸	🤡	🥶	💦	🤡	🤡	😈
🤡	✗	✓	✗	✗	✓	✓	✗
😈	✗	✗	✗	✗	✗	✗	✓
🤡	✗	✓	✗	✗	✓	✓	✗
🐸	✓	✗	✗	✗	✗	✗	✗
🥶	✗	✗	✓	✓	✗	✗	✗
💦	✗	✗	✓	✓	✗	✗	✗
🐸	✓	✗	✗	✗	✗	✗	✗

$O(n^2)$  work  
 $O(n^2)$  output size

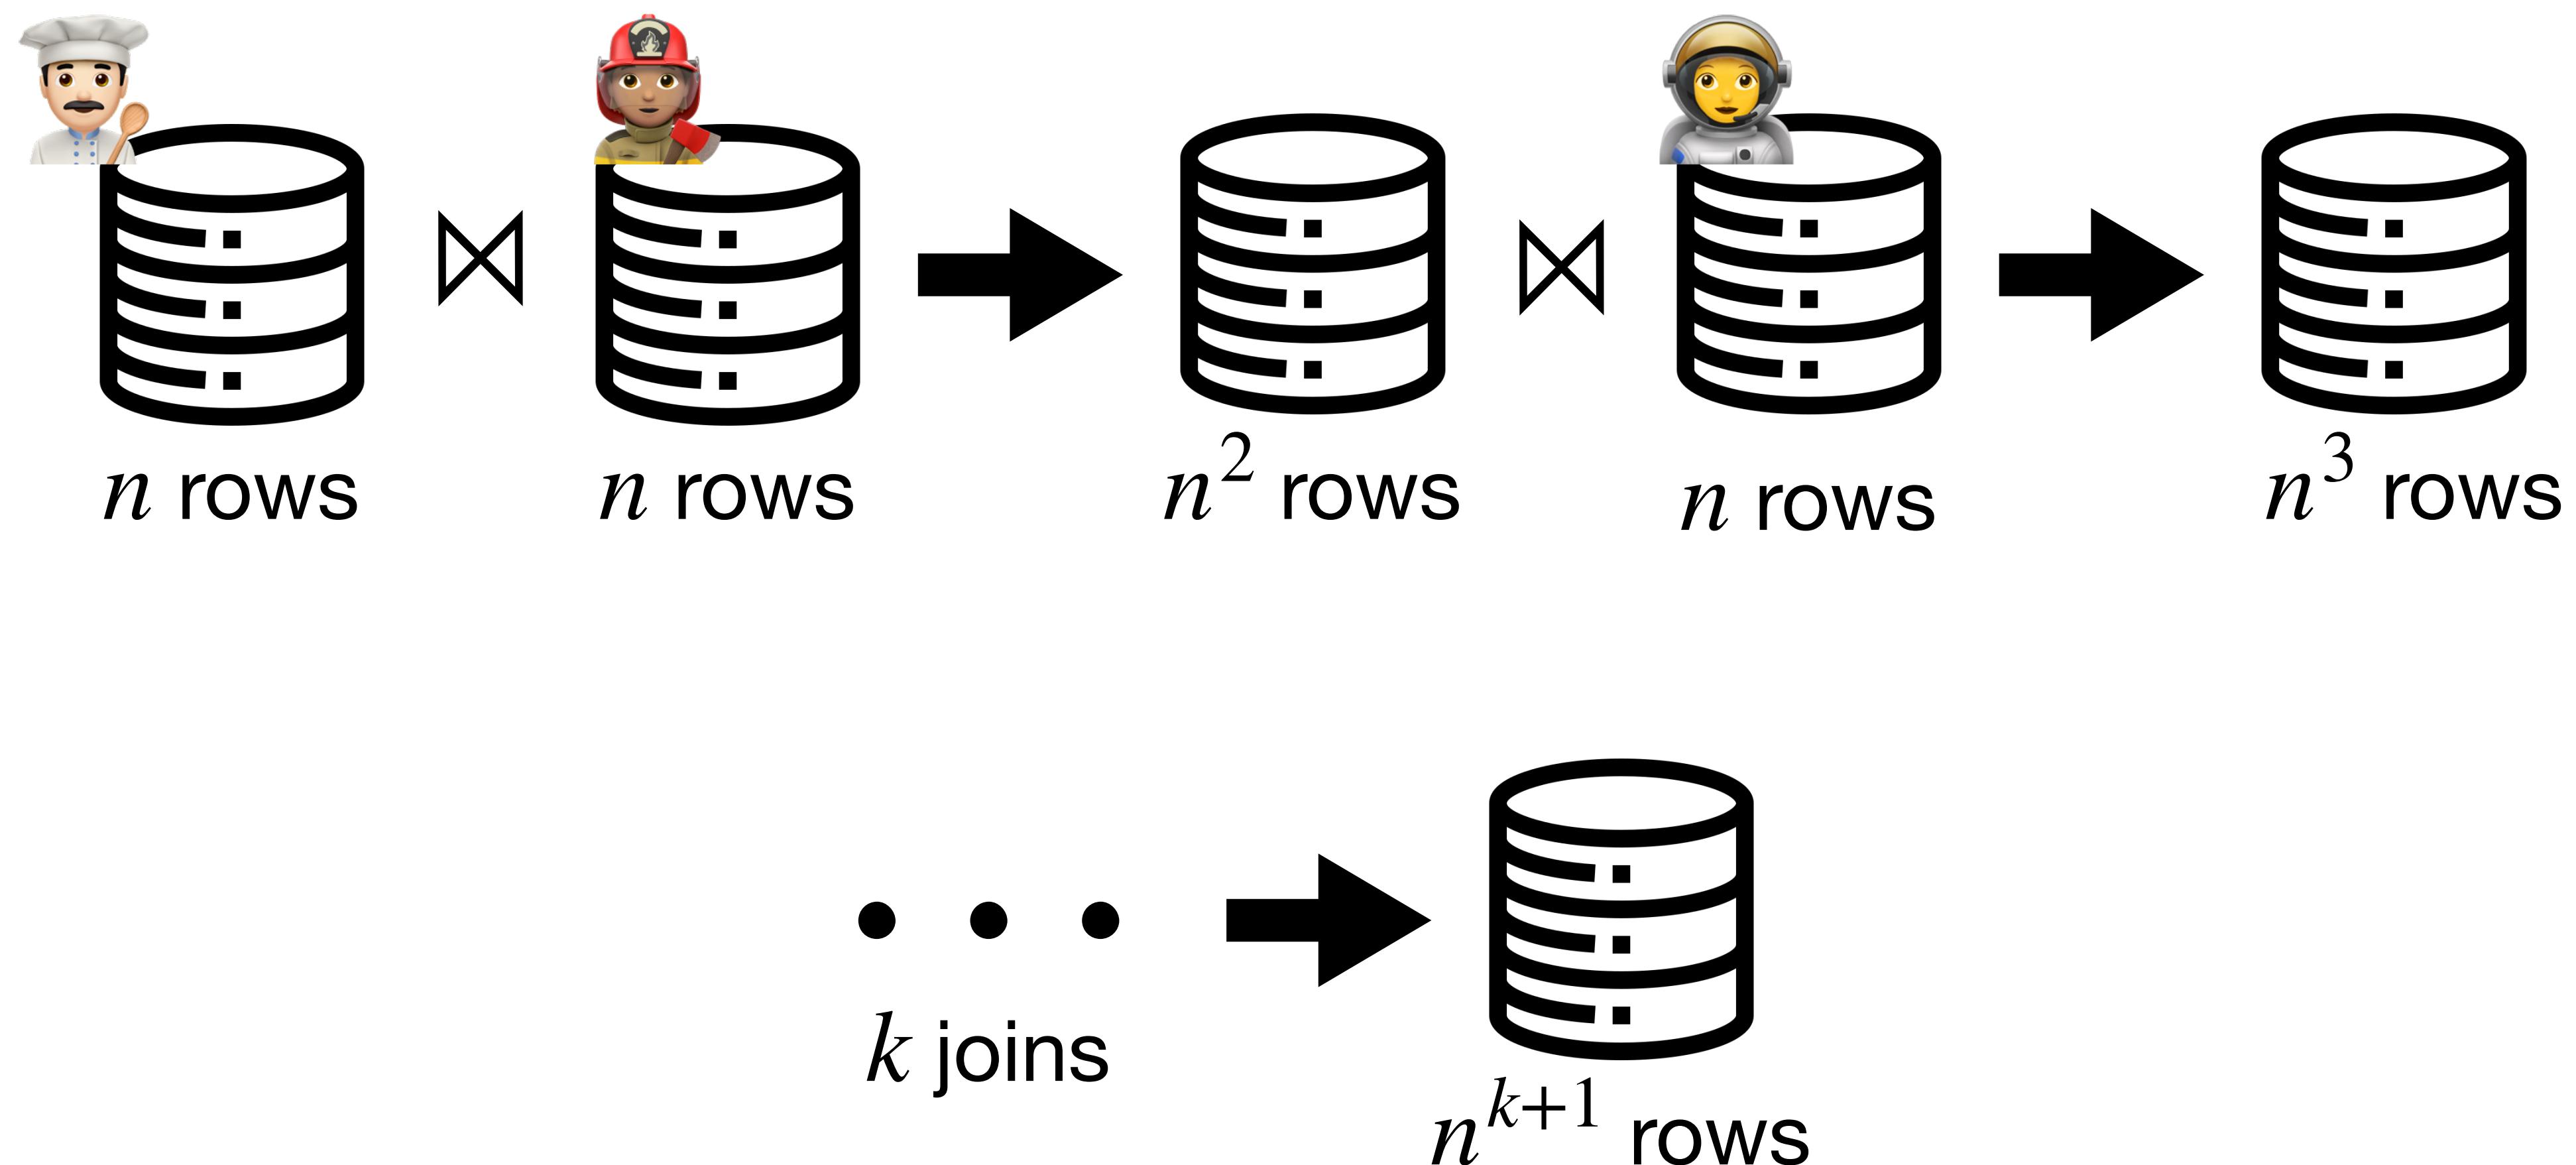
# The Problem



# The Problem

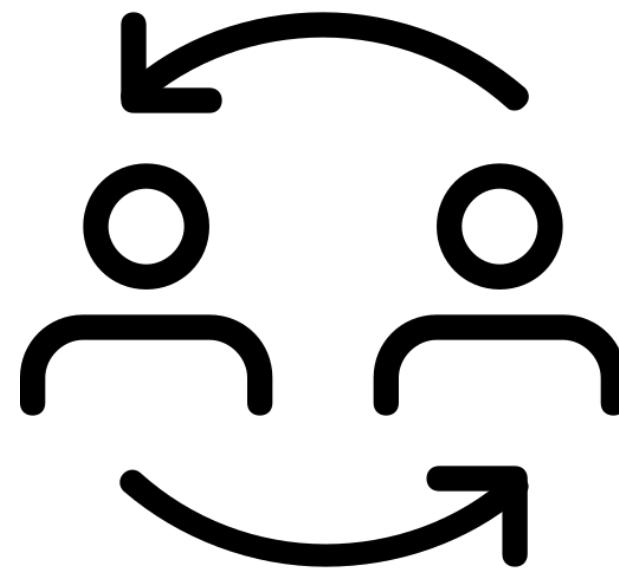


# The Problem



# Prior Approaches to Collaborative Analytics

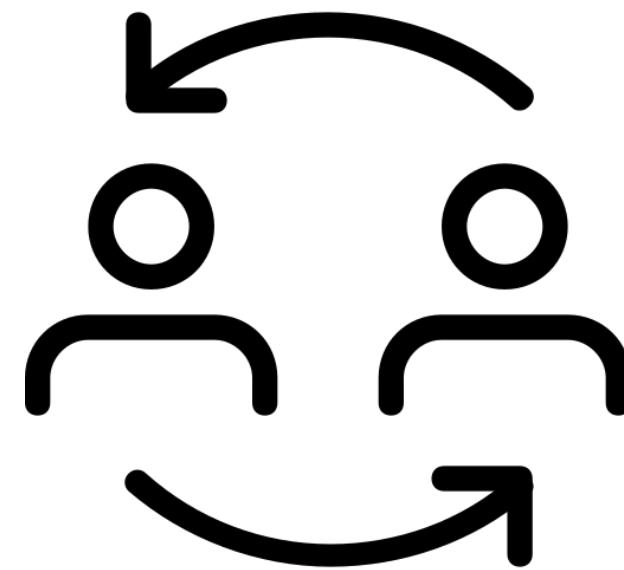
# Prior Approaches to Collaborative Analytics



## Peer-to-peer setting

e.g., Conclave (EuroSys'19),  
Senate (Security'21), SMCQL (VLDB'17),  
Shrinkwrap (VLDB'18)

# Prior Approaches to Collaborative Analytics

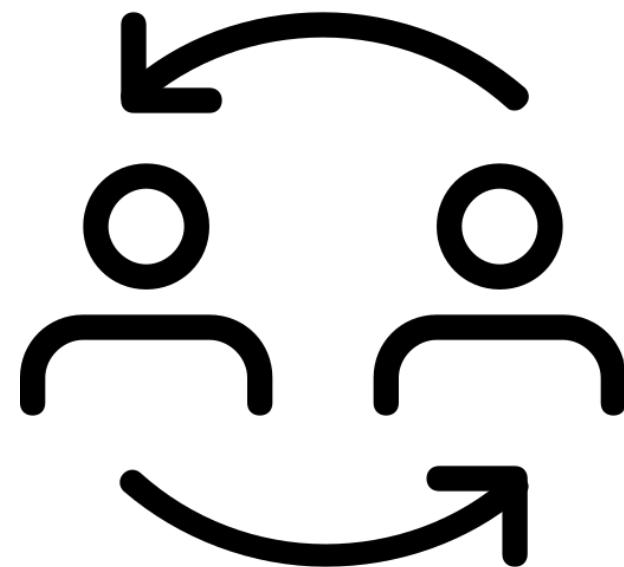


## Peer-to-peer setting

e.g., Conclave (EuroSys'19),  
Senate (Security'21), SMCQL (VLDB'17),  
Shrinkwrap (VLDB'18)

✗ Doesn't scale with more parties

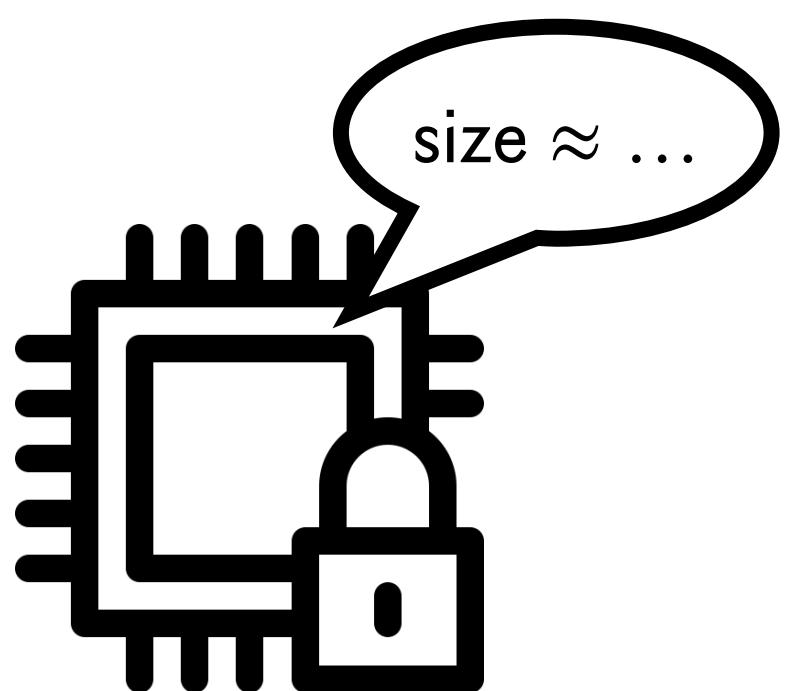
# Prior Approaches to Collaborative Analytics



## Peer-to-peer setting

e.g., Conclave (EuroSys'19),  
Senate (Security'21), SMCQL (VLDB'17),  
Shrinkwrap (VLDB'18)

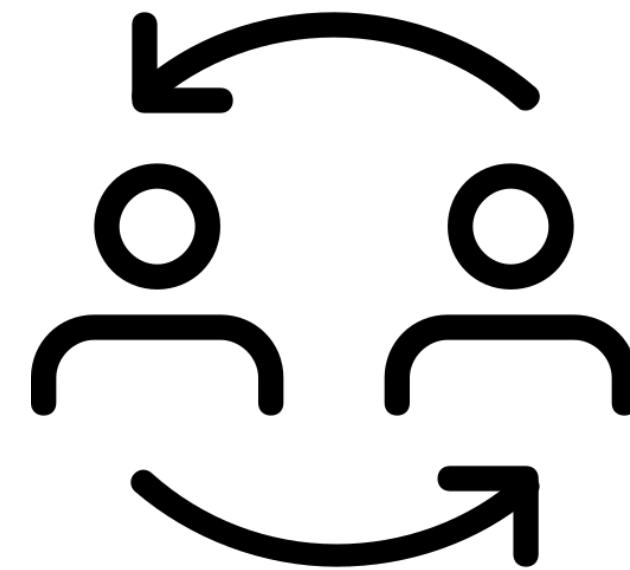
✗ Doesn't scale with more parties



## Leakage or upper bounds

e.g., Opaque (NSDI'17), Obliviator (Security'25),  
SAQE (VLDB'20)

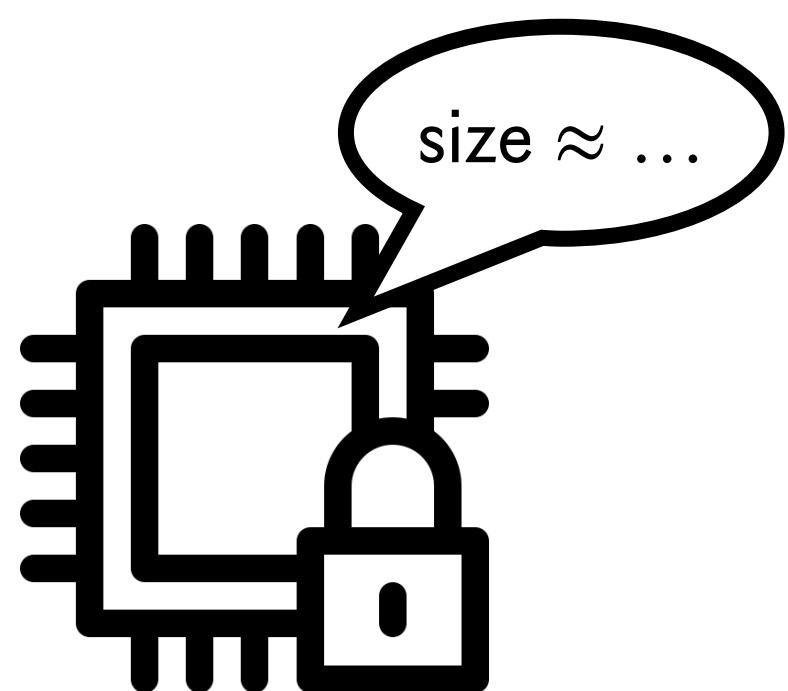
# Prior Approaches to Collaborative Analytics



## Peer-to-peer setting

e.g., Conclave (EuroSys'19),  
Senate (Security'21), SMCQL (VLDB'17),  
Shrinkwrap (VLDB'18)

✗ Doesn't scale with more parties

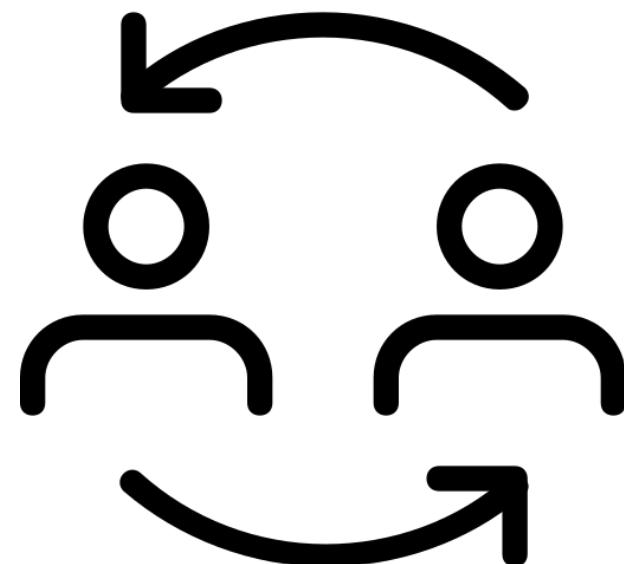


## Leakage or upper bounds

e.g., Opaque (NSDI'17), Obliviator (Security'25),  
SAQE (VLDB'20)

✗ Only provides partial security

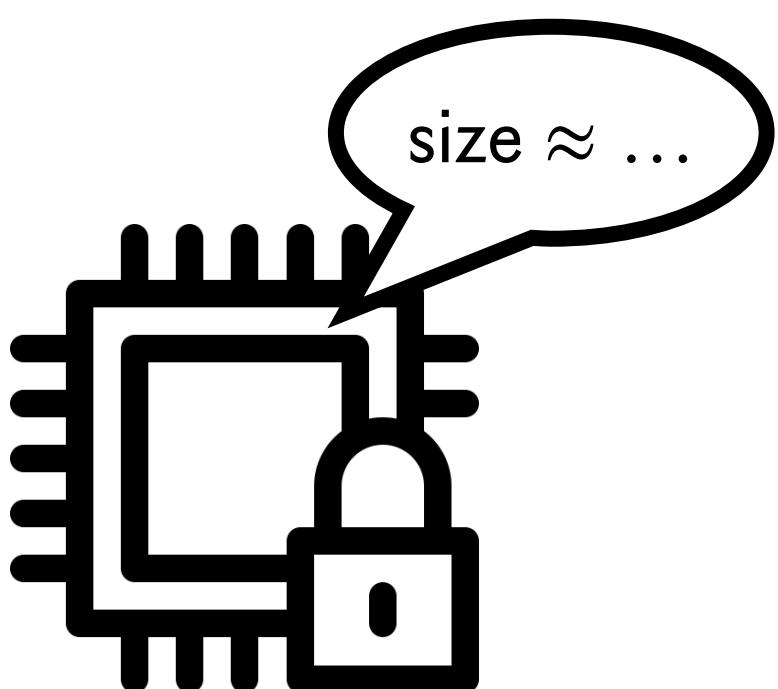
# Prior Approaches to Collaborative Analytics



## Peer-to-peer setting

e.g., Conclave (EuroSys'19),  
Senate (Security'21), SMCQL (VLDB'17),  
Shrinkwrap (VLDB'18)

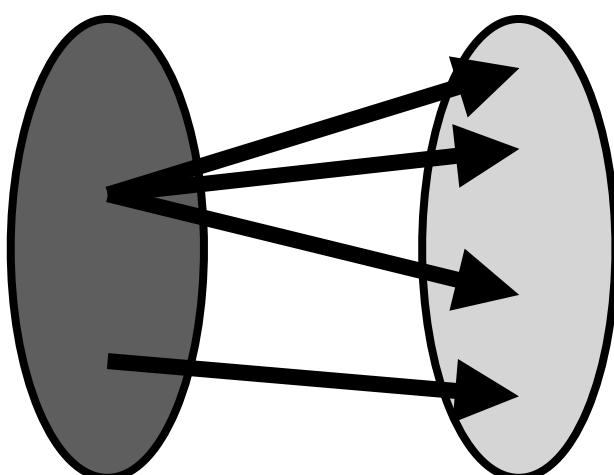
✗ Doesn't scale with more parties



## Leakage or upper bounds

e.g., Opaque (NSDI'17), Obliviator (Security'25),  
SAQE (VLDB'20)

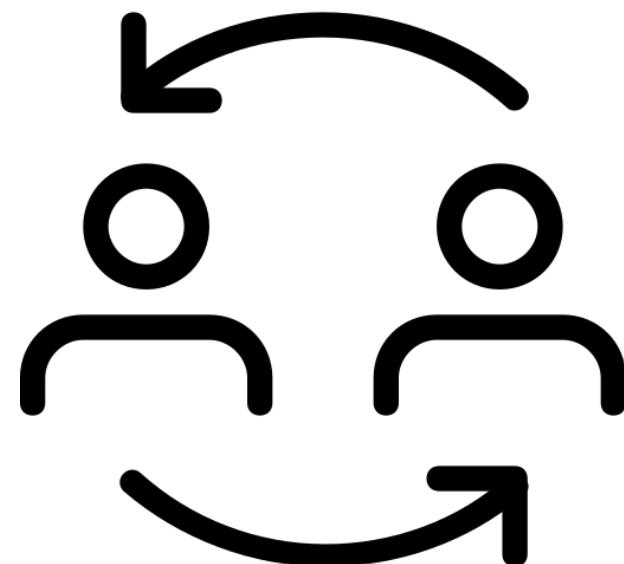
✗ Only provides partial security



## Individual join operators

e.g., [MRR] (CCS'20), [BDG+] (CCS'22),  
[AHK+] (CCS'23)

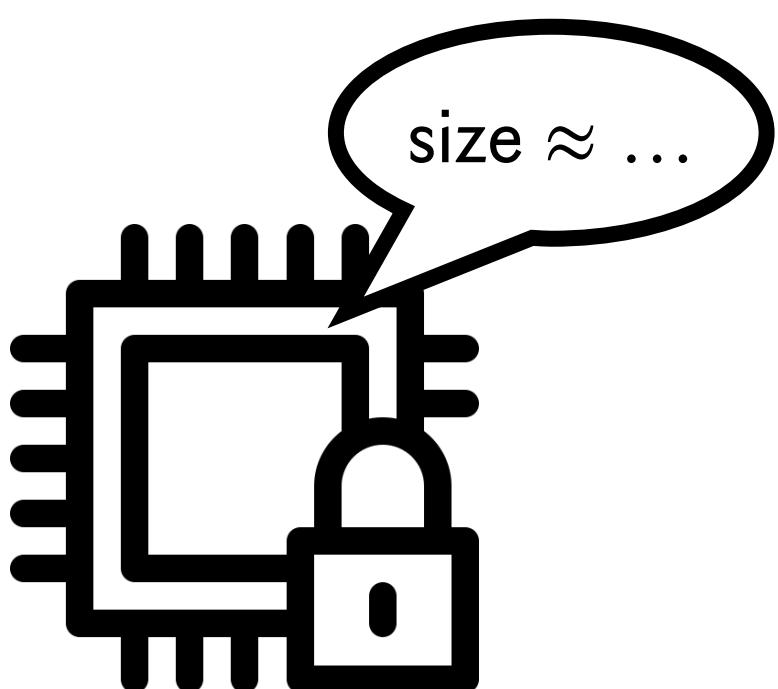
# Prior Approaches to Collaborative Analytics



## Peer-to-peer setting

e.g., Conclave (EuroSys'19),  
Senate (Security'21), SMCQL (VLDB'17),  
Shrinkwrap (VLDB'18)

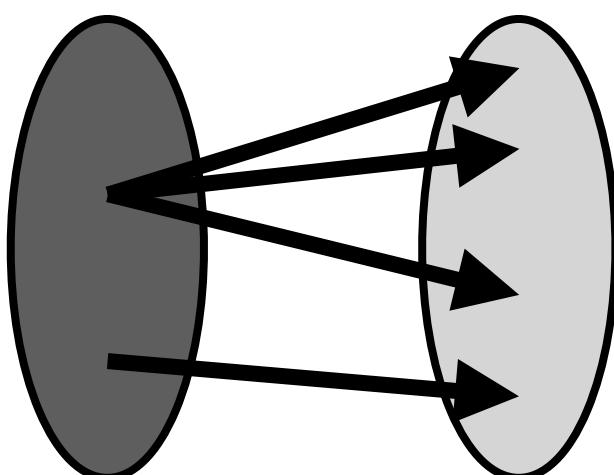
✗ Doesn't scale with more parties



## Leakage or upper bounds

e.g., Opaque (NSDI'17), Obliviator (Security'25),  
SAQE (VLDB'20)

✗ Only provides partial security



## Individual join operators

e.g., [MRR] (CCS'20), [BDG+] (CCS'22),  
[AHK+] (CCS'23)

✗ Cannot evaluate queries

*Can we efficiently evaluate complex analytics  
without compromising security?*

*Can we efficiently evaluate complex analytics  
without compromising security?*

**Yes!**

*Can we efficiently evaluate complex analytics  
without compromising security?*

**Yes!**

We are the first to show **multiple joins**  
with no security compromises

...and run **all queries** from prior work,  
with **orders-of-magnitude** larger inputs

# What do real-world queries look like?

# What do real-world queries look like?

We collect all **31 queries** from related work in MPC analytics

# What do real-world queries look like?

We collect all **31 queries** from related work in MPC analytics

All have output size bound  $O(n)$  instead of  $O(n^k)$ !

Why?

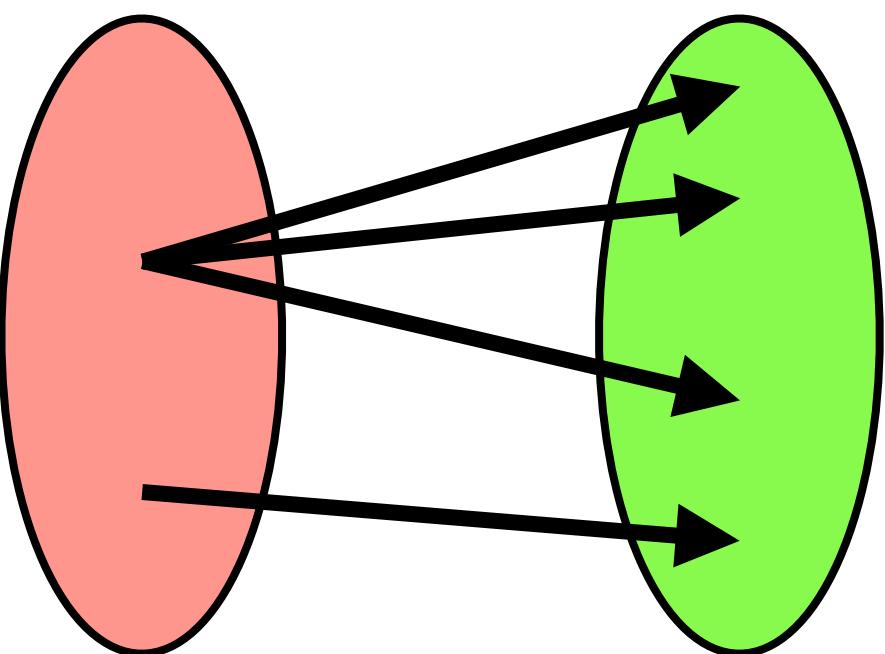
# What do real-world queries look like?

We collect all **31 queries** from related work in MPC analytics

All have output size bound  $O(n)$  instead of  $O(n^k)$ !

Why?

one-to-many



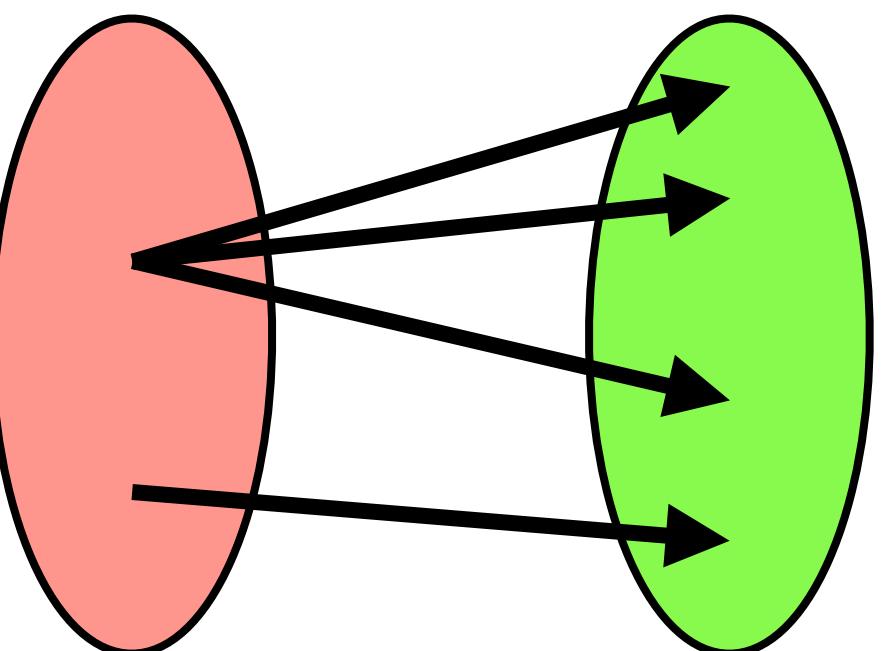
# What do real-world queries look like?

We collect all **31 queries** from related work in MPC analytics

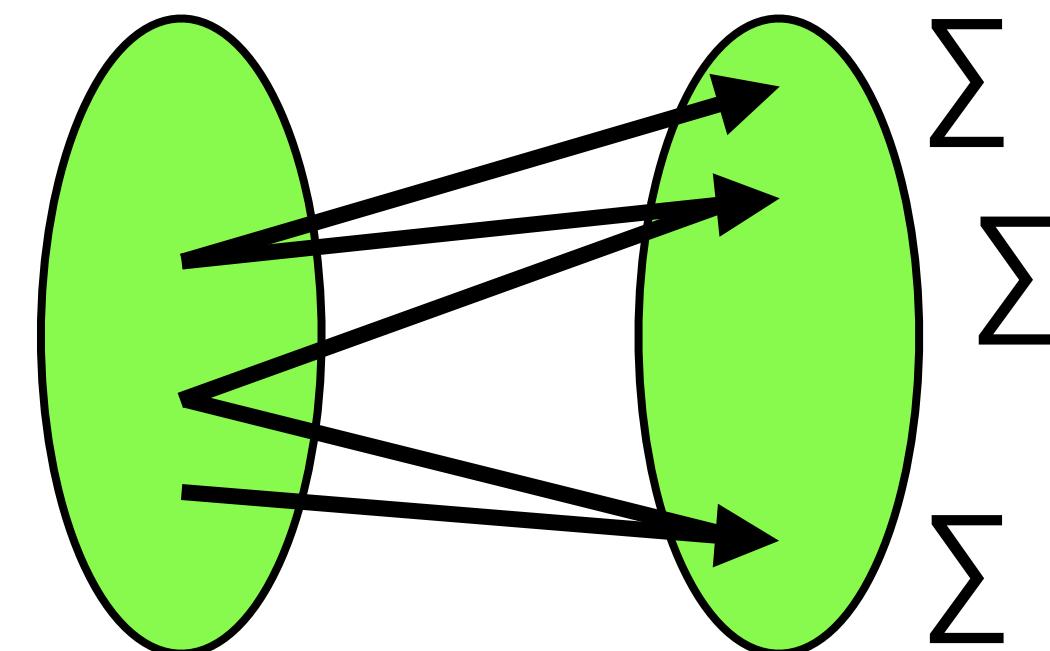
All have output size bound  $O(n)$  instead of  $O(n^k)$ !

Why?

one-to-many



many-to-many + aggregation



# What do real-world queries look like?

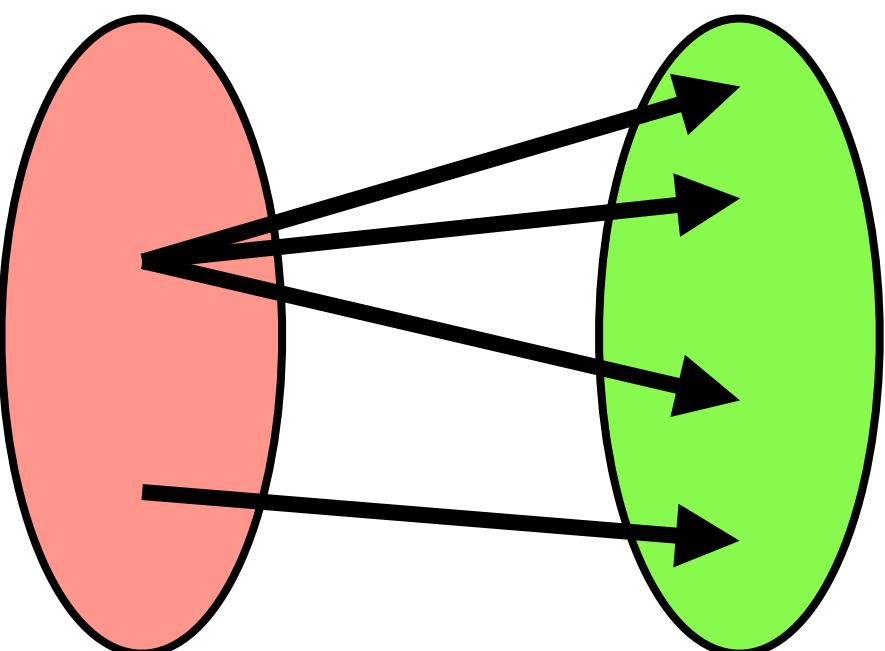
We collect all **31 queries** from related work in MPC analytics

All have output size bound  $O(n)$  instead of  $O(n^k)$ !

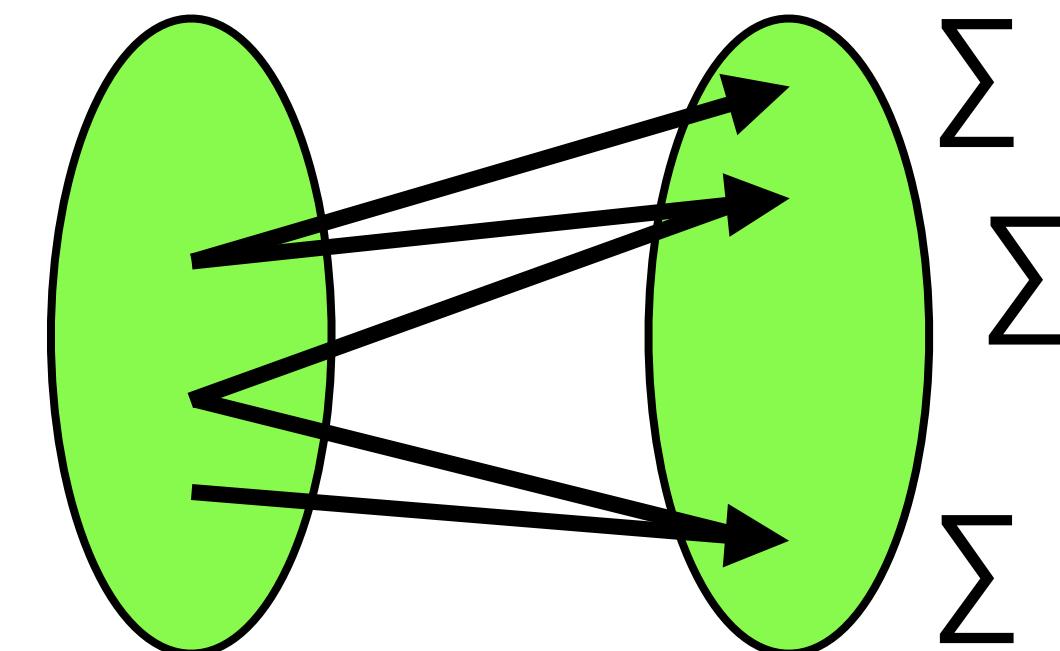
Why?

decomposable

one-to-many



many-to-many + aggregation



# Our Insight

# Our Insight

The general problem is impractically difficult

# Our Insight

**The general problem is impractically difficult**

**Practical instances are not!**

# Our Insight

The general problem is impractically difficult

Practical instances are not!

Common query semantics impose an  
 $O(n)$  bound on the output

# Our Insight

The general problem is impractically difficult

Practical instances are not!

Common query semantics impose an  
 $O(n)$  bound on the output

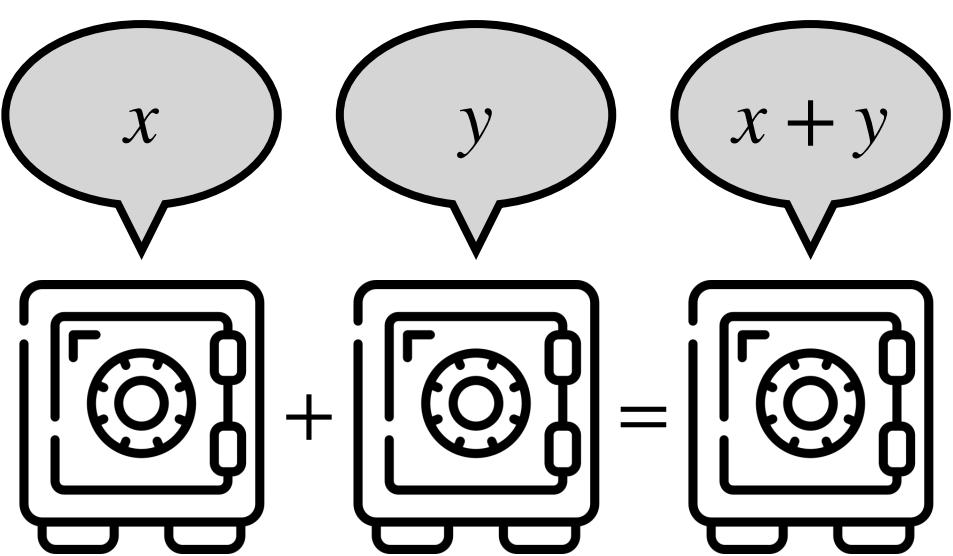
Decompose joins into sort + aggregation

# Our Insight

The general problem is impractically difficult

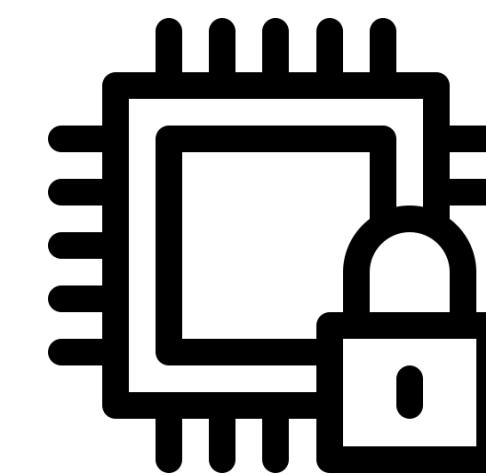
Practical instances are not!

Common query semantics impose an  
 $O(n)$  bound on the output



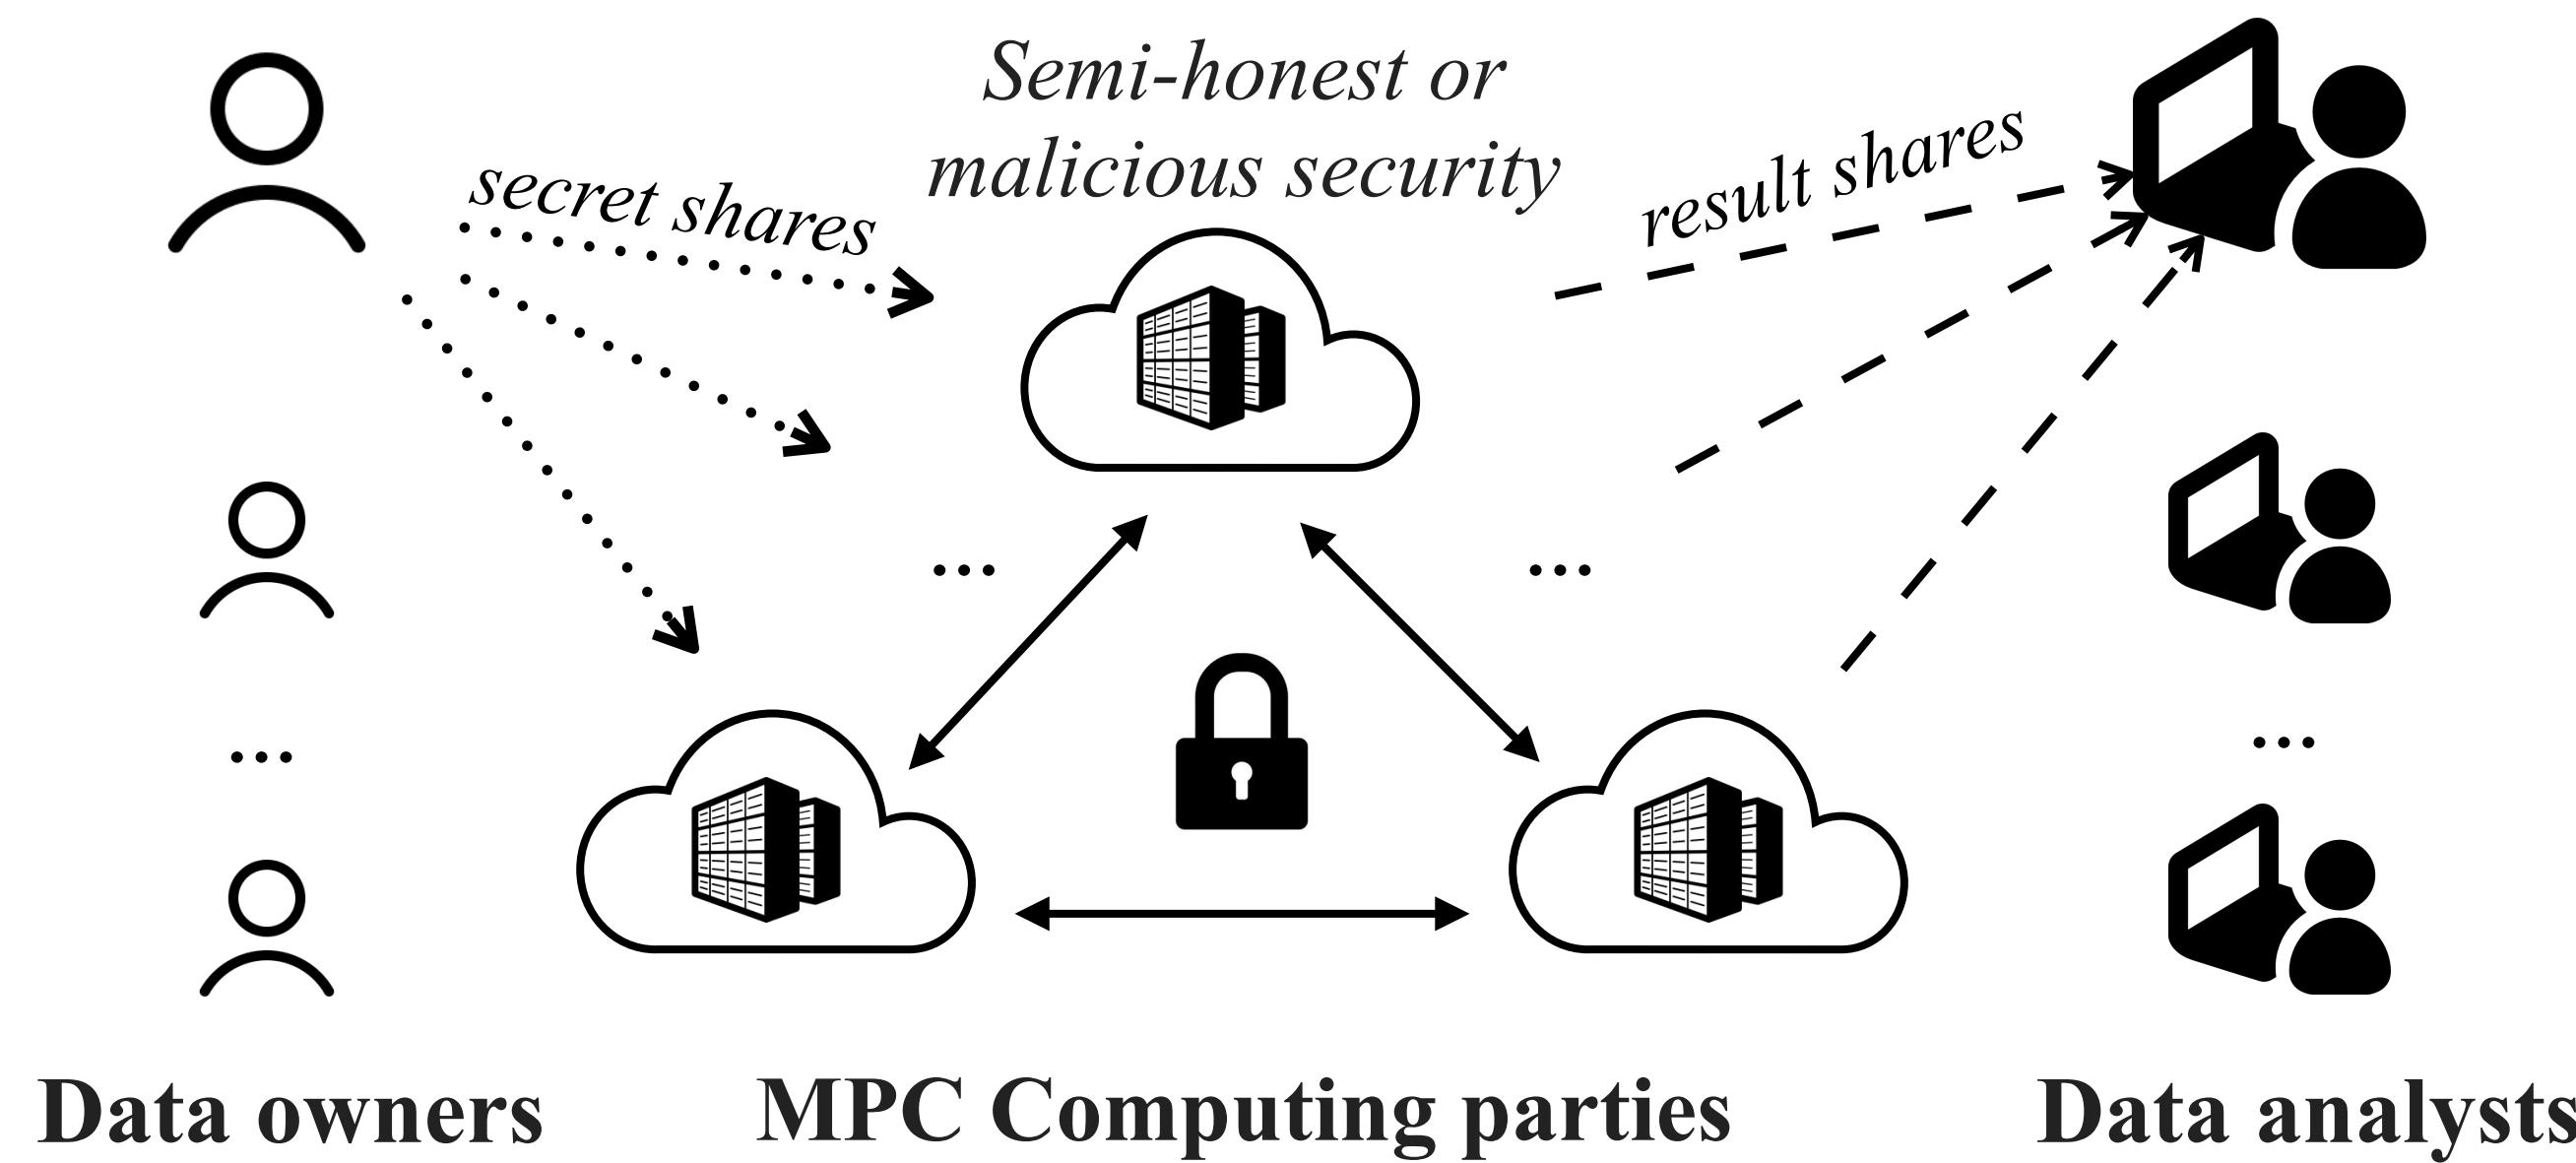
Decompose joins into sort + aggregation

Applies beyond multiparty computation!



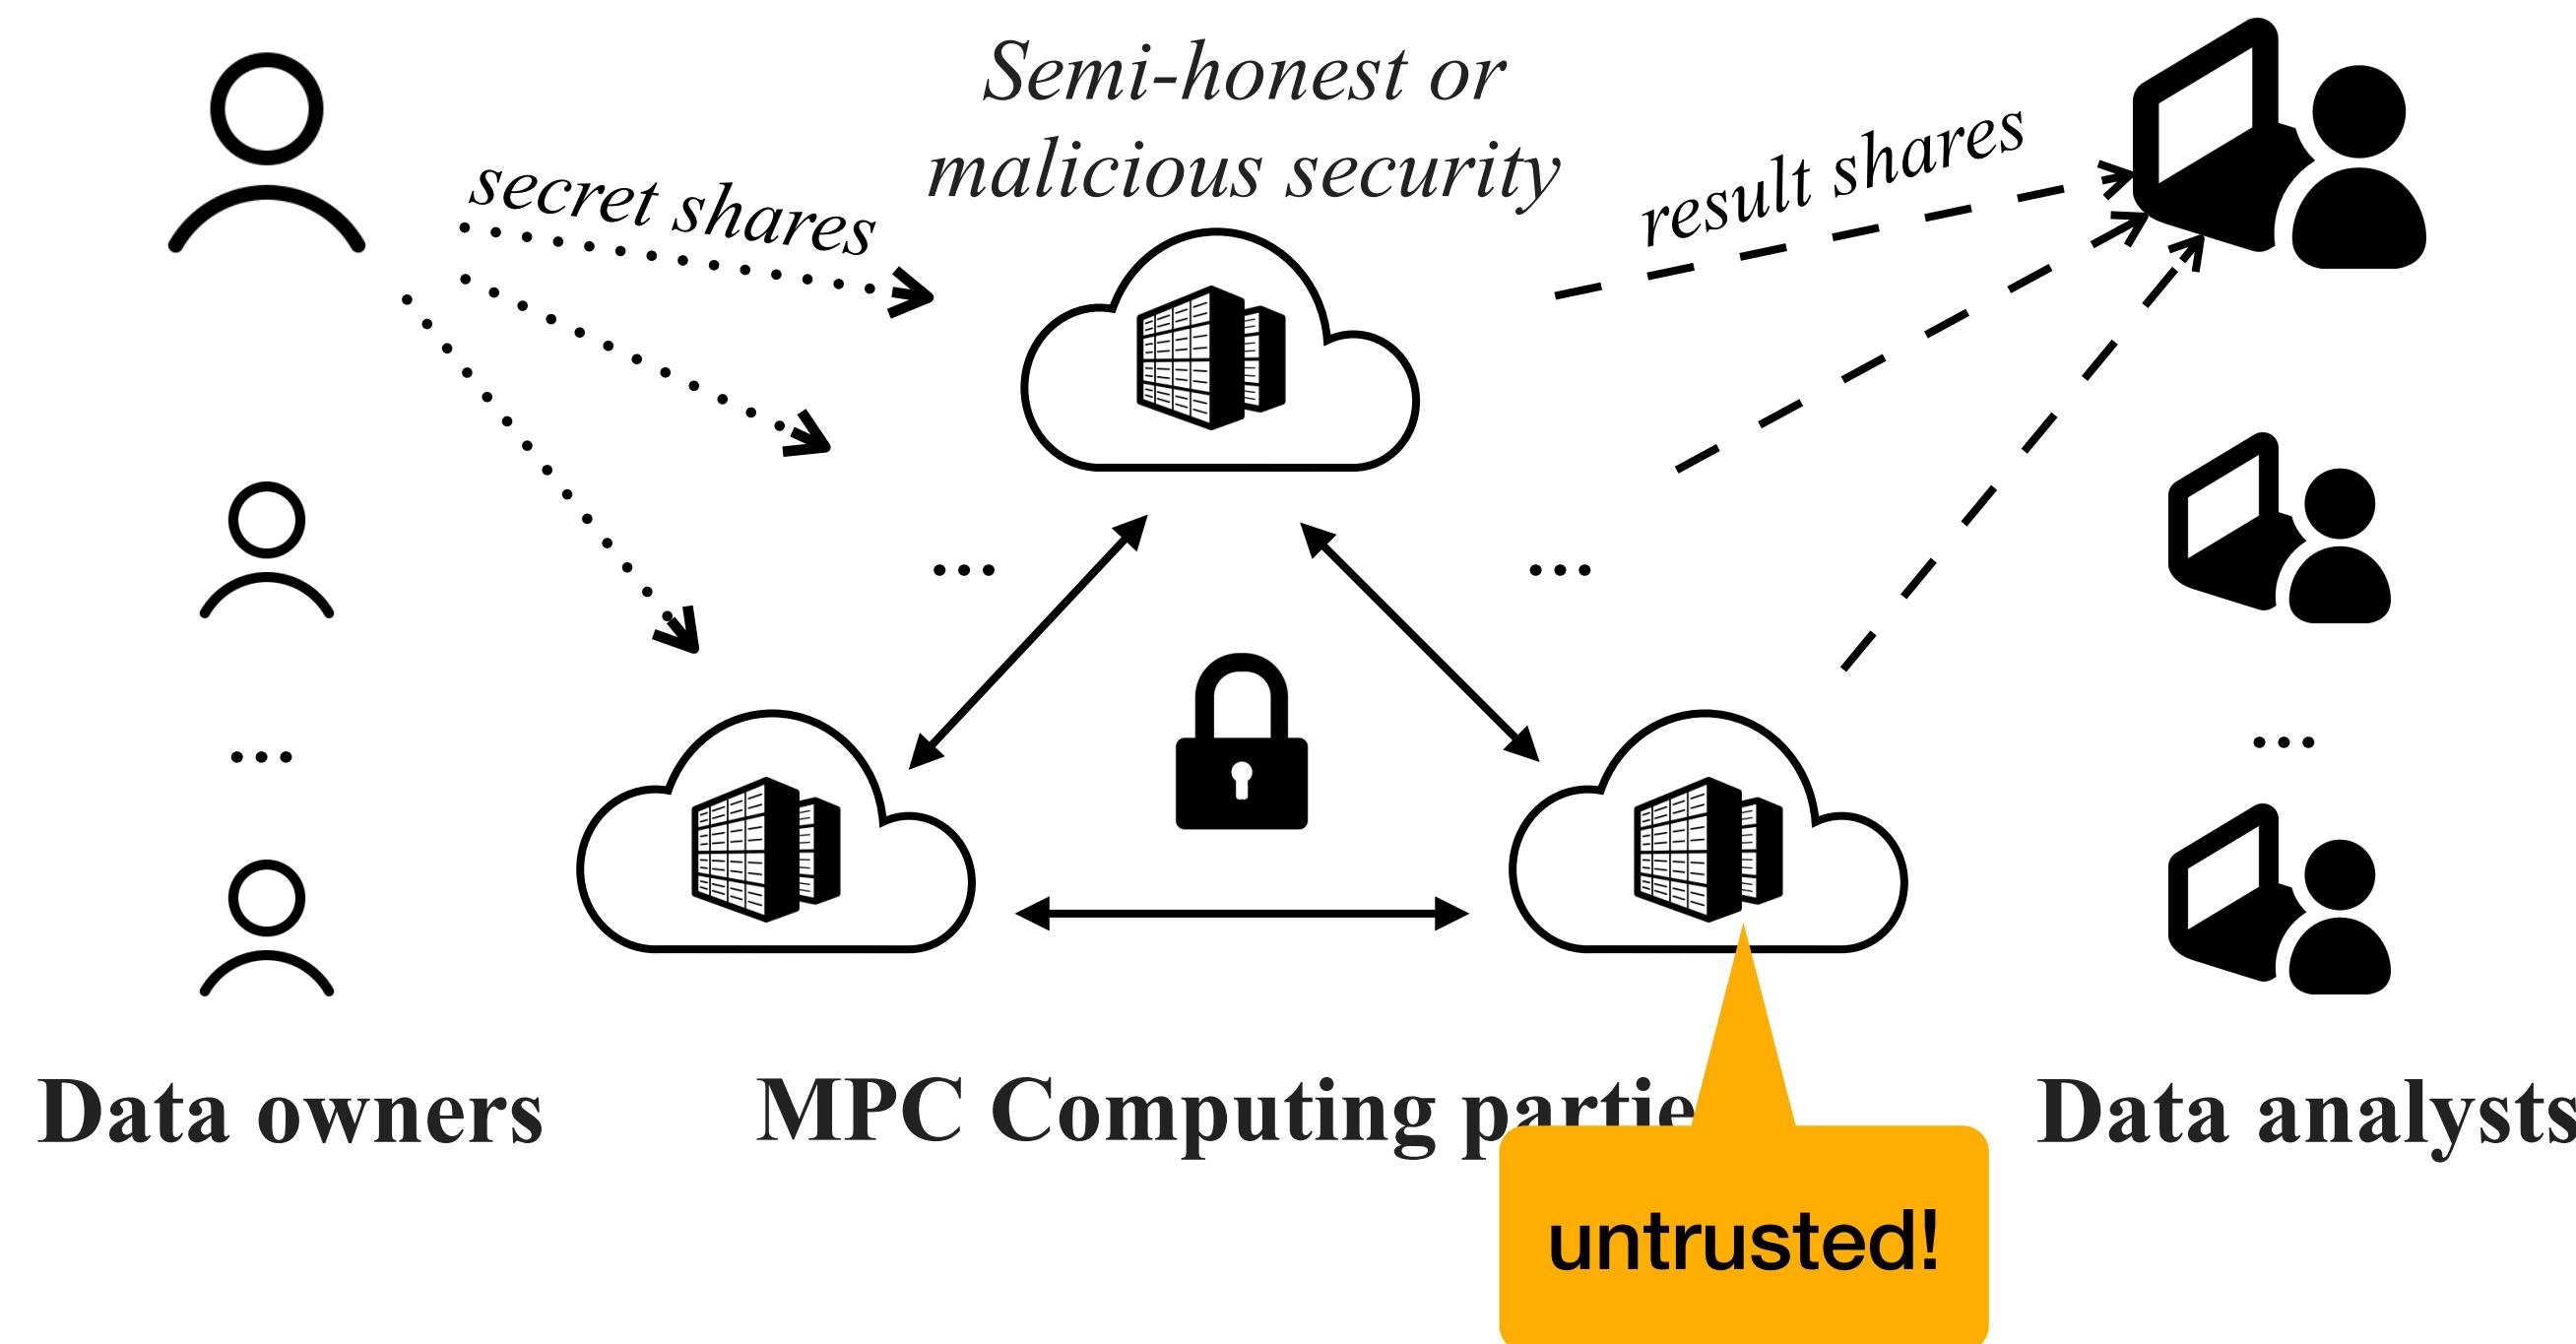
# Secure Multiparty Computation (MPC)

## Outsourced Setting



# Secure Multiparty Computation (MPC)

## Outsourced Setting



# Secret Sharing: Data Owners



ID	Balance
1	\$100
4	\$50
9	\$25
7	-\$13



ID	Balance
3	\$814
2	\$312
7	\$60
6	-\$366



ID	Balance
14	\$82
12	\$14
0	-\$8
3	\$160

# Secret Sharing: Data Owners



Pr	ID	Balance
N1	1799JtRkU	NV9MVqRf
nk	qT	qicpdFpxn ZgaoE7Hx
R	FV	JtRkUmc5 FVwp7XTF
P	2g	qSiJ3nBG vhnCMlo7j

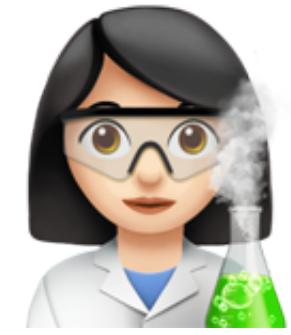


CB	ID	Balance
j8k	D	UW6m9tx EABqfgpRi
3jf	ay	m6vQFTgb 9u19eWyV
0tv	Vy	qhq2JFbe hM2hOg8
uR	MQeNGgyI	q1uon8xM

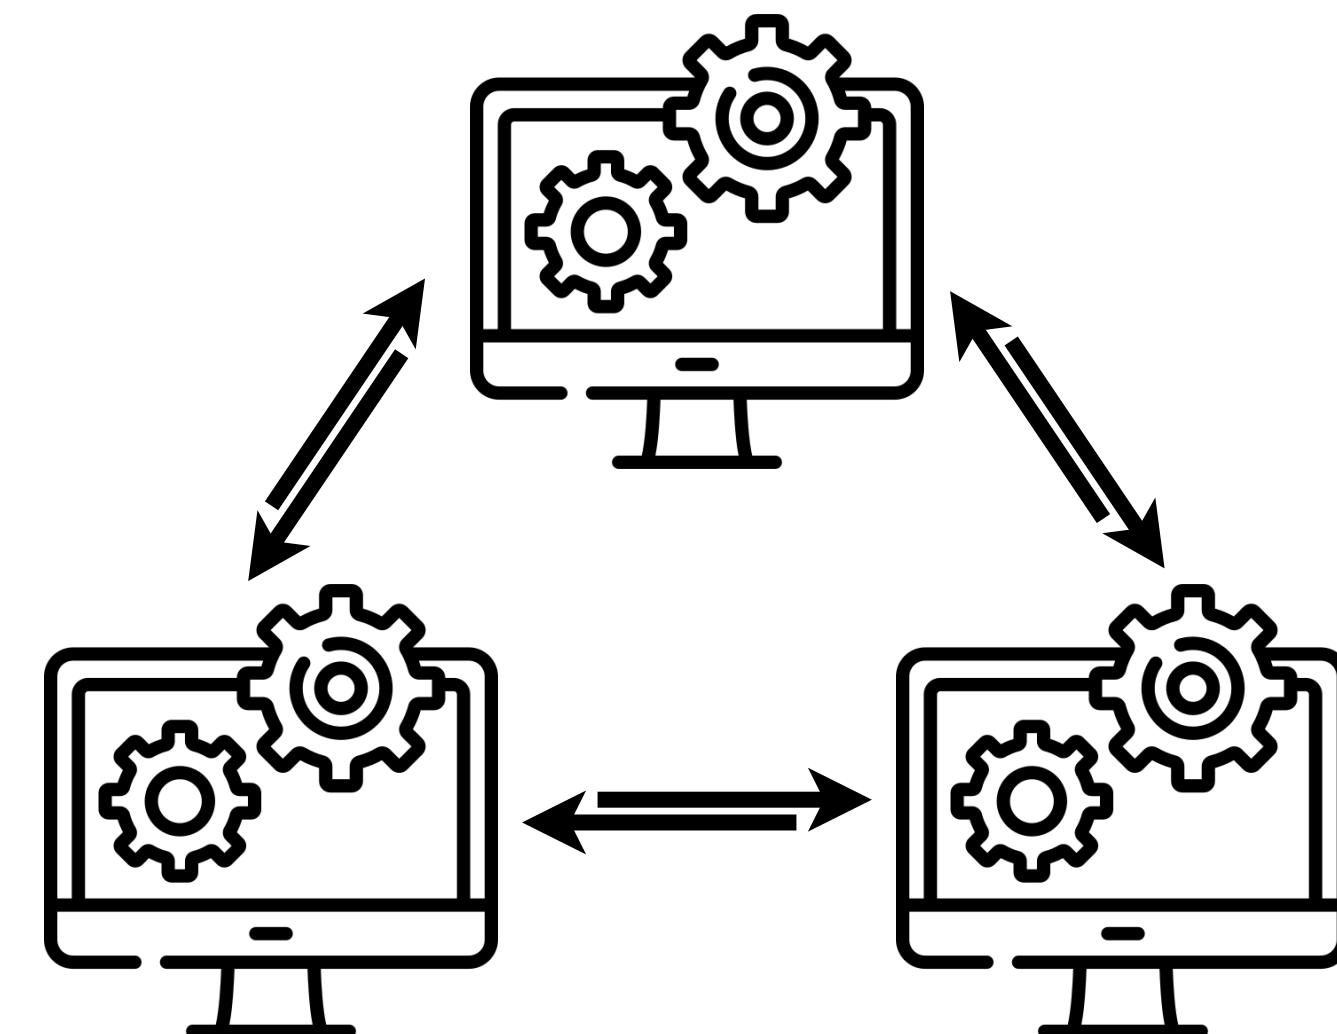


Q8	ID	Balance
MN	as	hsIcAgWLy OUPgb2V
PQ	91	MVAJi4Vnt FzI0j2hH8
ajC	4n	8KsbVnN YlaNhg5py
am	ehbzAtm6j	LmWp415

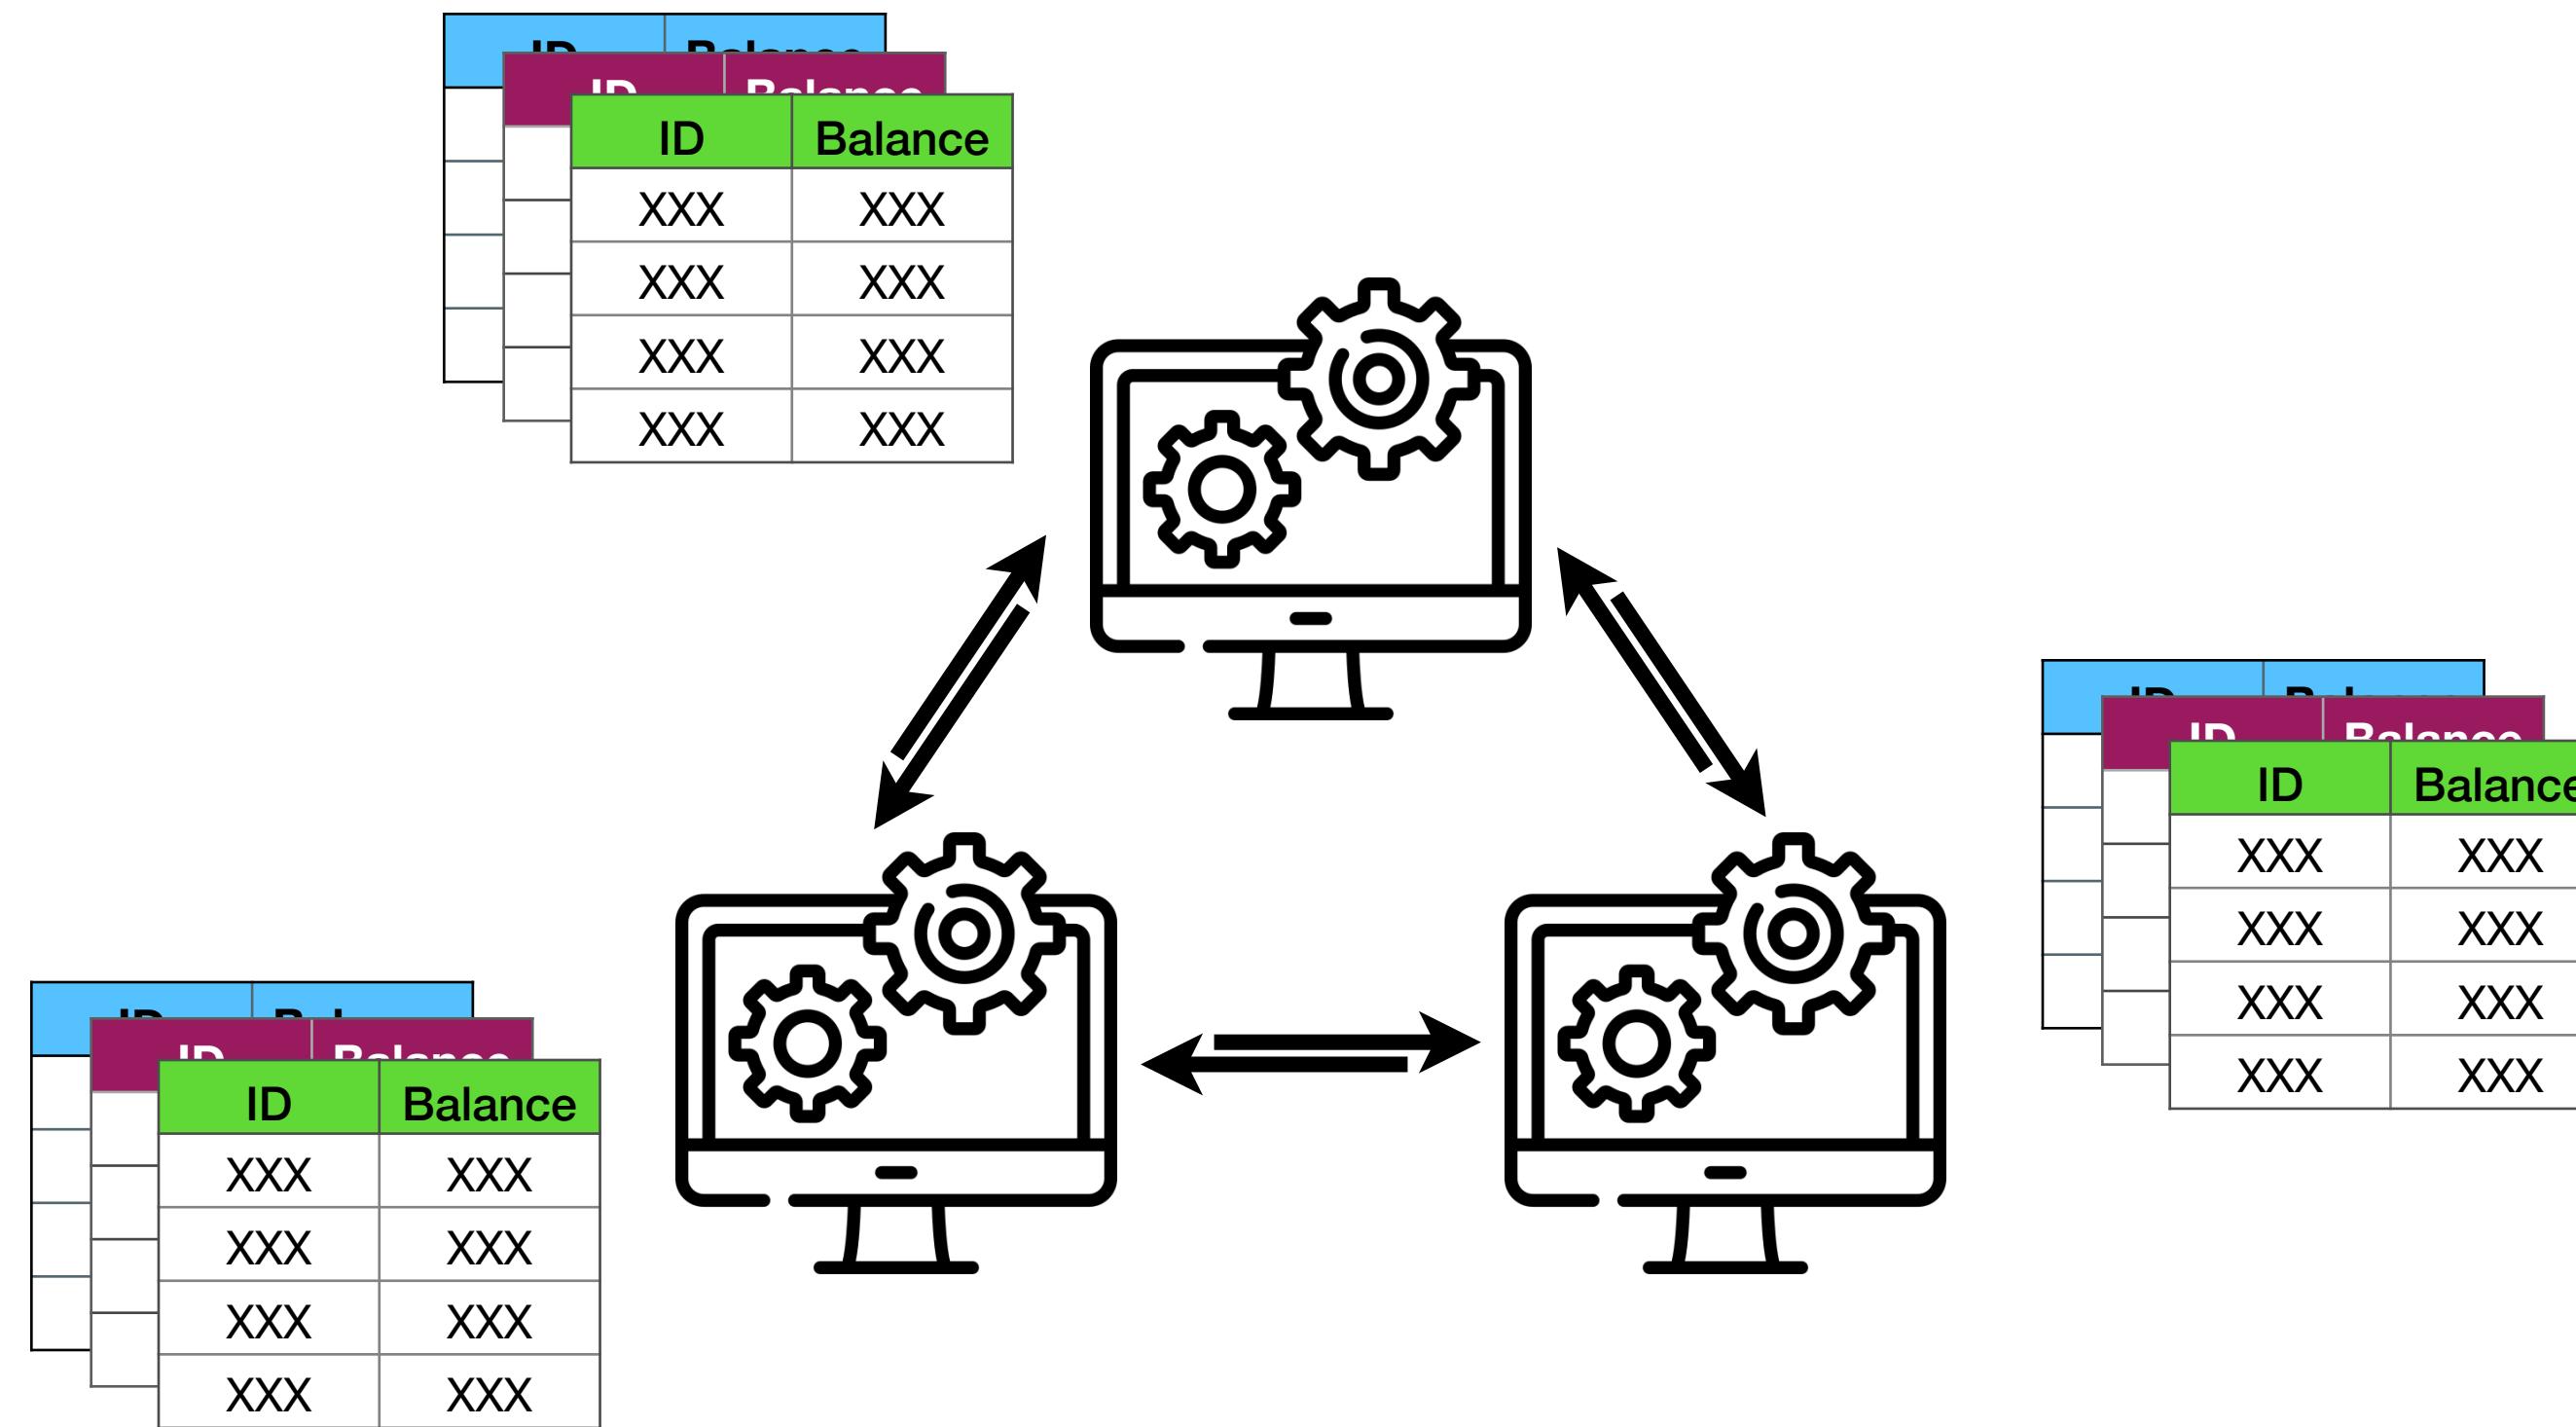
# Secret Sharing: Data Owners



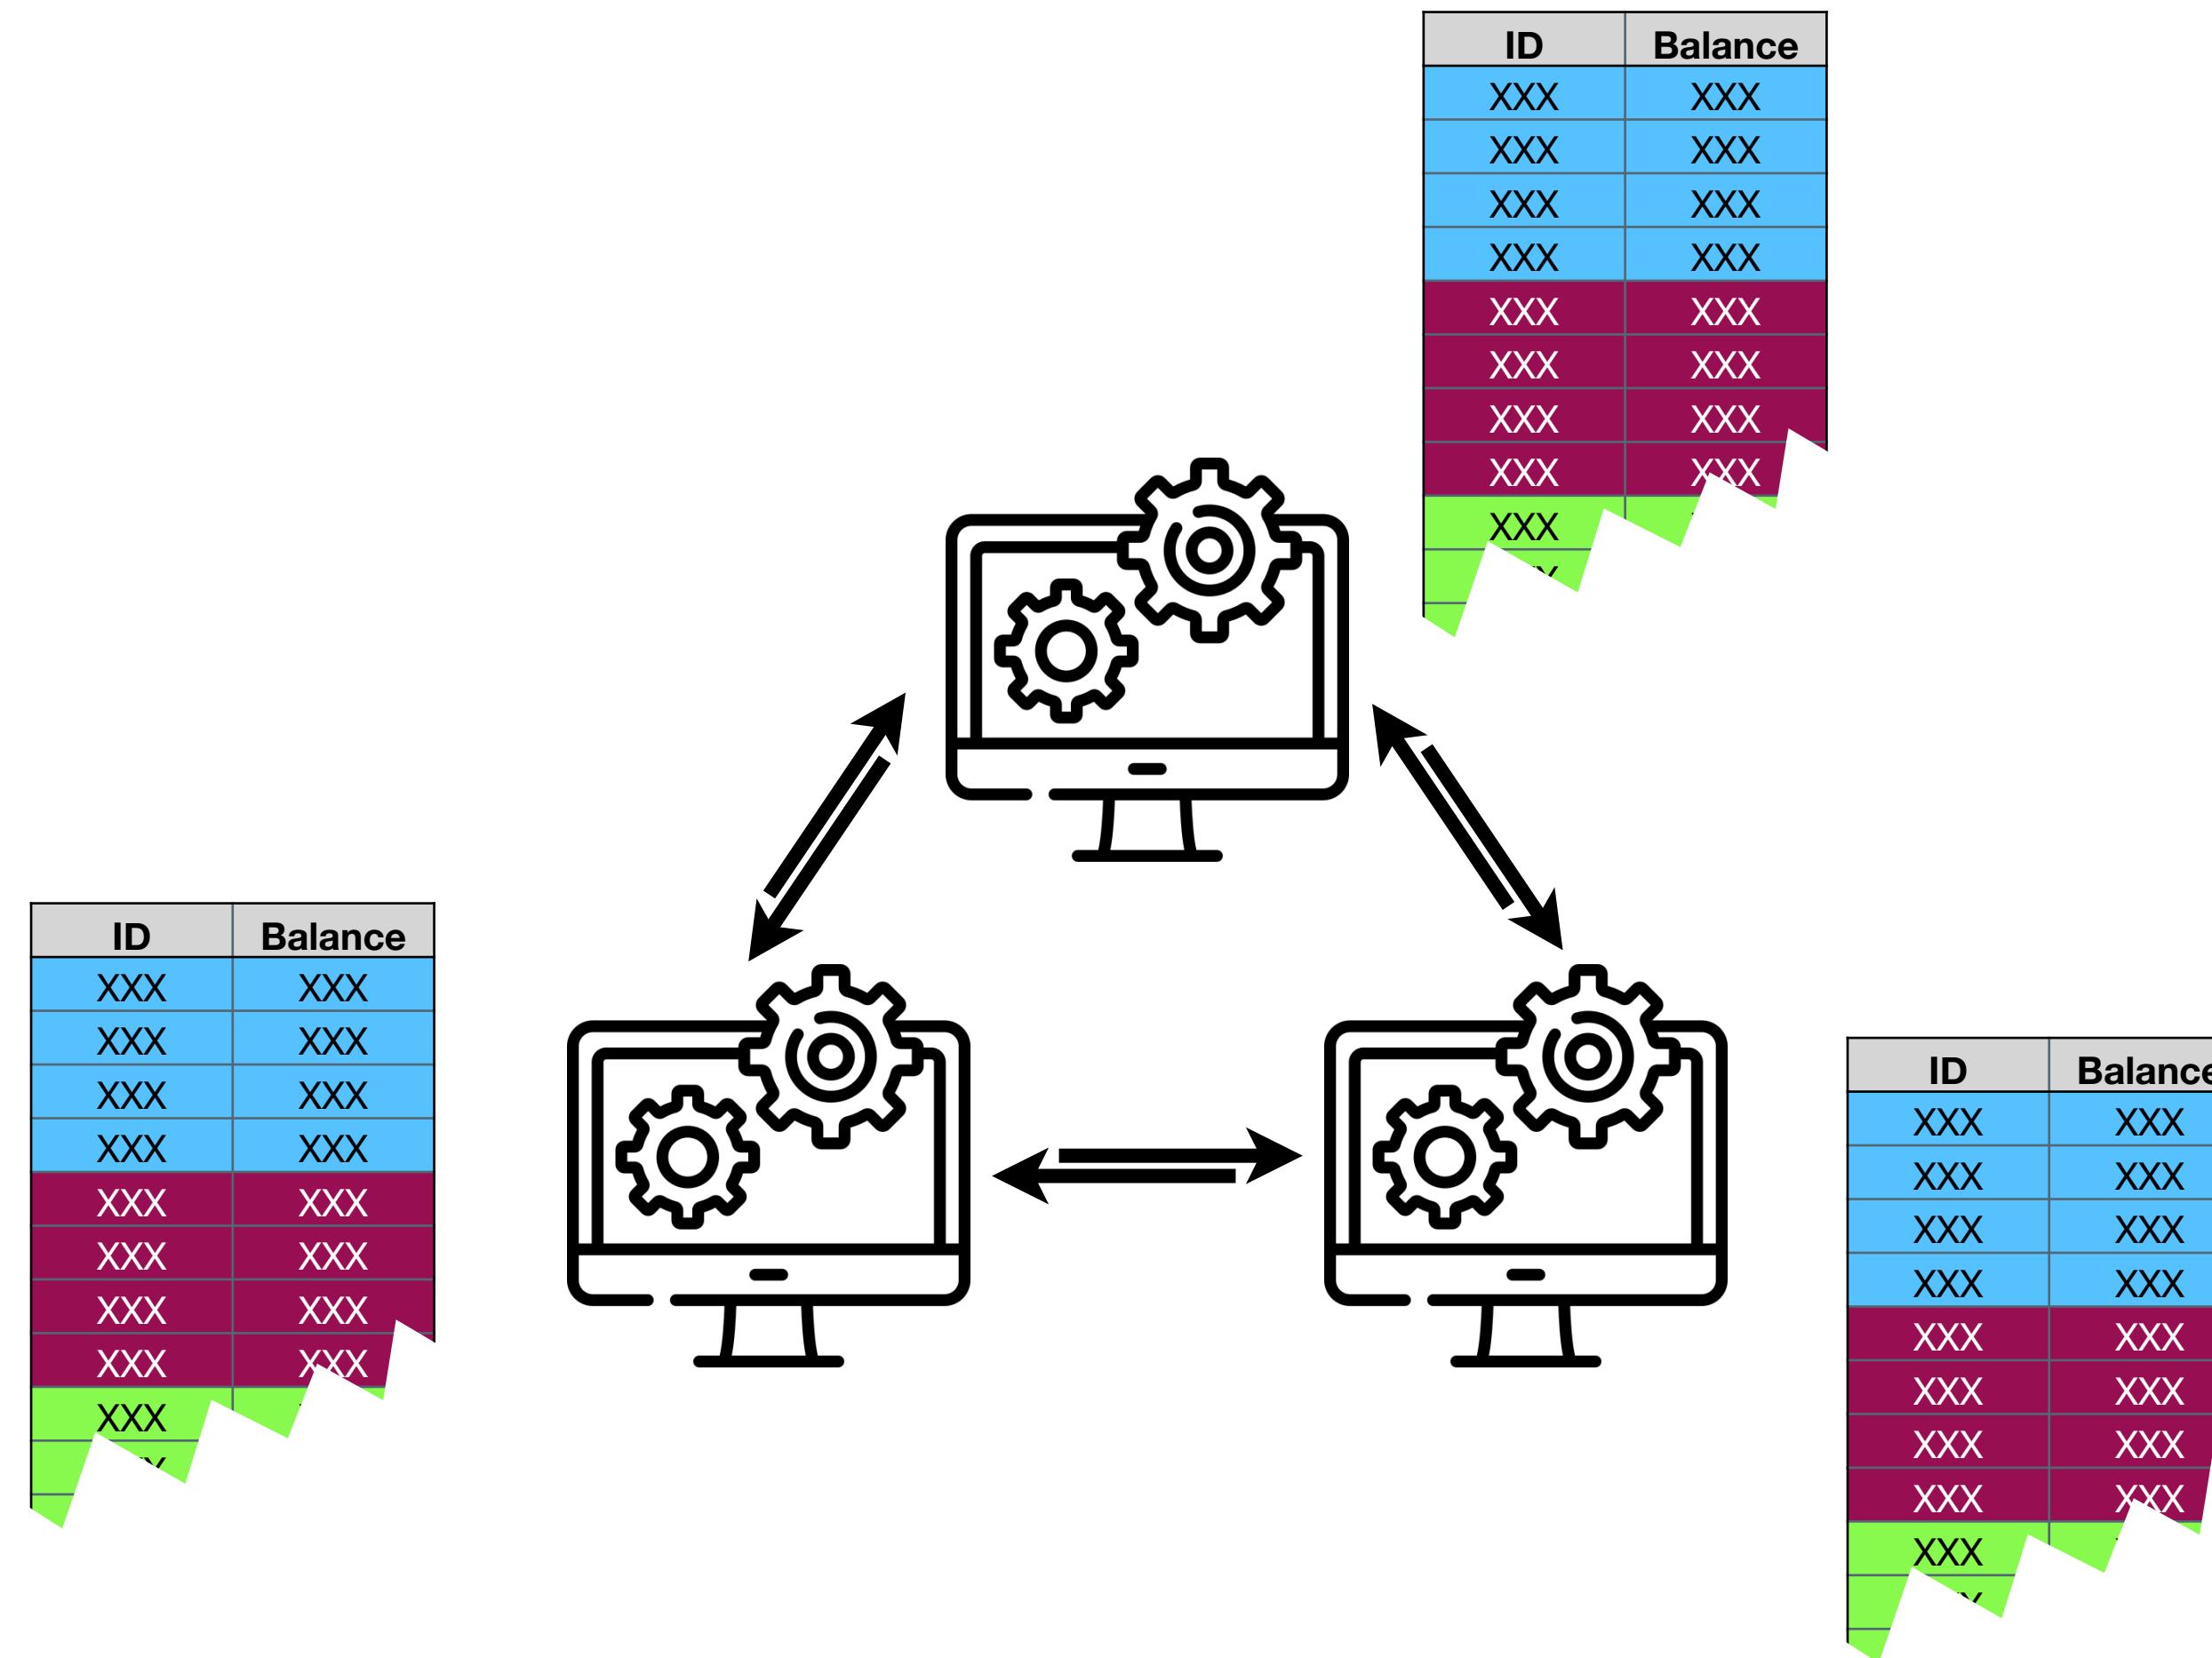
# Secret Sharing: Computing Parties



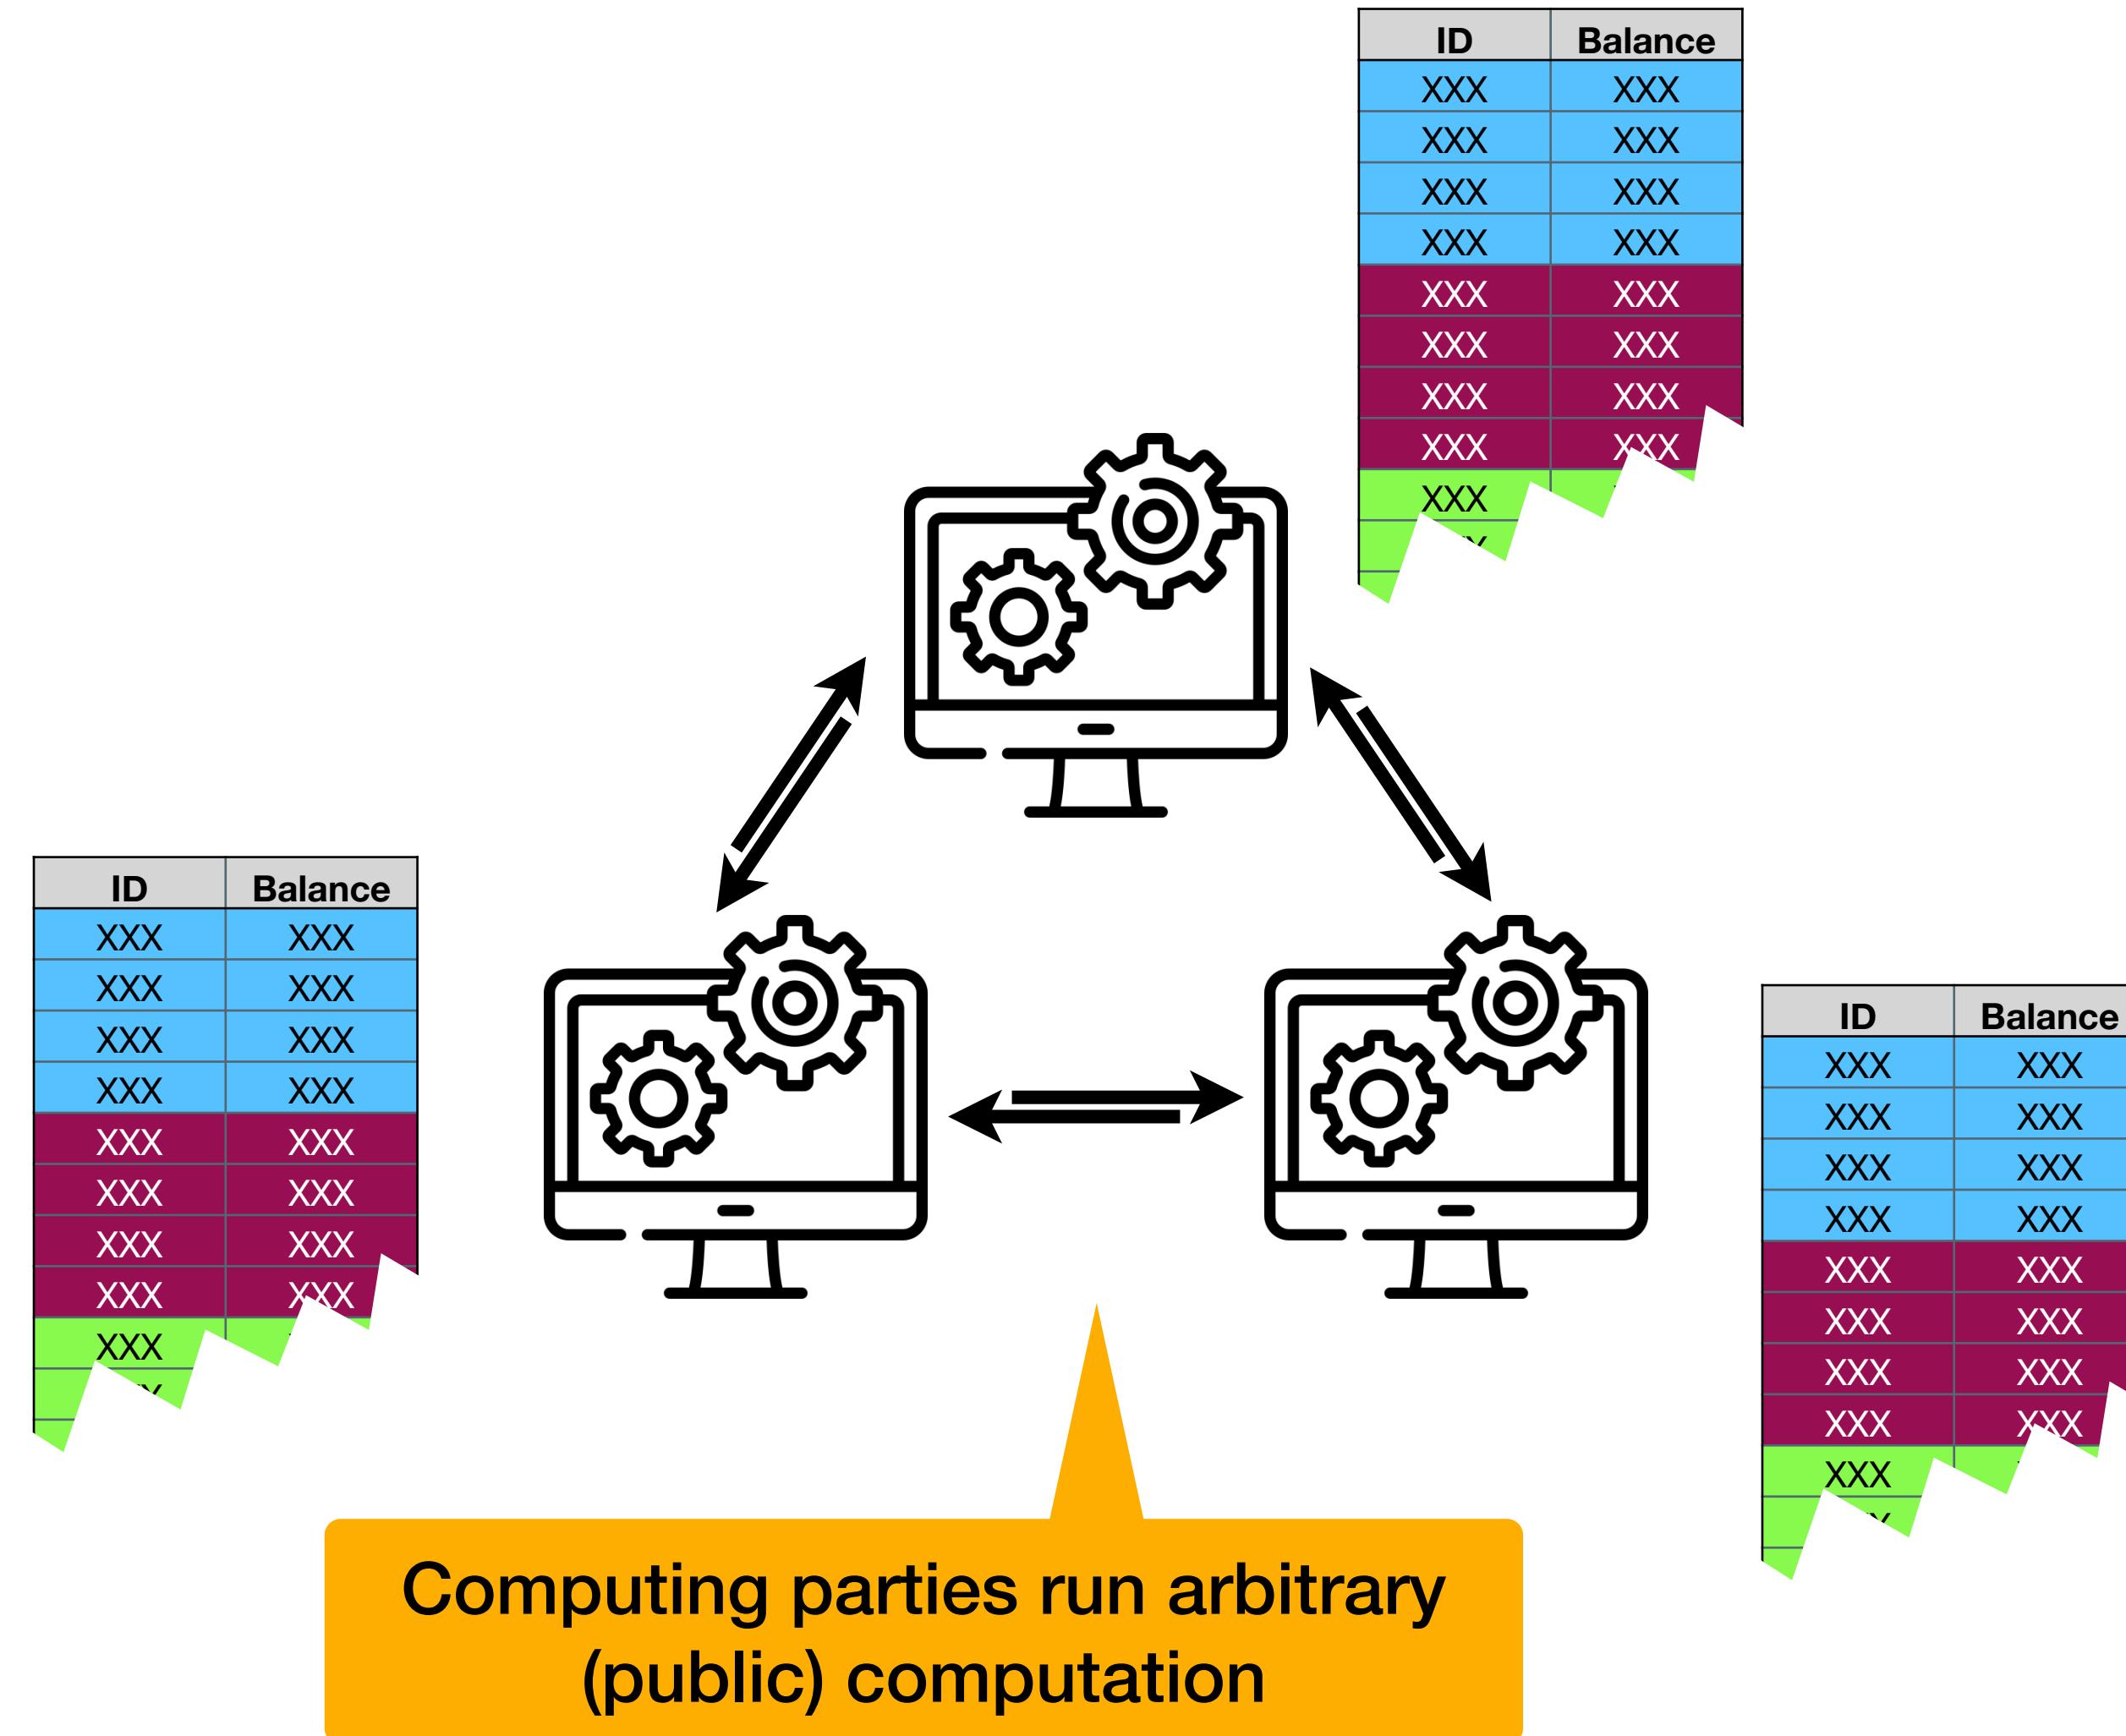
# Secret Sharing: Computing Parties



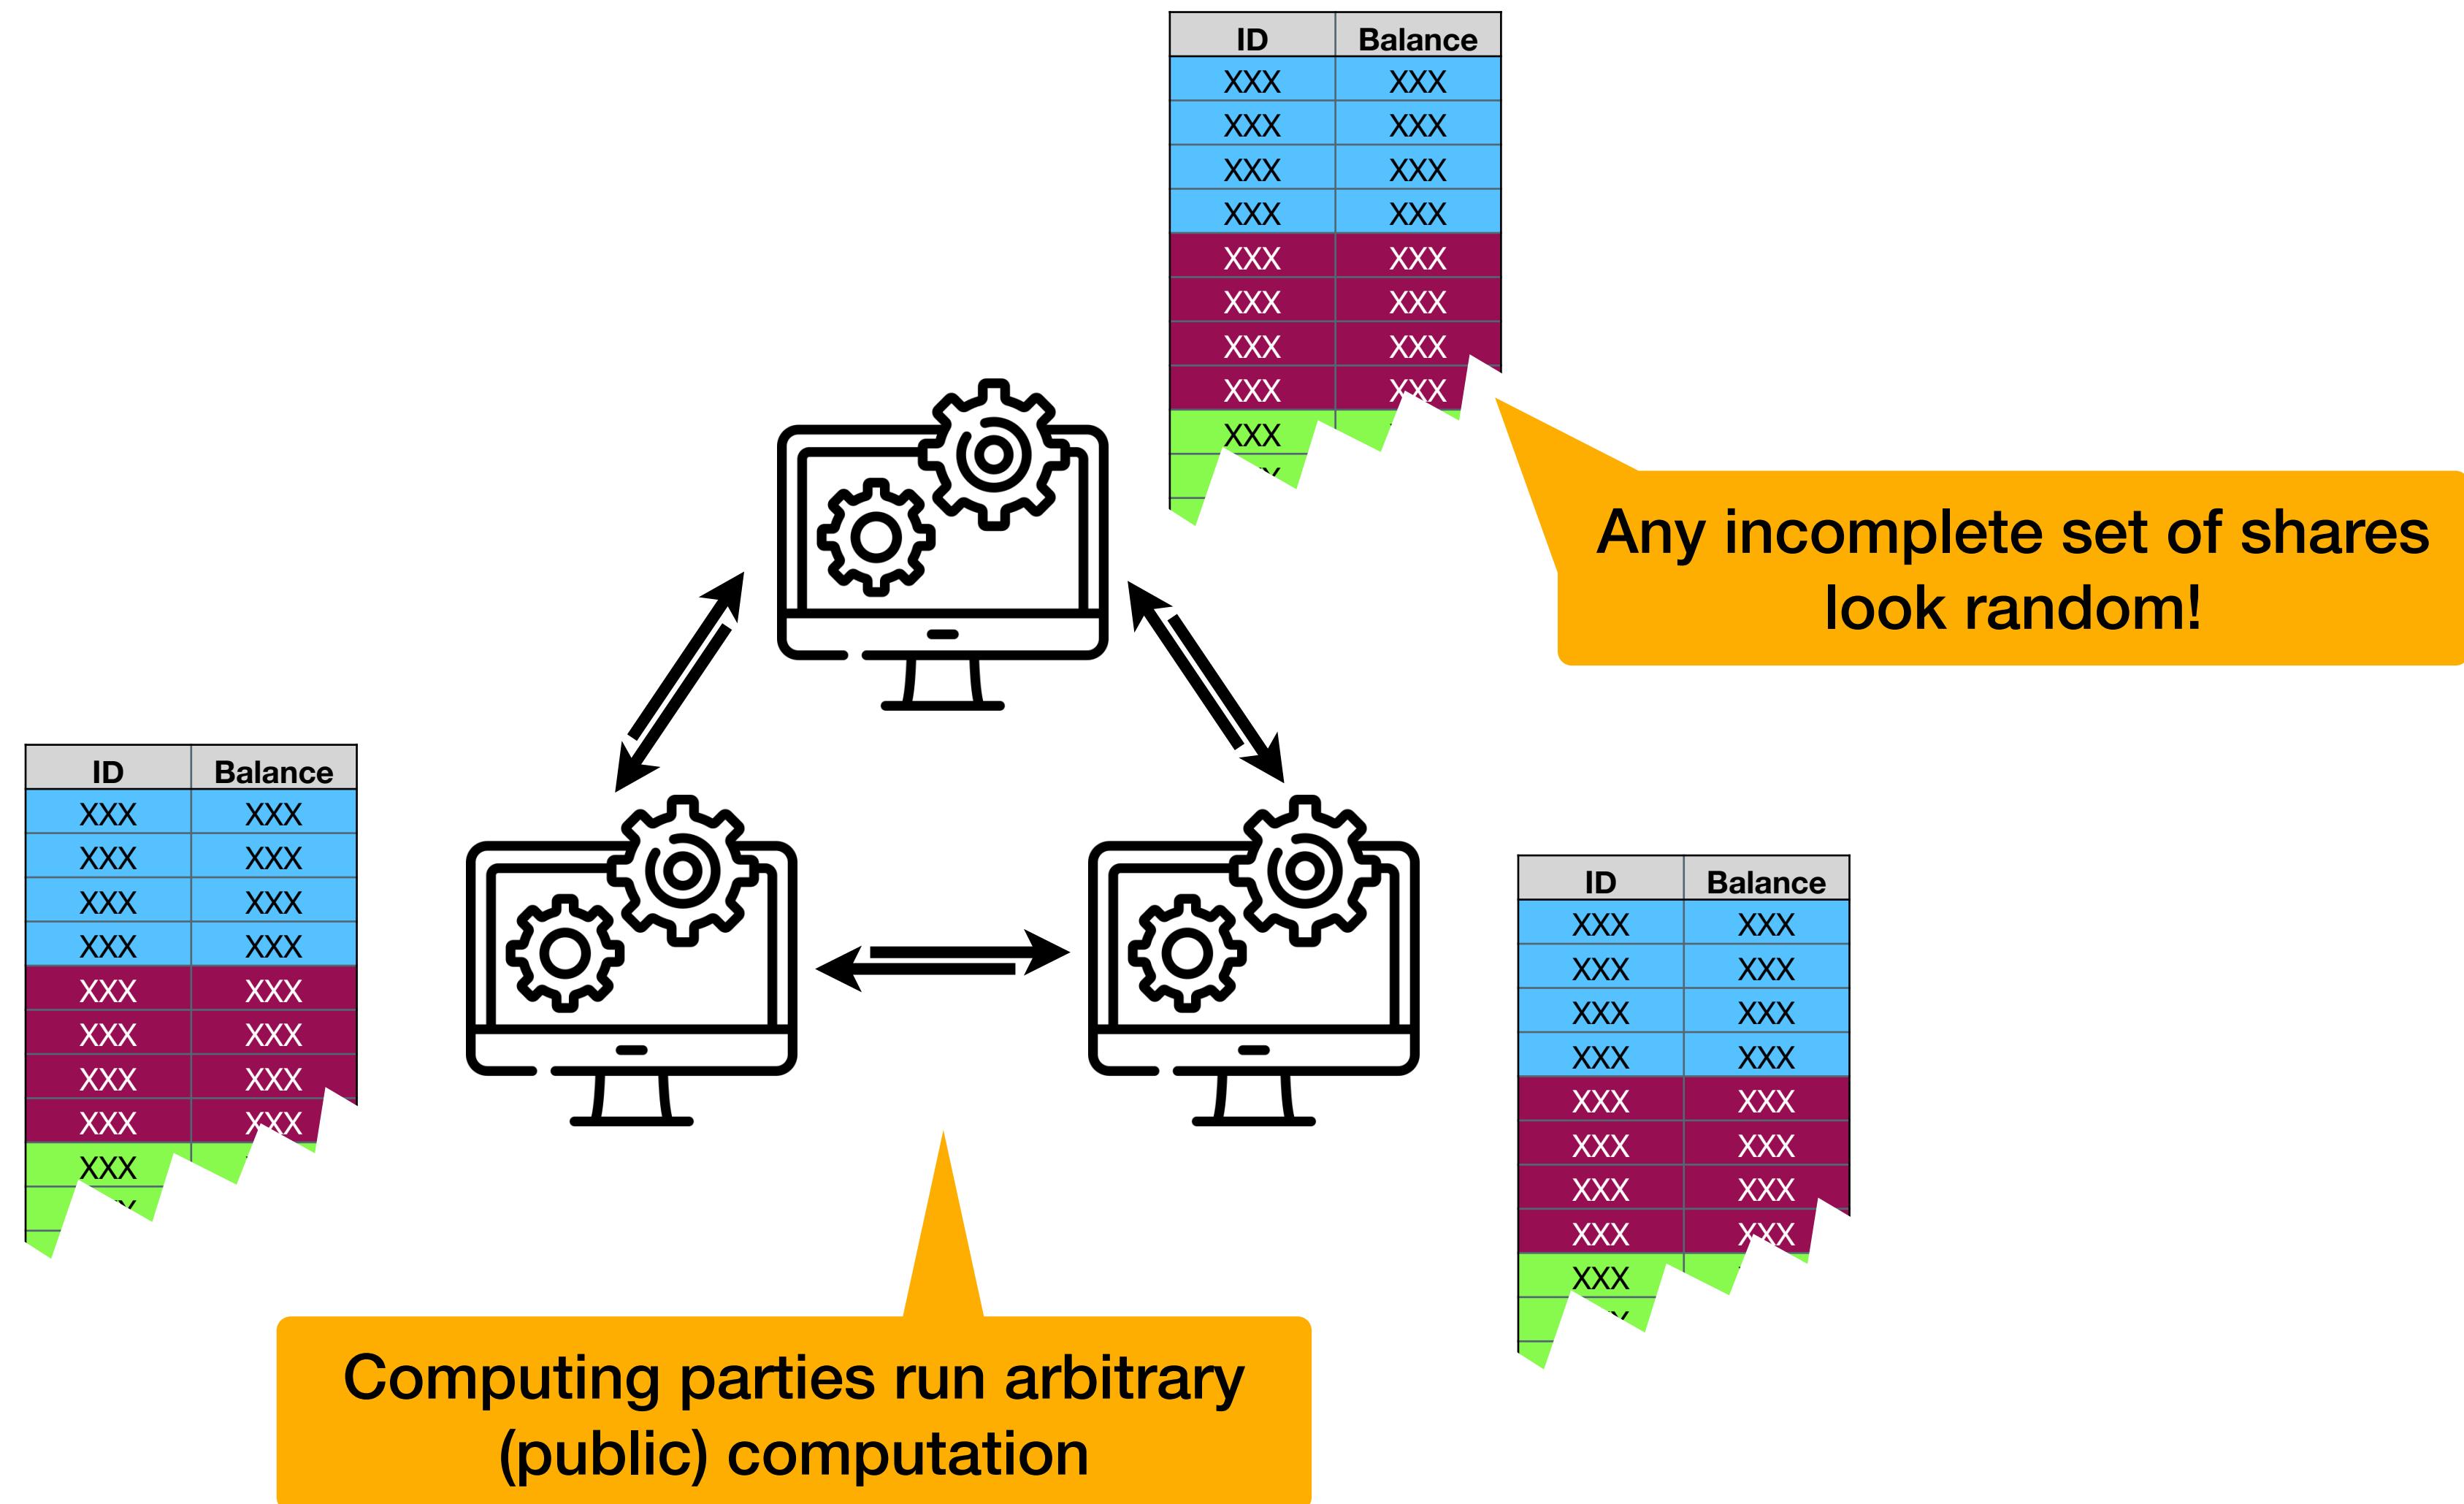
# Secret Sharing: Computing Parties



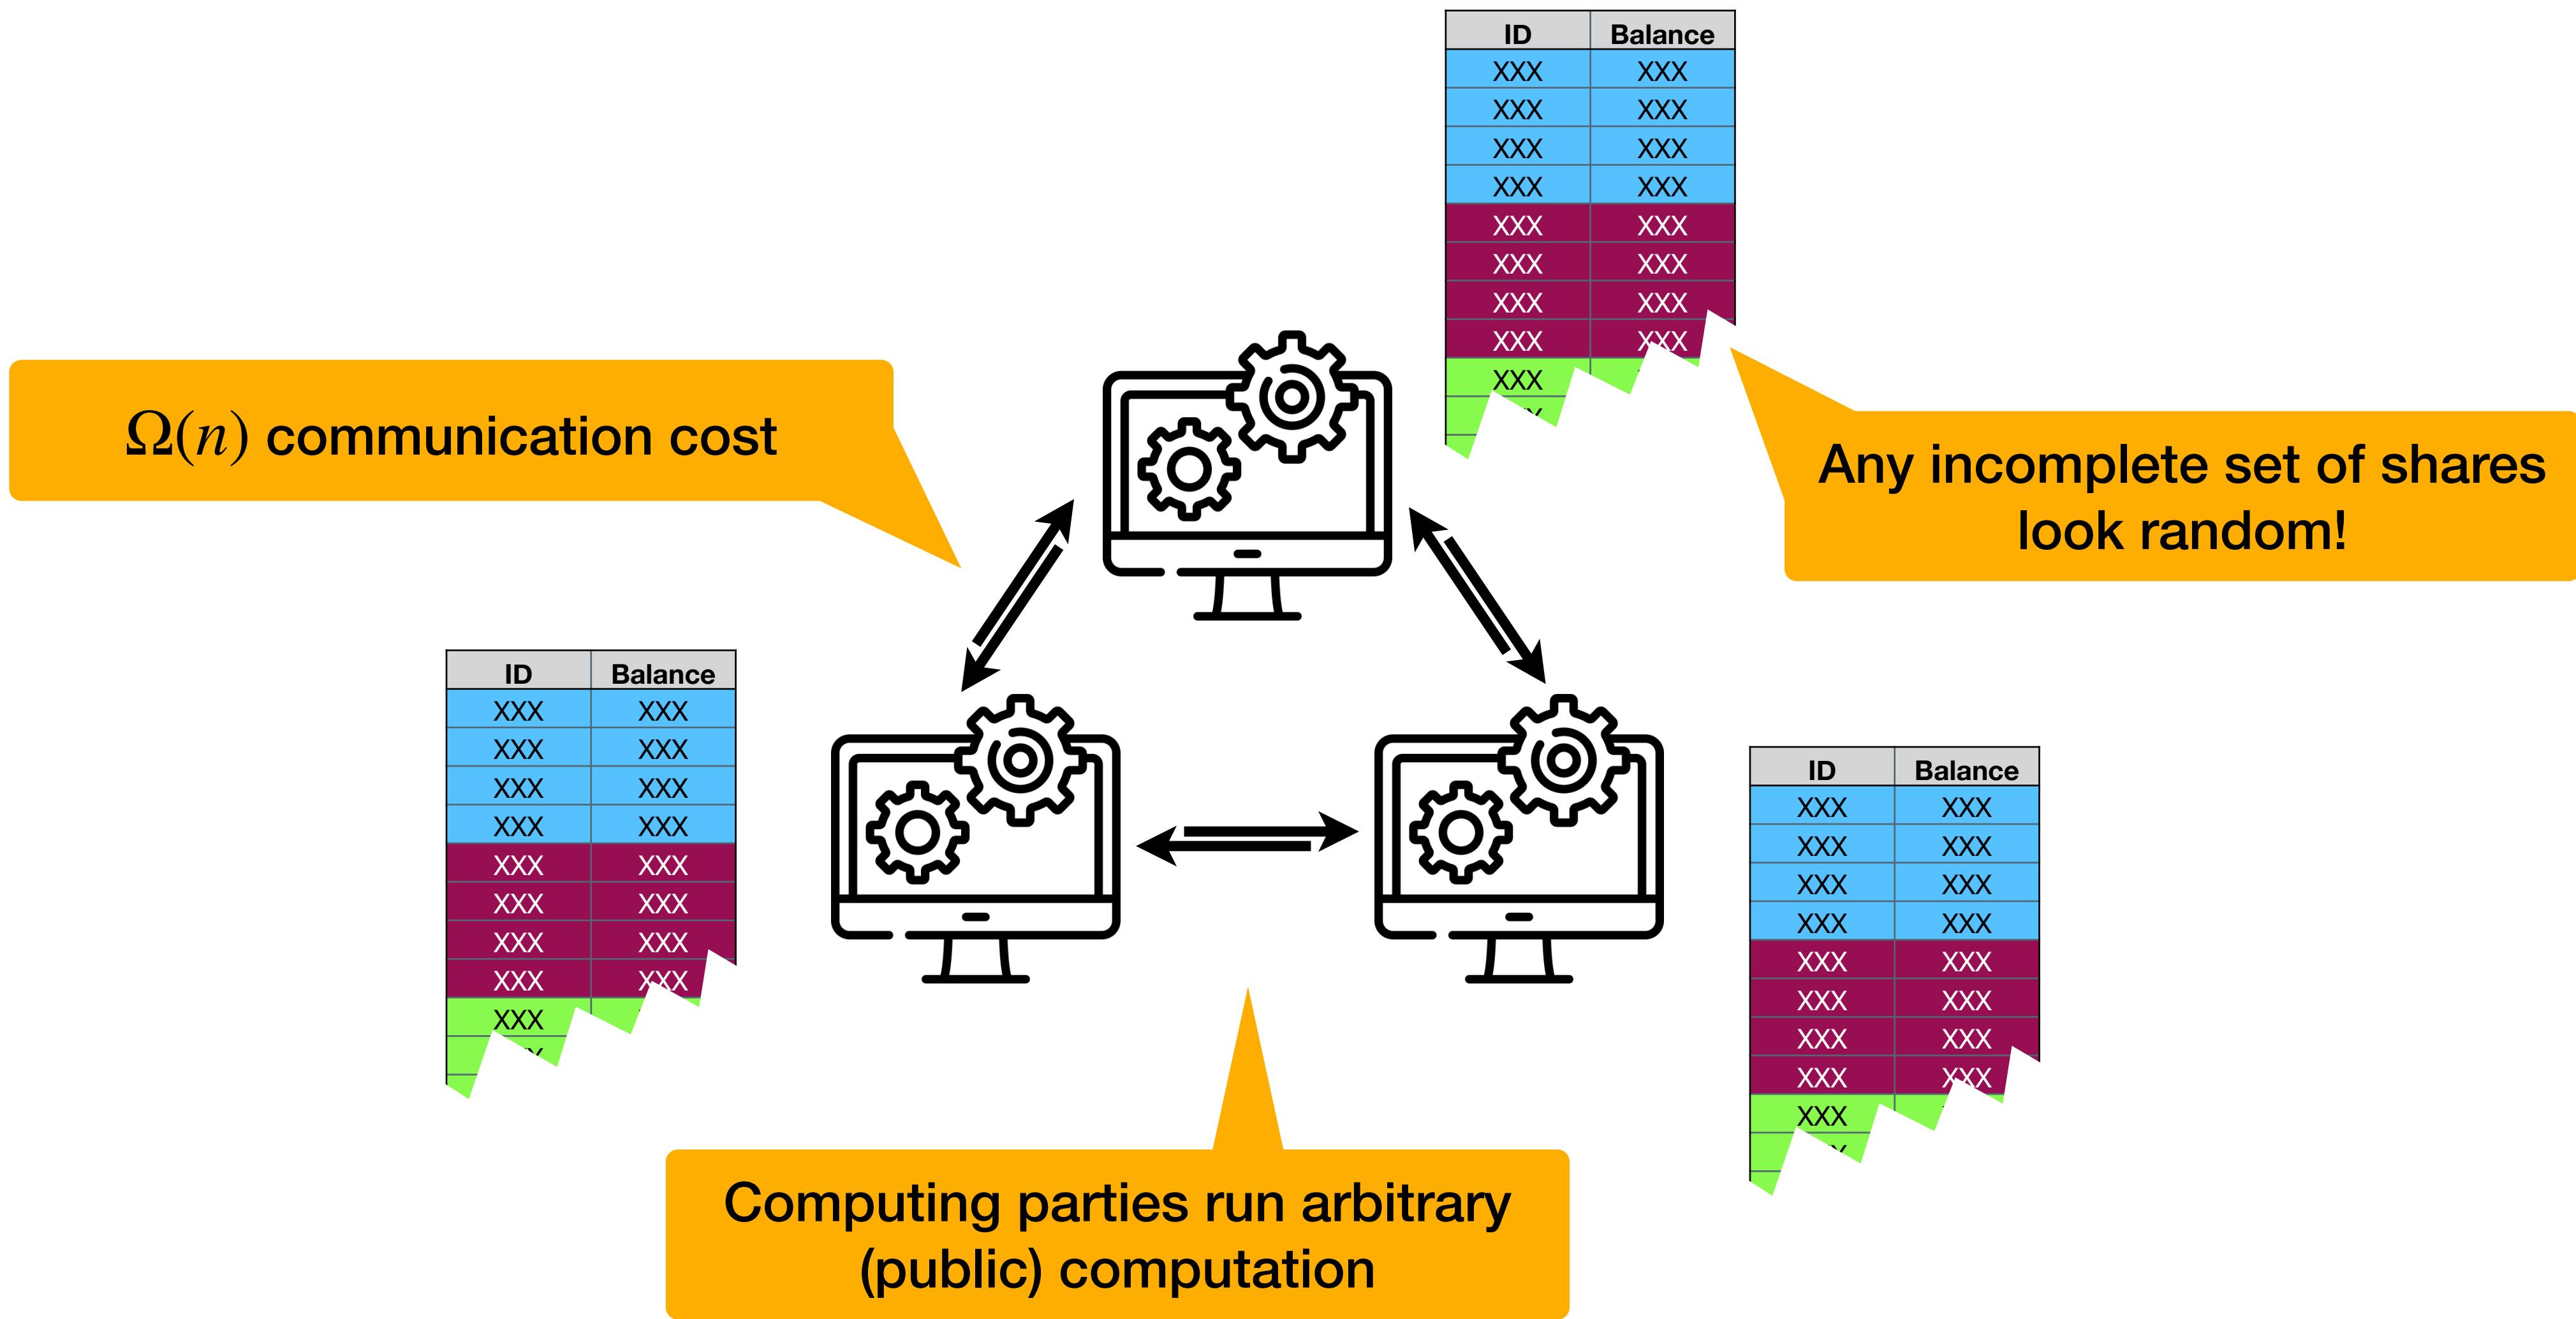
# Secret Sharing: Computing Parties



# Secret Sharing: Computing Parties



# Secret Sharing: Computing Parties



# Secret Shared Tables



Pr	ID	Balance
N1	1799JtRkU	NV9MVqRf
nk	qT	qicpdFpxn ZgaoE7Hx
R	FV	JtRkUmc5 FVwp7XTF
P	2g	qSiJ3nBG vhnCMlo7j

CB	ID	Balance
j8k	Df	UW6m9tx EABqfgpRi
3jf	Otv	m6vQFTgb 9u19eWyV
ay	Vy	qhq2JFbe hM2hOg8
uR	MQeNGgyI	q1uon8xM

Q8	ID	Balance
MN	hsIcAgWLy	OUPgb2V
as	PQ	
91	MVAJi4Vnt	FzI0j2hH8
ajC	8KsbVnN	YlaNhg5py
4n	am.	
ehbzAtm6j	LmWp415	

# Secret Shared Tables



ID	Balance
1	\$100
4	\$50
9	\$25
7	-\$13

ID	Balance
3	\$814
2	\$312
7	\$60
6	-\$366

ID	Balance
14	\$82
12	\$14
0	-\$8
3	\$160

# Improving on Quadratic: Sort & Aggregate

## One table must have unique keys

```
func join(L, R, keys) :  
    T := concat*(L, R)  
    T.sort(keys)  
    T.prefix_aggregate(...)  
    return T
```

# Improving on Quadratic: Sort & Aggregate

## One table must have unique keys

$O(n \log n)$

```
func join(L, R, keys) :  
    T := concat*(L, R)  
    T.sort(keys)  
    T.prefix_aggregate(...)  
    return T
```

$O(n \log n)$

# Improving on Quadratic: Sort & Aggregate

## One table must have unique keys

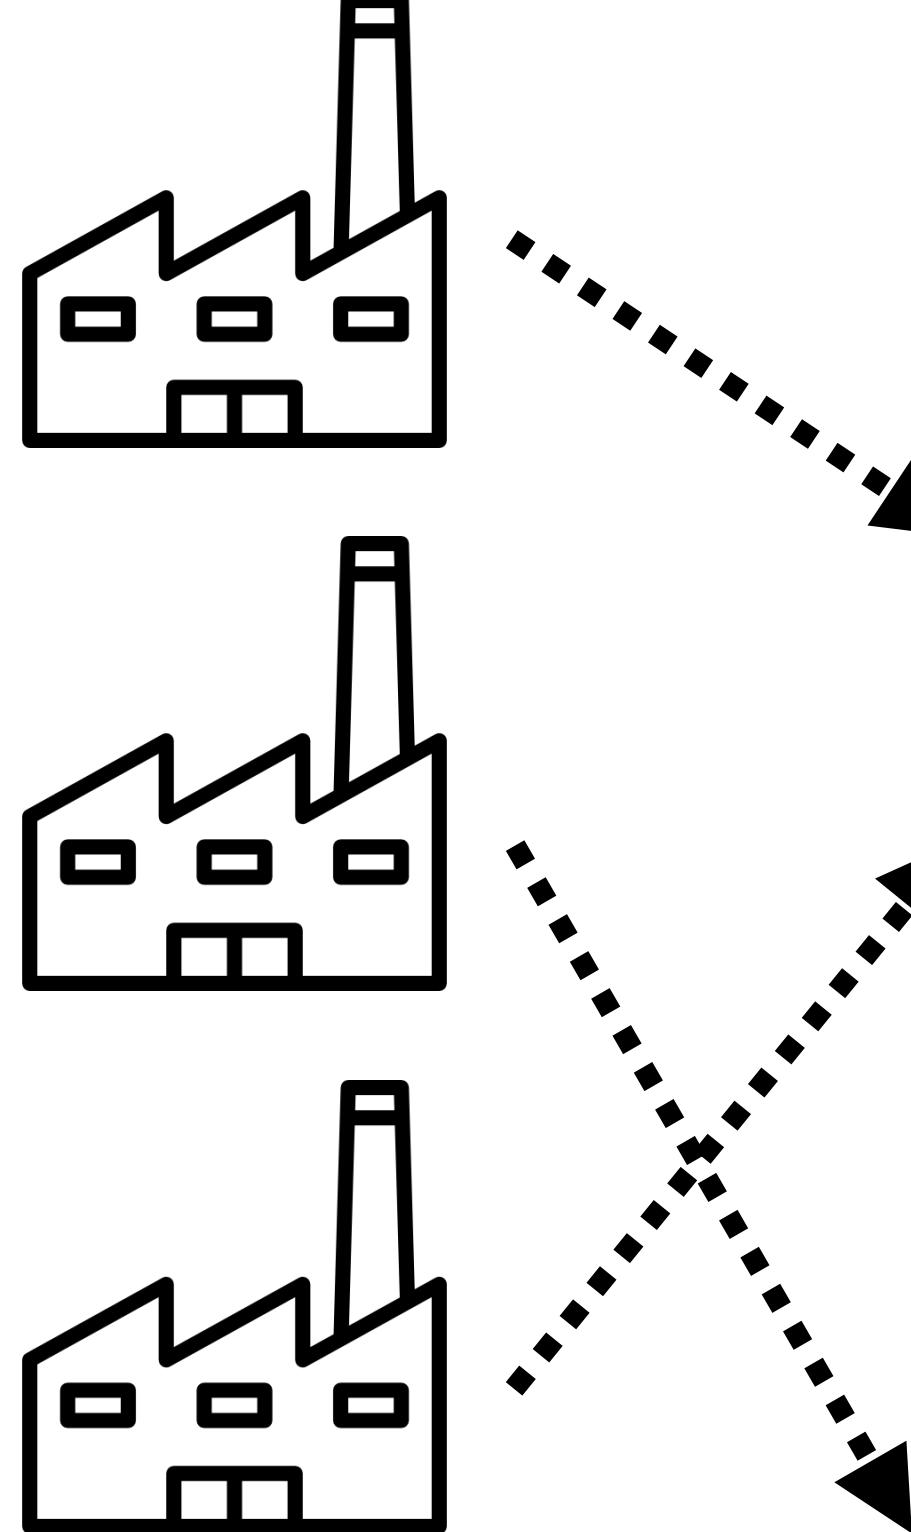
$O(n \log n)$

```
func join(L, R, keys) :  
    T := concat*(L, R)  
    T.sort(keys)  
    T.prefix_aggregate(...)  
    return T
```

$O(n \log n)$

Extends to **inner**, **outer**, **anti**, and **semi** joins.

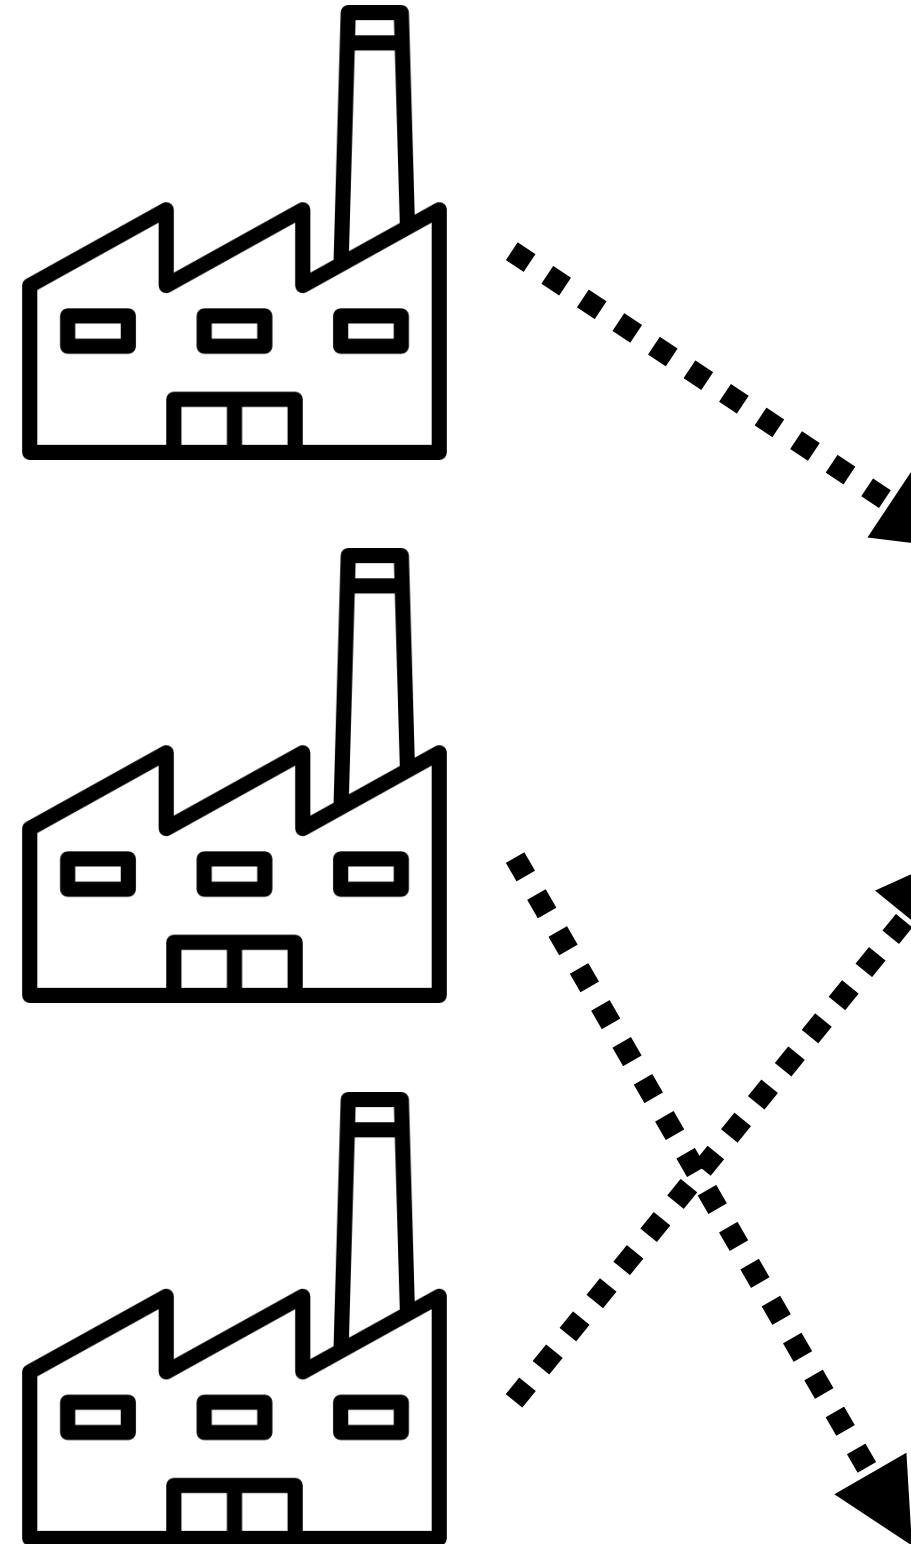
# Example with Our Approach



Customers	
Customer	AcctBalance
🤡	\$1000
😈	\$2000
🤡	\$500
🐸	\$60
😭	\$8000
😭	\$700
🐸	-\$200

Orders	
Customer	OrdCost
🐸	\$30
🤡	\$20
😭	\$55
😭	\$18
🤡	\$70
🤡	\$12
😈	\$8

# Example with Our Approach



Customers

Customer	AcctBalance
🤡	\$1000
😈	\$2000
🤡	\$500
🐸	\$60
😭	\$8000
😭	\$700
🐸	-\$200

Orders

Customer	OrdCost
🐸	\$30
🤡	\$20
😭	\$55
😭	\$18
🤡	\$70
🤡	\$12
😈	\$8

duplicates!

# Example with Our Approach

```
SELECT  
    SPECIES, SUM(ACCTBALANCE) / SUM(ORDCOST)  
FROM  
    ORDERS, CUSTOMERS  
WHERE  
    ORDER.CUSTOMER = CUSTOMERS.CUSTOMER  
GROUP BY  
    SPECIES
```

*Per customer group, what is the ratio of account balance to order total?*

Customers		Orders	
Customer	AcctBalance	Customer	OrdCost
🤡	\$1000	🐸	\$30
😈	\$2000	🤡	\$20
🤡	\$500	🥶	\$55
🐸	\$60	🥶	\$18
🥶	\$8000	🤡	\$70
🥶	\$700	🤡	\$12
🐸	-\$200	😈	\$8

# Pre-Aggregate

Customers		Orders	
Customer	AcctBalance	Customer	OrdCost
🤡	\$1000	🐸	\$30
😈	\$2000	🤡	\$20
🤡	\$500	🥶	\$55
🐸	\$60	🥶	\$18
🥶	\$8000	🤡	\$70
🥶	\$700	🤡	\$12
🐸	-\$200	😈	\$8

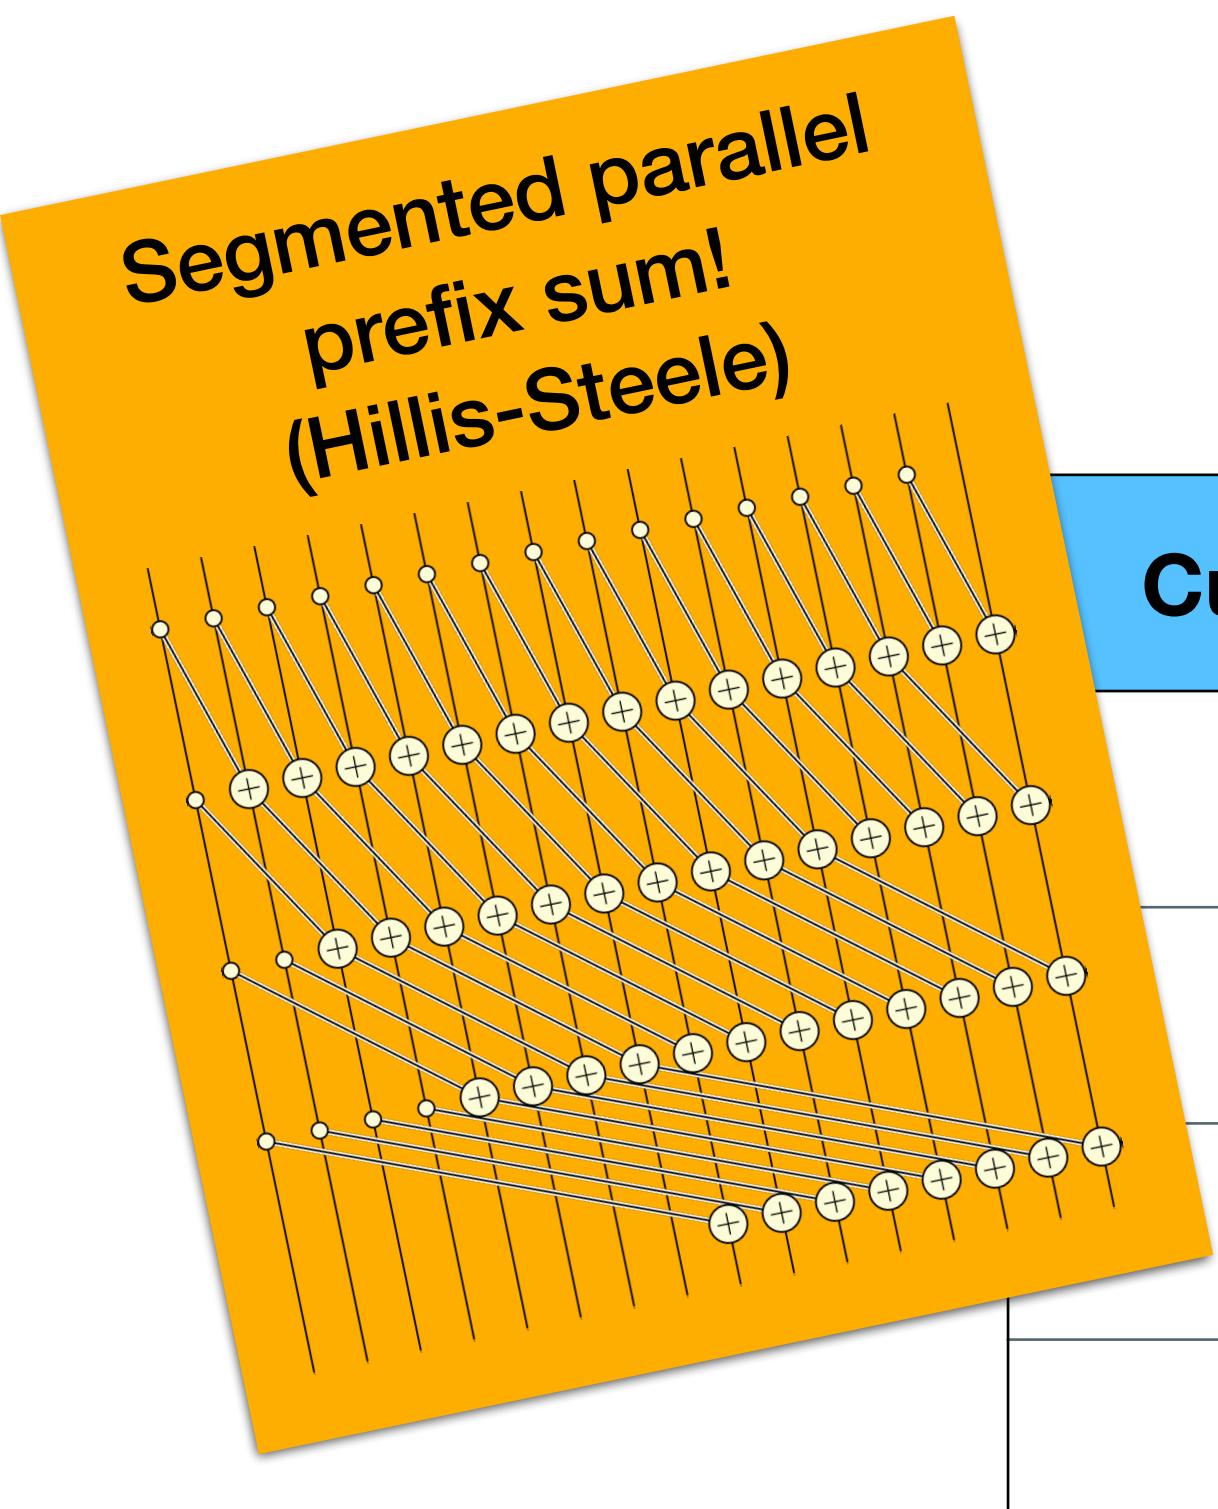
# Pre-Aggregate

Customers	
Customer	AcctBalance
🤡	\$1000
😈	\$2000
🤡	\$500
🐸	\$60
🥶	\$8000
🥶	\$700
🐸	-\$200

Orders	
Customer	OrdCost
🐸	\$30
🤡	\$20
🥶	\$55
🥶	\$18
🤡	\$70
🤡	\$12
😈	\$8

# After Pre-aggregation

## One-to-many join

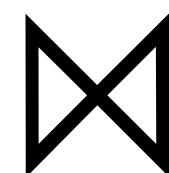


Customers

Customer	TotalAcctBalance
🤡	\$1500
😈	\$2000
🐸	-\$140
🥶	\$8700

*unique keys!*

*pre-aggregated values*



Orders

Customer	OrdCost
🐸	\$30
🤡	\$20
🥶	\$55
🥶	\$18
🤡	\$70
🤡	\$12
😈	\$8

# After Pre-aggregation

## One-to-many join

The diagram illustrates a one-to-many join operation between two tables: 'Customers' and 'Orders'. The 'Customers' table (left) contains four rows with data: a clown emoji and a balance of \$1500, a devil emoji and a balance of \$2000, a frog emoji and a balance of -\$140, and a blue face emoji and a balance of \$8700. A small padlock icon is positioned at the top right of this table. The 'Orders' table (right) contains seven rows, each with a customer emoji and an order cost: a frog emoji (\$30), a clown emoji (\$20), a blue face emoji (\$55), a blue face emoji (\$18), a clown emoji (\$70), a clown emoji (\$12), and a devil emoji (\$8). A small padlock icon is also at the top right of the 'Orders' table. A large 'X' symbol is centered between the two tables, indicating the join operation.

Customers		Orders	
Customer	TotalAcctBalance	Customer	OrdCost
🤡	\$1500	🐸	\$30
😈	\$2000	🤡	\$20
🐸	-\$140	🥶	\$55
🥶	\$8700	🥶	\$18
		🤡	\$70
		🤡	\$12
		😈	\$8

# Concatenate

```
func join(L, R, keys):  
    T := concat*(L, R)  
    T.sort(keys)  
    T.prefix_aggregate(...)  
    return T
```

Customer	C.Total	O.Cost
🤡	\$1500	
😈	\$2000	
🐸	-\$140	
🥶	\$8700	
🐸		\$30
🤡		\$20
🥶		\$55
🥶		\$18
🤡		\$70
🤡		\$12
😈		\$8

# Sort

```
func join(L, R, keys):  
    T := concat*(L, R)  
    T.sort(keys)  
    T.prefix_aggregate(...)  
    return T
```

Customer	C.Balance	O.Cost
🤡	\$1500	
🤡		\$20
🤡		\$70
🤡		\$12
😈	\$2000	
😈		\$8
🐸	-\$140	
🐸		\$30
🥶	\$8700	
🥶		\$55
🥶		\$18

# Group Aggregate



```
func join(L, R, keys):
    T := concat*(L, R)
    T.sort(keys)
    T.prefix_aggregate(...)
return T
```

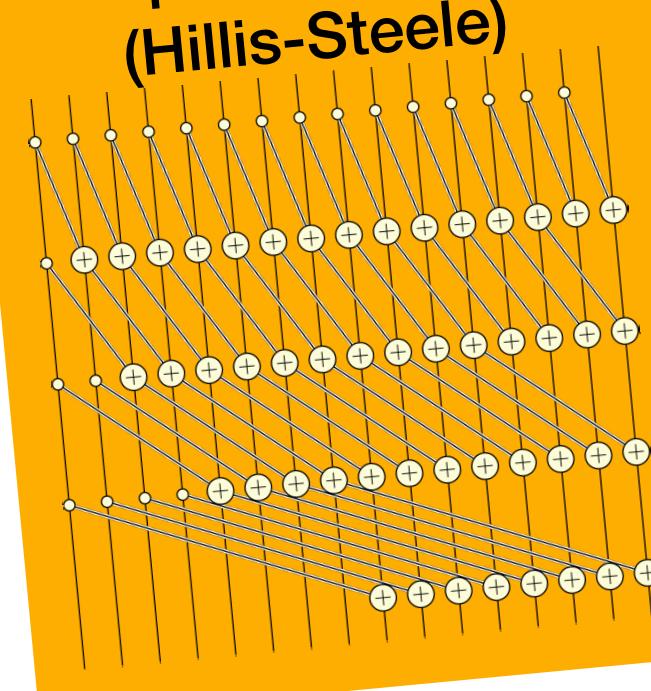
Customer	C.Balance	O.Cost	TotalCost
🤡	\$1500	∅	
🤡	\$1500	\$20	
🤡	\$1500	\$70	
🤡	\$1500	\$12	\$102
😈	\$2000	∅	
😈	\$2000	\$8	\$8
🐸	-\$140	∅	
🐸	-\$140	\$30	\$30
🥶	\$8700	∅	
🥶	\$8700	\$55	
🥶	\$8700	\$18	\$73

# Group Aggregate



```
func join(L, R, keys):
    T := concat*(L, R)
    T.sort(keys)
    T.prefix_aggregate(...)
    return T
```

Segmented parallel  
prefix sum!  
(Hillis-Steele)



Customer	C.Balance	O.Cost	TotalCost
🤡	\$1500	∅	
🤡	\$1500	\$20	
🤡	\$1500	\$70	
🤡	\$1500	\$12	\$102
😈	\$2000	∅	
😈	\$2000	\$8	\$8
🐸	-\$140	∅	
🐸	-\$140	\$30	\$30
🥶	\$8700	∅	
🥶	\$8700	\$55	
🥶	\$8700	\$18	\$73

# Mark Invalid Rows



```
func join(L, R, keys):
    T := concat*(L, R)
    T.sort(keys)
    T.prefix_aggregate(...)
    T.mark_invalid()
    return T
```

Customer	C.Balance	O.Cost	TotalCost
∅	∅	∅	∅
∅	∅	∅	∅
∅	∅	∅	∅
🤡	\$1500	\$12	\$102
∅	∅	∅	∅
😈	\$2000	\$8	\$8
∅	∅	∅	∅
🐸	-\$140	\$30	\$30
∅	∅	∅	∅
∅	∅	∅	∅
😭	\$8700	\$18	\$73

# Join Output

```
SELECT  
    SPECIES, SUM(ACCTBALANCE) / SUM(ORDCOST)  
FROM  
    ORDERS, CUSTOMERS  
WHERE  
    ORDER.CUSTOMER = CUSTOMERS.CUSTOMER  
GROUP BY  
    SPECIES
```

Customer	C.Balance	TotalCost	Species
🤡	\$1500	\$102	Human
🥶	\$8700	\$73	Human
😈	\$2000	\$8	NonHuman
🐸	-\$140	\$30	NonHuman



# Final Aggregation

```
SELECT  
    SPECIES, SUM(ACCTBALANCE) / SUM(ORDCOST)  
FROM  
    ORDERS, CUSTOMERS  
WHERE  
    ORDER.CUSTOMER = CUSTOMERS.CUSTOMER  
GROUP BY  
    SPECIES
```

Species	Customer	C.Balance	TotalCost
Human	∅	∅	∅
Human	∅	\$10,200	\$175
NonHuman	∅	∅	∅
NonHuman	∅	\$1,860	\$38



# Compute Ratio

```
SELECT  
    SPECIES, SUM(ACCTBALANCE) / SUM(ORDCOST)  
FROM  
    ORDERS, CUSTOMERS  
WHERE  
    ORDER.CUSTOMER = CUSTOMERS.CUSTOMER  
GROUP BY  
    SPECIES
```

Species	C.Balance	TotalCost	Ratio
Human	\$10,200	\$174	<b>58</b>
NonHuman	\$1,860	\$38	<b>48</b>



# Final Table

```
SELECT  
    SPECIES, SUM(ACCTBALANCE) / SUM(ORDCOST)  
FROM  
    ORDERS, CUSTOMERS  
WHERE  
    ORDER.CUSTOMER = CUSTOMERS.CUSTOMER  
GROUP BY  
    SPECIES
```

Species	Ratio
Human	58
NonHuman	48

# Final Table

```
SELECT  
    SPECIES, SUM(ACCTBALANCE) / SUM(ORDCOST)  
FROM  
    ORDERS, CUSTOMERS  
WHERE  
    ORDER.CUSTOMER = CUSTOMERS.CUSTOMER  
GROUP BY  
    SPECIES
```

Still secret shared! More joins...

Species	Ratio
Human	58
NonHuman	48

# Final Table

```
SELECT  
    SPECIES, SUM(ACCTBALANCE) / SUM(ORDCOST)  
FROM  
    ORDERS, CUSTOMERS  
WHERE  
    ORDER.CUSTOMER = CUSTOMERS.CUSTOMER  
GROUP BY  
    SPECIES
```

Species	Ratio
Human	58
NonHuman	48

Still secret shared! More joins...



Oblivious output bound:  
 $| \text{Orders} |$

# ORQ: an Oblivious Relational Query engine

# ORQ: an Oblivious Relational Query engine

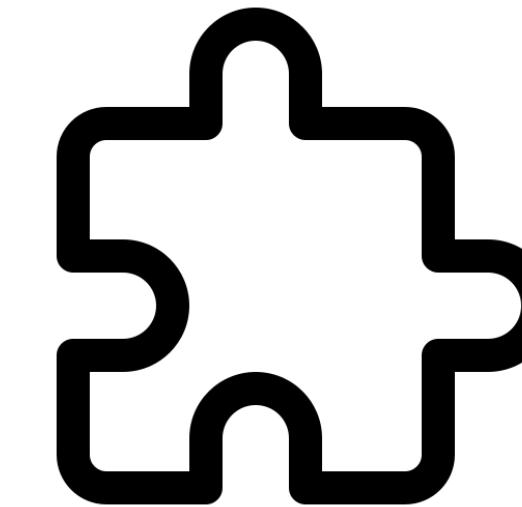


Familiar APIs &  
simple orchestration

# ORQ: an Oblivious Relational Query engine



Familiar APIs &  
simple orchestration

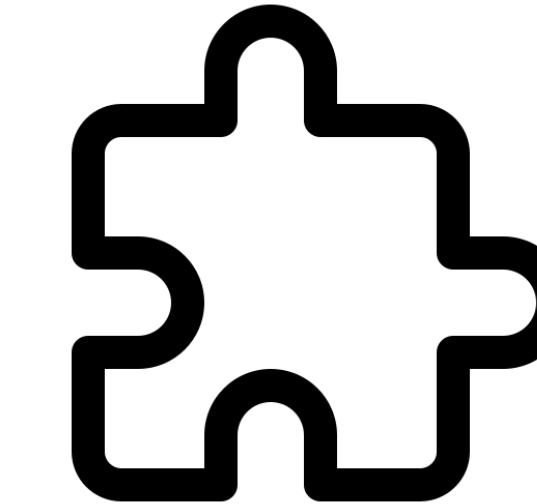


Support for multiple  
protocols & threat models

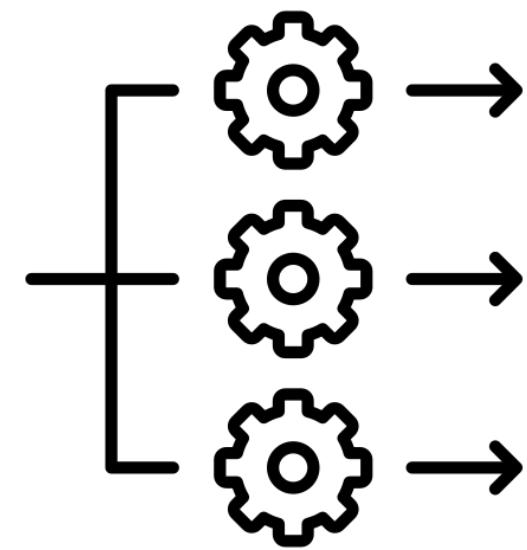
# ORQ: an Oblivious Relational Query engine



Familiar APIs &  
simple orchestration



Support for multiple  
protocols & threat models

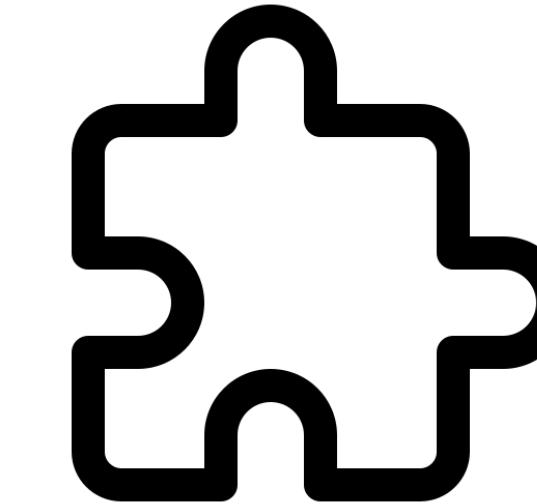


Data-parallel runtime

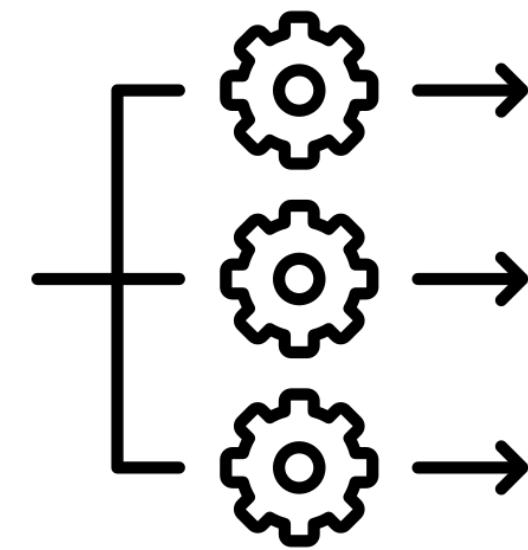
# ORQ: an Oblivious Relational Query engine



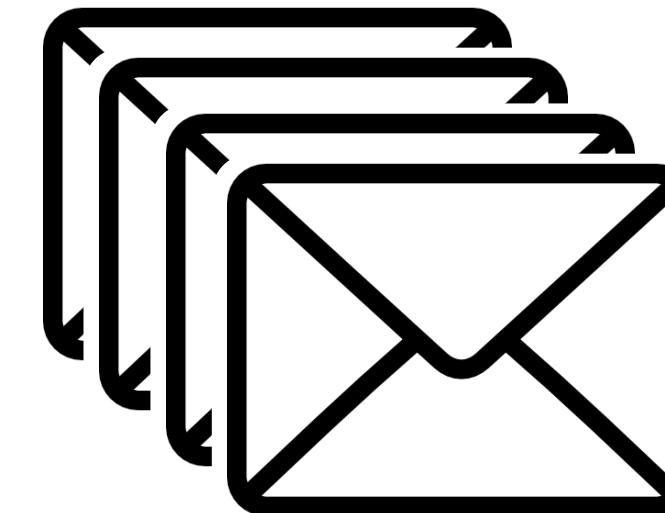
Familiar APIs &  
simple orchestration



Support for multiple  
protocols & threat models



Data-parallel runtime

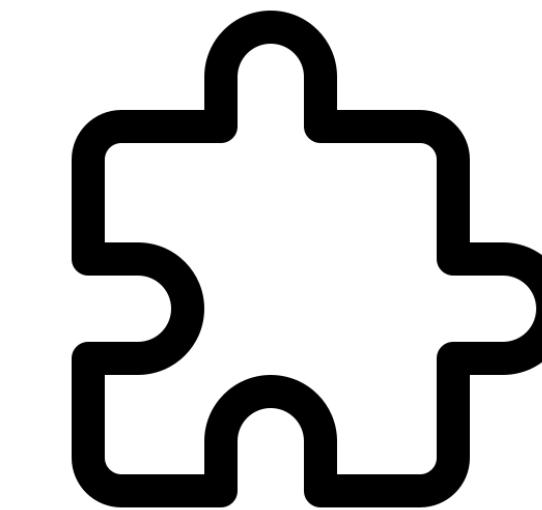


Vectorization &  
message batching

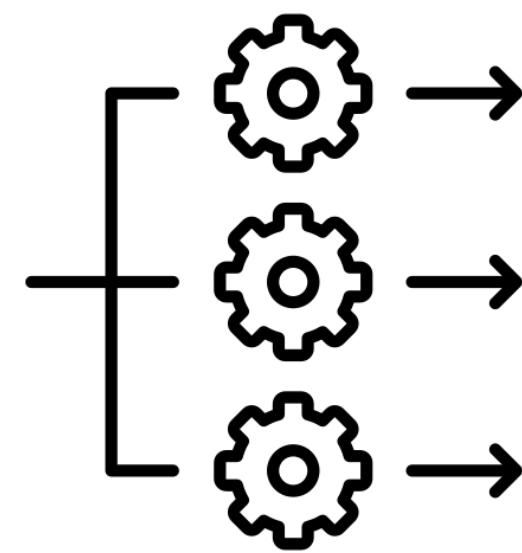
# ORQ: an Oblivious Relational Query engine



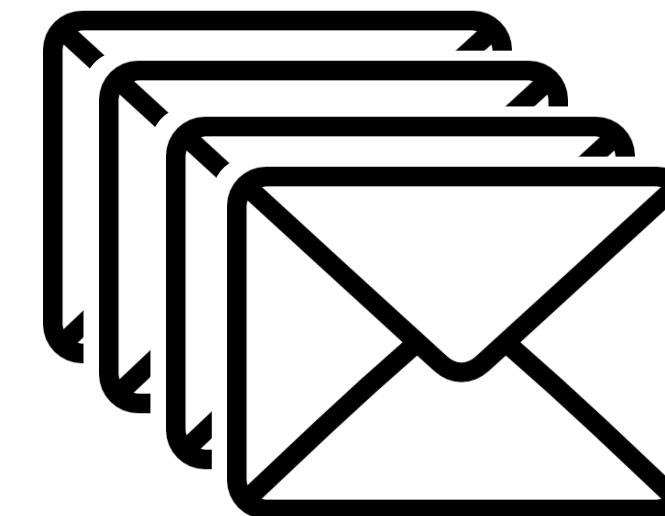
Familiar APIs & simple orchestration



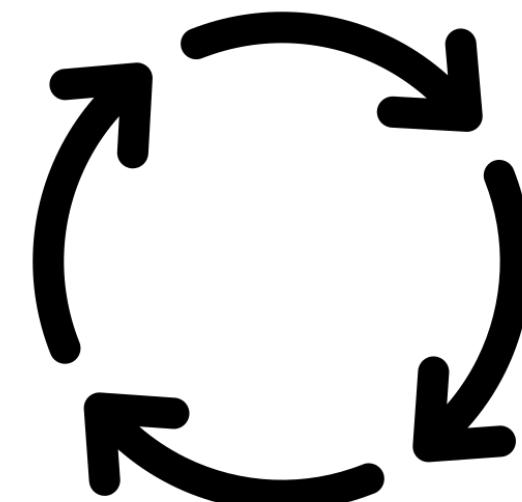
Support for multiple protocols & threat models



Data-parallel runtime



Vectorization & message batching

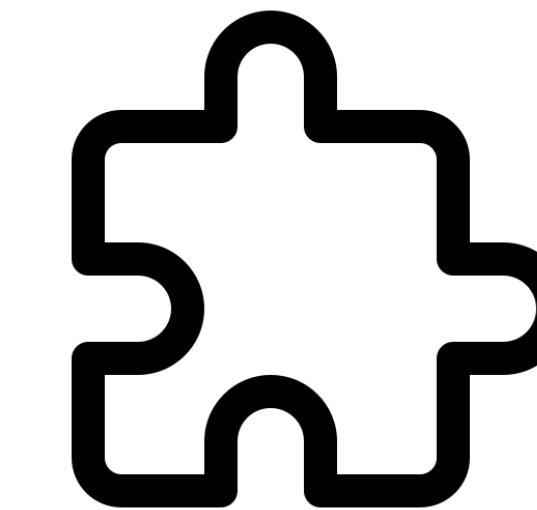


Custom TCP communicator

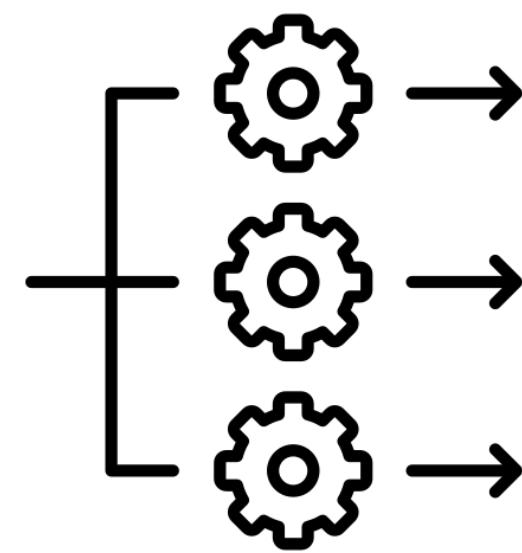
# ORQ: an Oblivious Relational Query engine



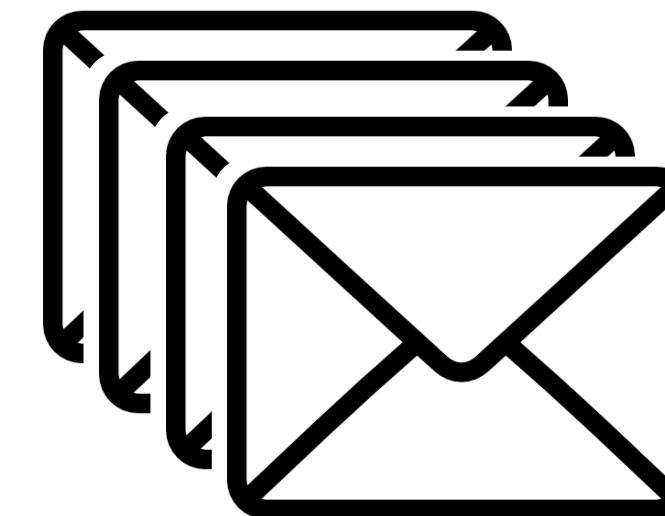
Familiar APIs & simple orchestration



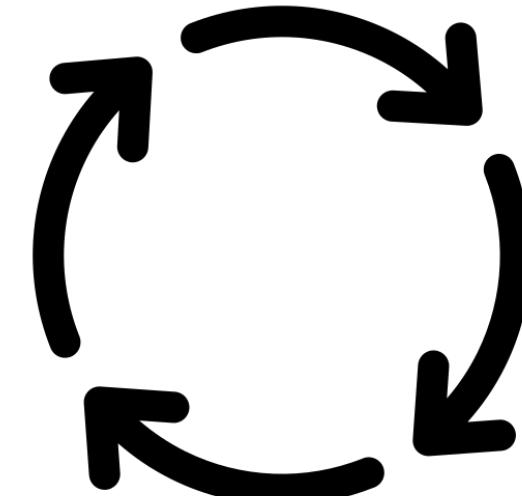
Support for multiple protocols & threat models



Data-parallel runtime



Vectorization & message batching



Custom TCP communicator

14,000 C++ LoC

# ORQ Stack

## Distributed Parallel Runtime

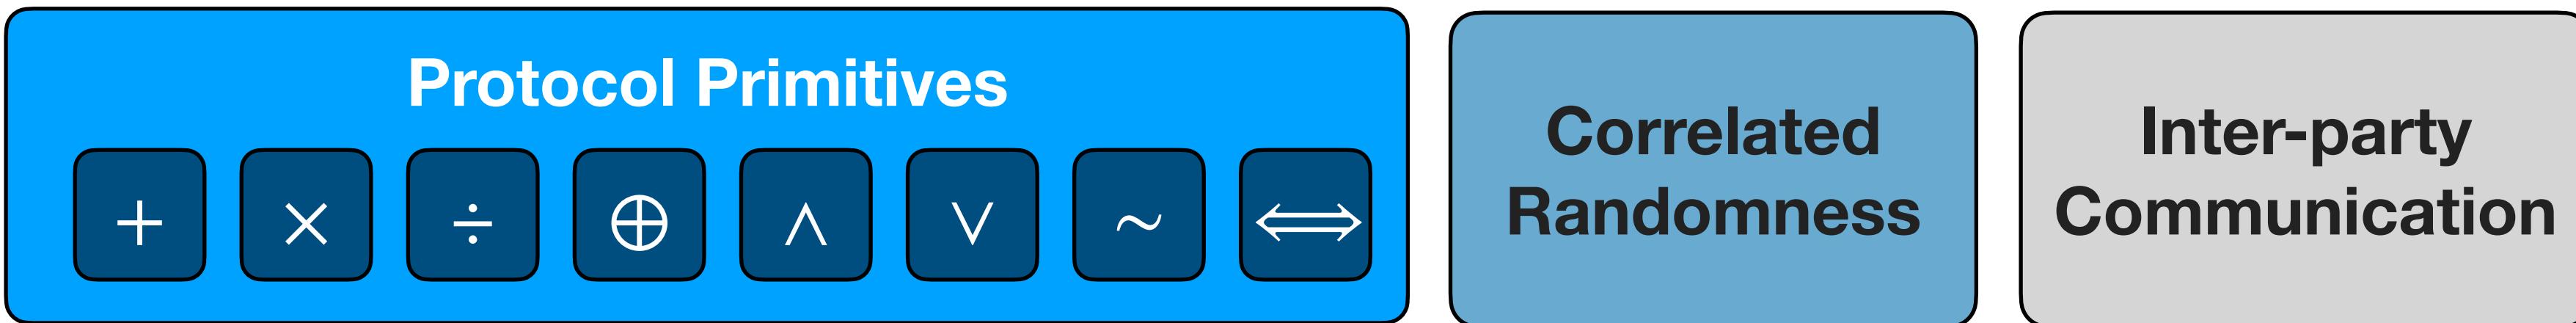


Untrusted Infrastructure  
CloudLab NERC



# ORQ Stack

MPC  
protocol  
layer



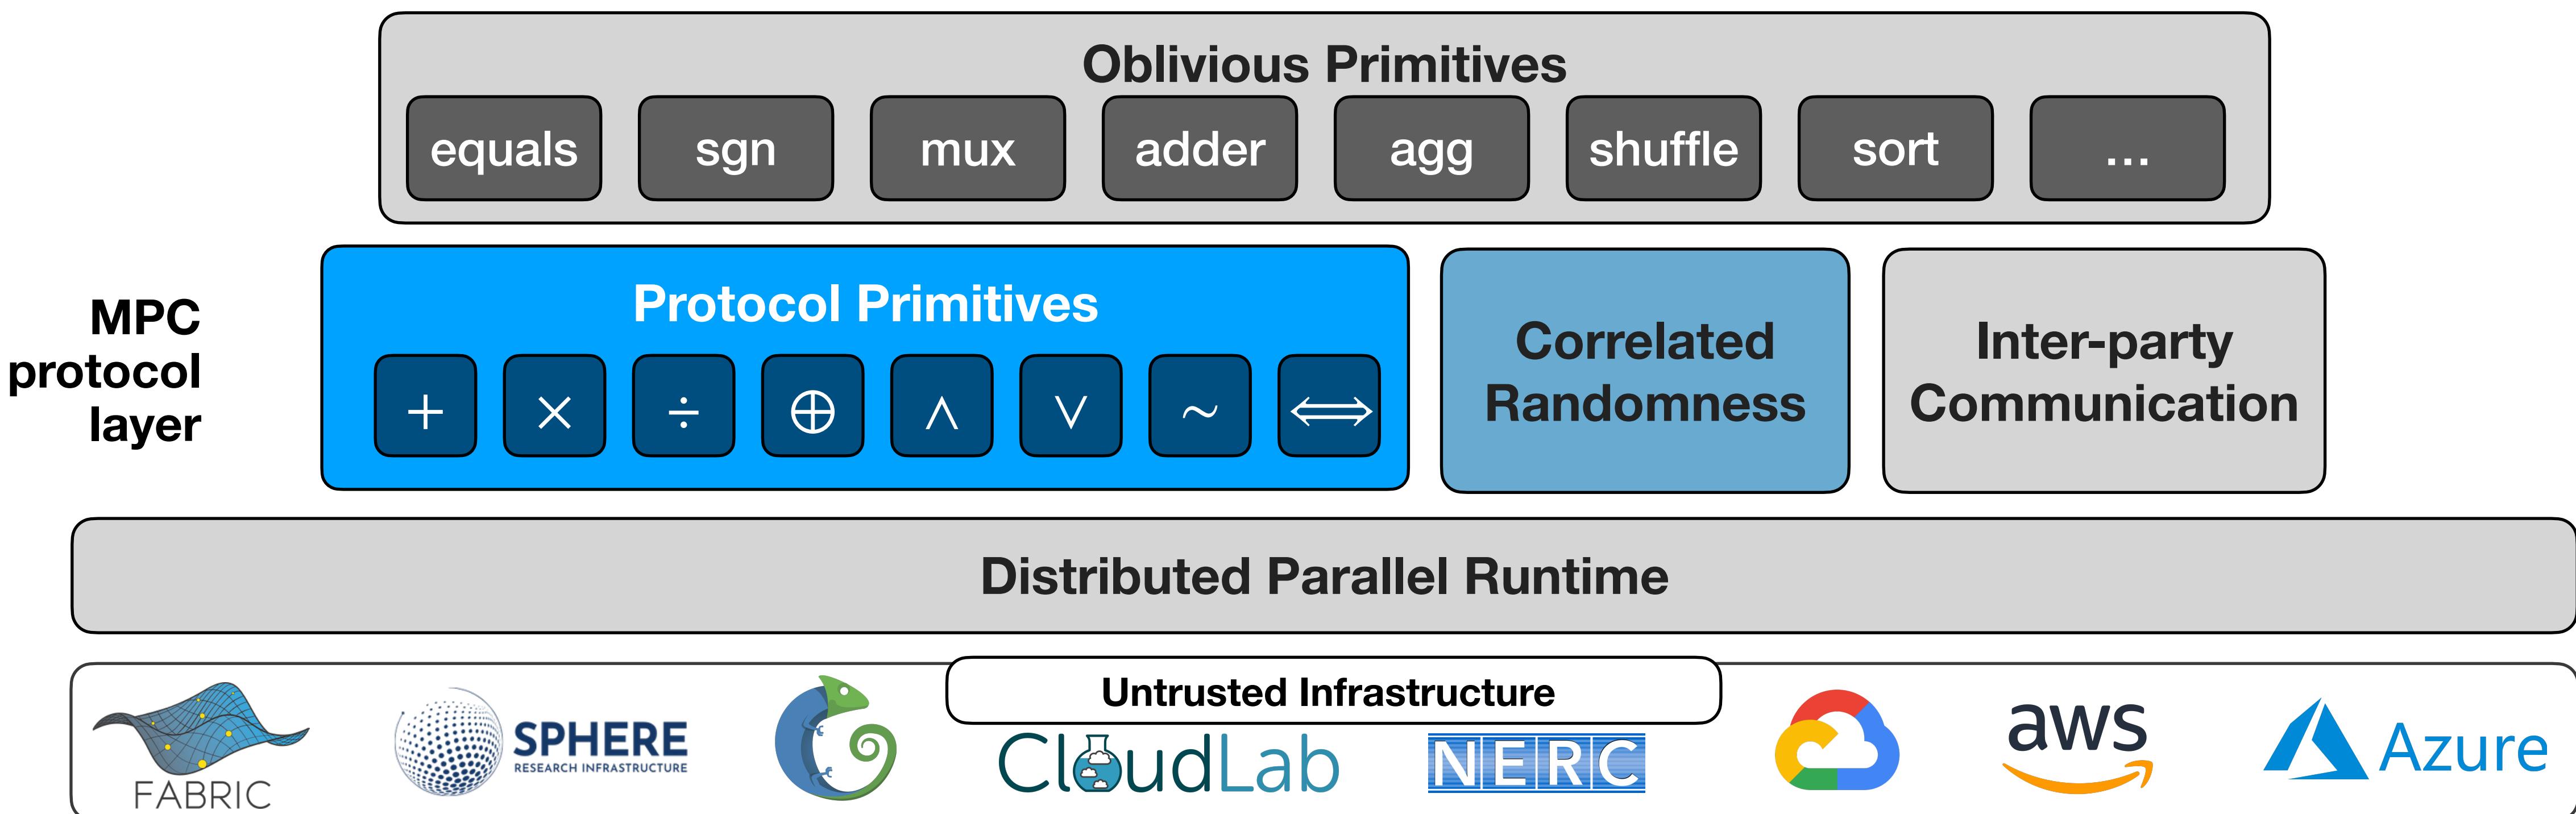
Distributed Parallel Runtime



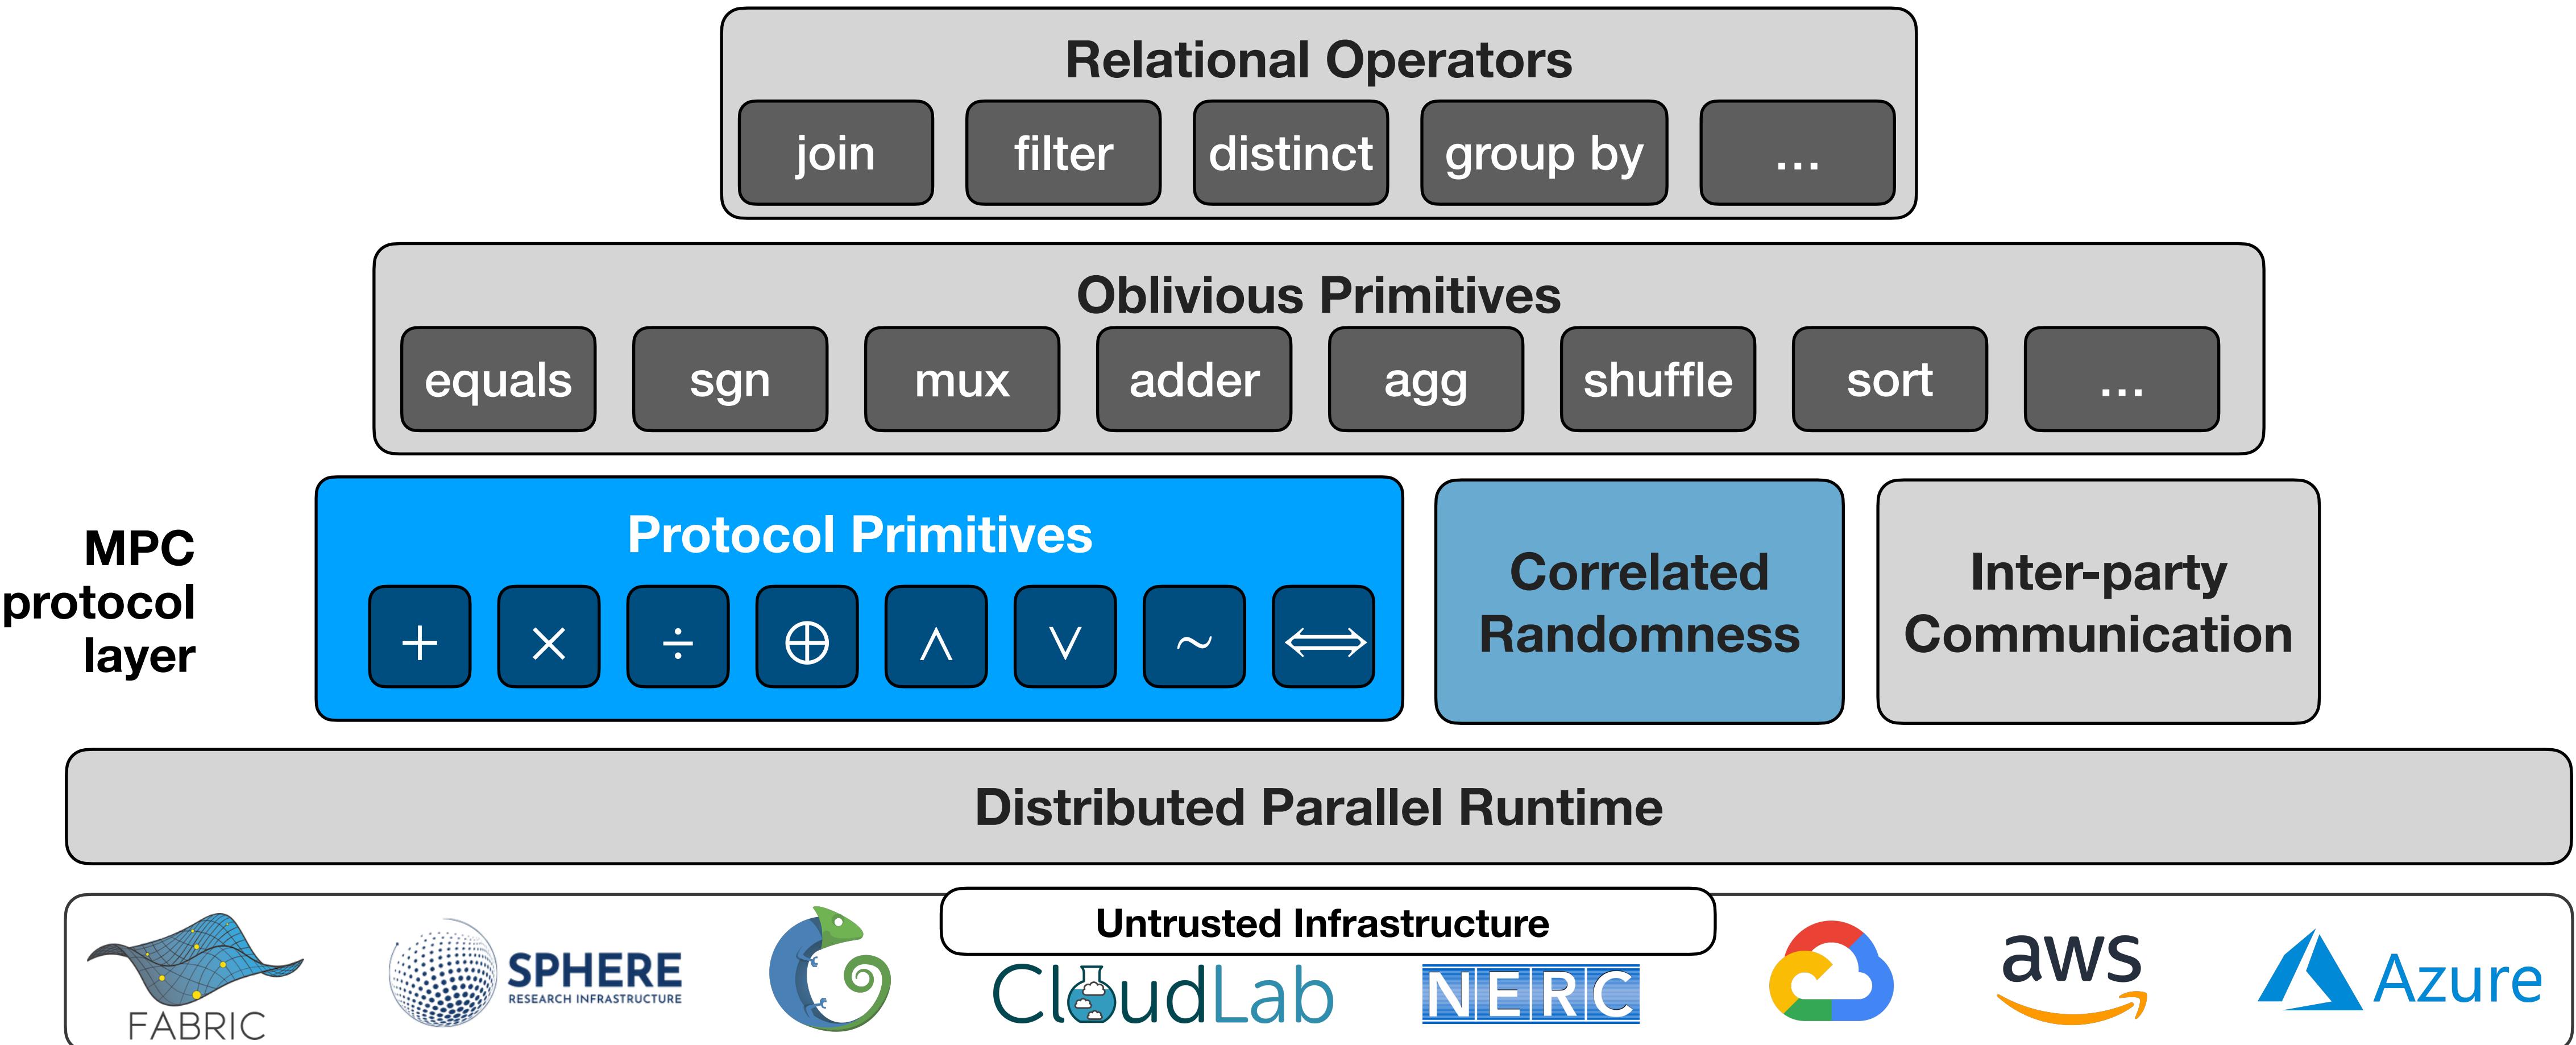
Untrusted Infrastructure  
CloudLab NERC



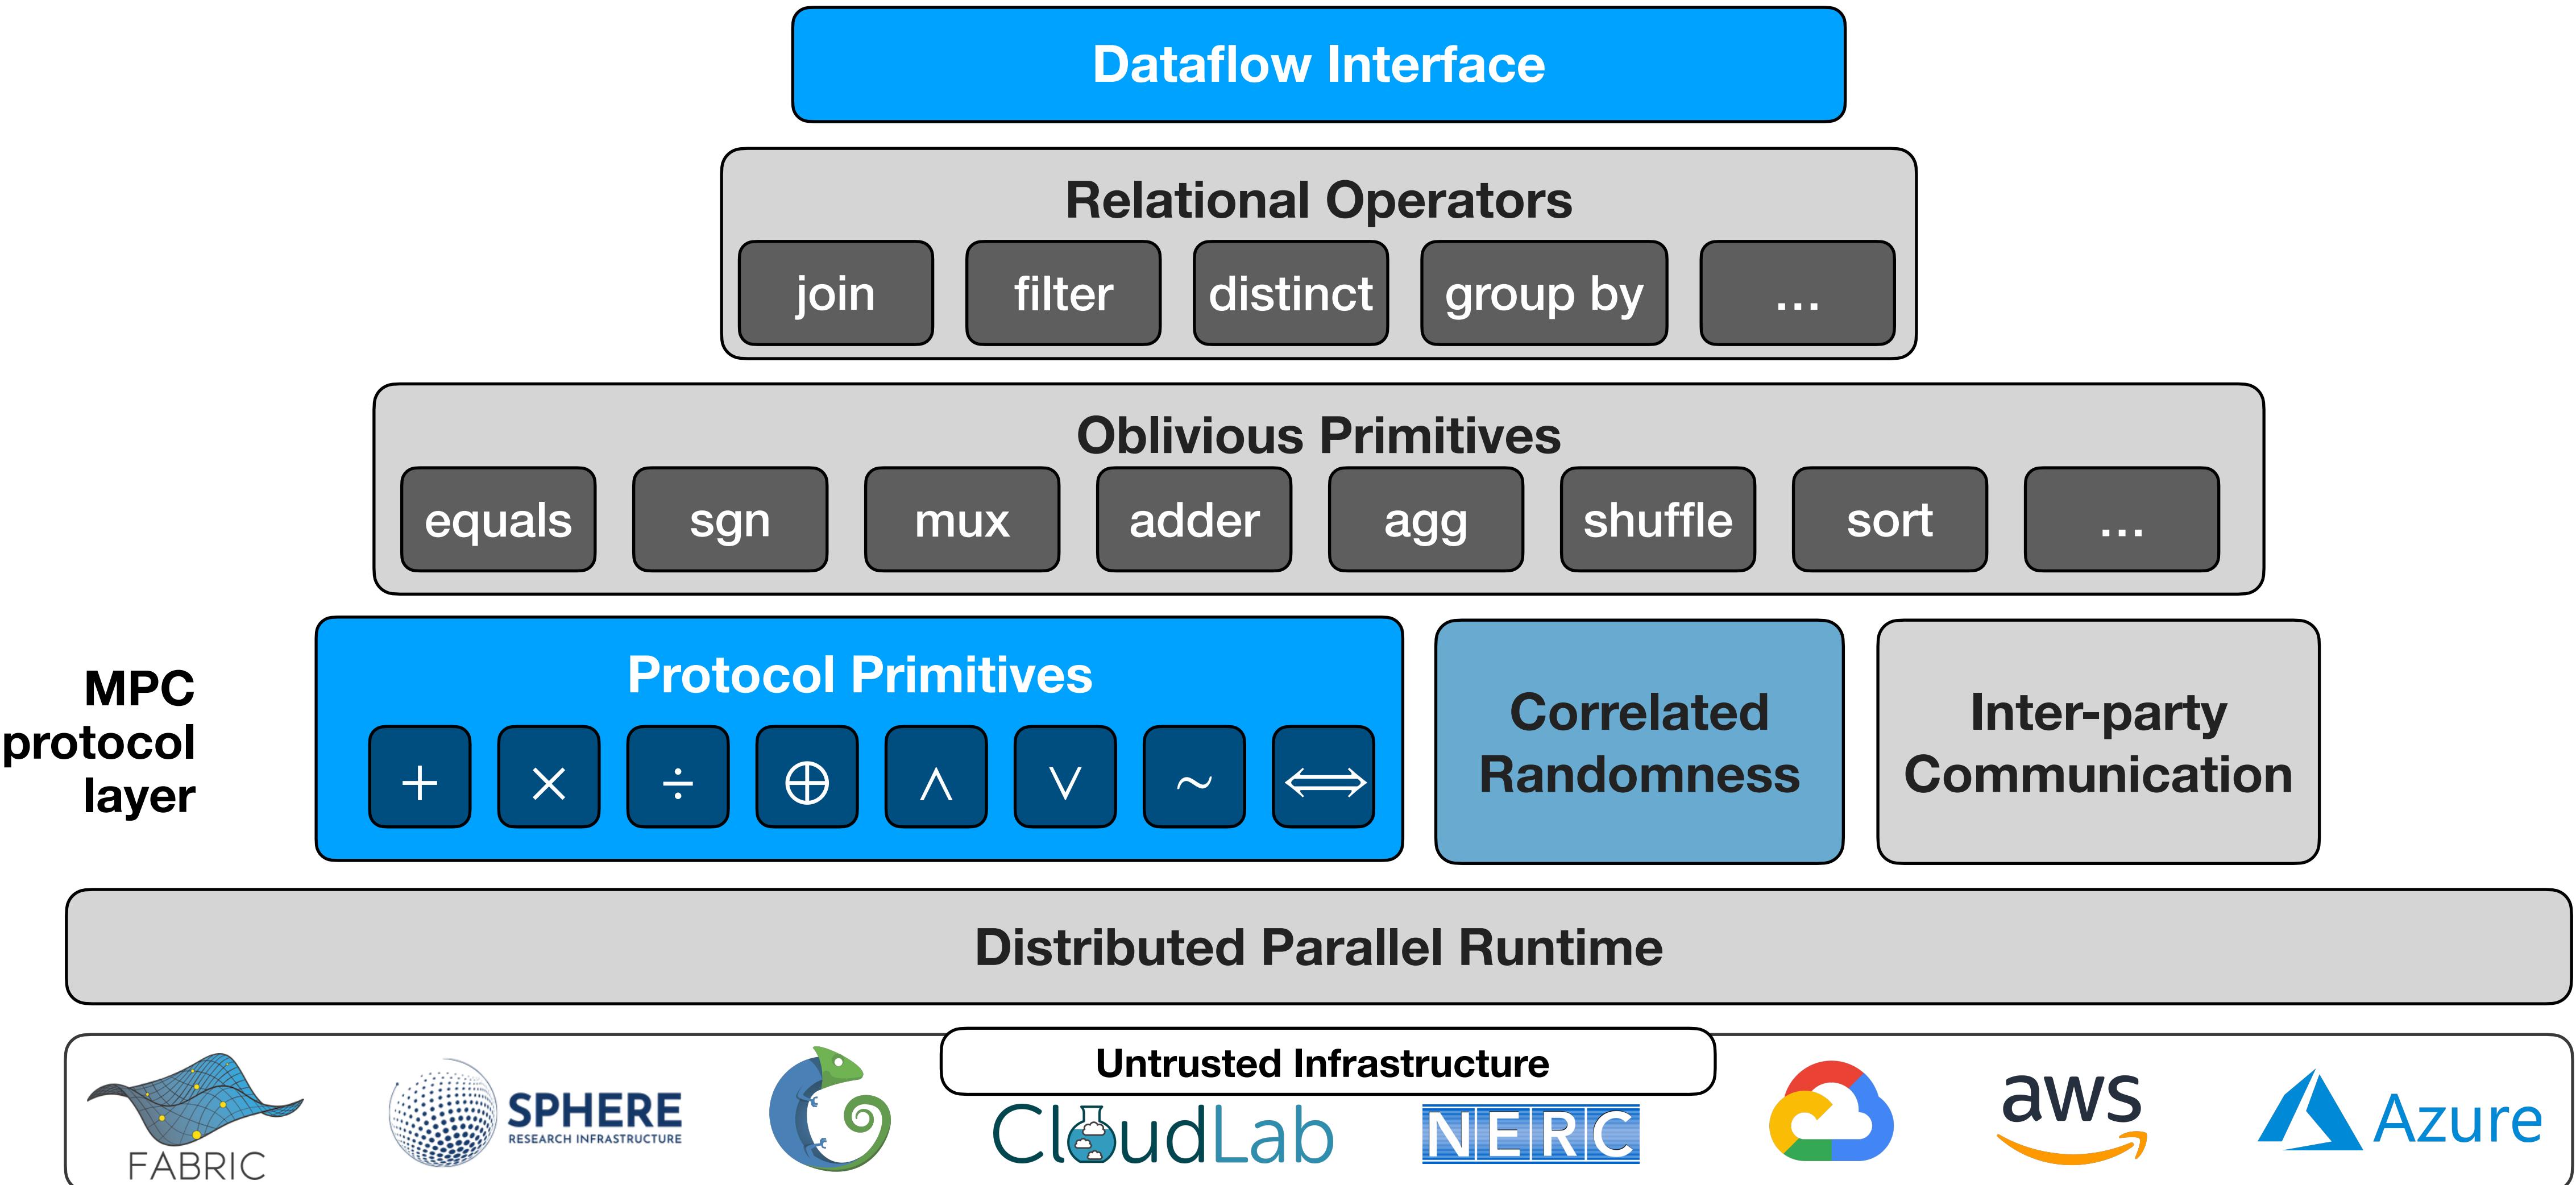
# ORQ Stack



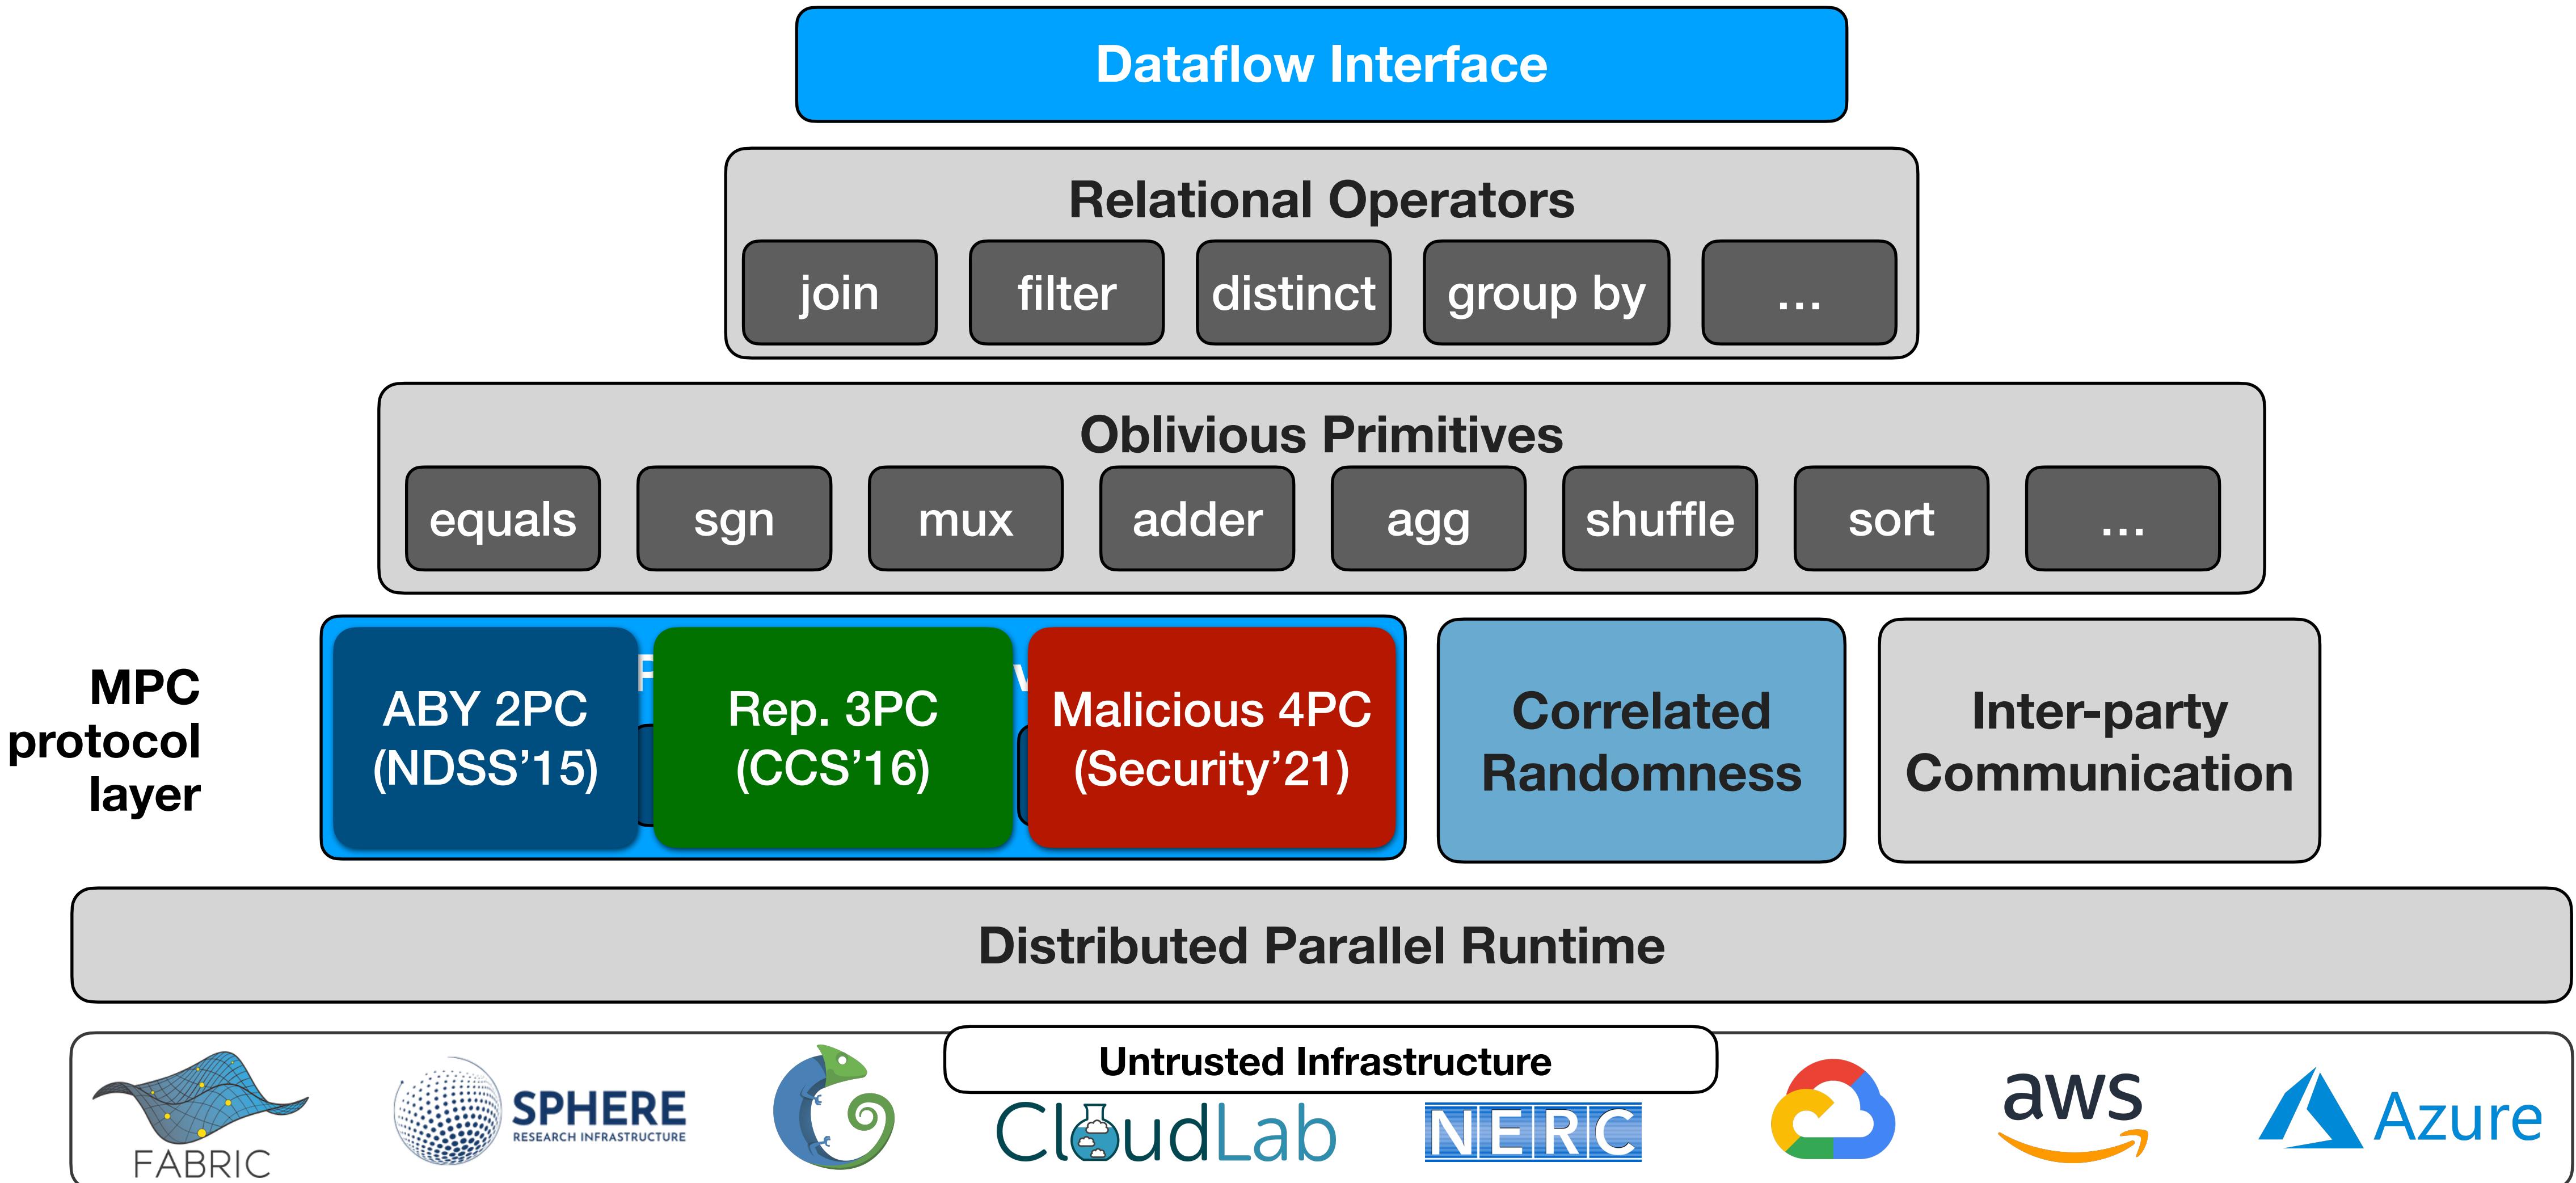
# ORQ Stack



# ORQ Stack



# ORQ Stack



# API Example

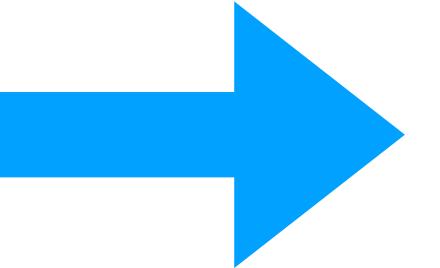
## TPC-H Q13: SQL vs. ORQ Dataflow API

```
select
    c_count, count(*) as custdist
from (
    select
        c_custkey,
        count(o_orderkey)
    from
        customer left outer join orders on
            c_custkey = o_custkey
            and o_comment != ''
    group by
        c_custkey
    ) as c_orders (c_custkey, c_count)
group by
    c_count
order by
    custdist desc,
    c_count desc;
```

# API Example

## TPC-H Q13: SQL vs. ORQ Dataflow API

```
select
    c_count, count(*) as custdist
from (
    select
        c_custkey,
        count(o_orderkey)
    from
        customer left outer join orders on
            c_custkey = o_custkey
            and o_comment != ''
    group by
        c_custkey
    ) as c_orders (c_custkey, c_count)
group by
    c_count
order by
    custdist desc,
    c_count desc;
```



```
auto C = db.getCustomersTable();
auto O = db.getOrdersTable();

O.filter(O["[Comment]"] != '');

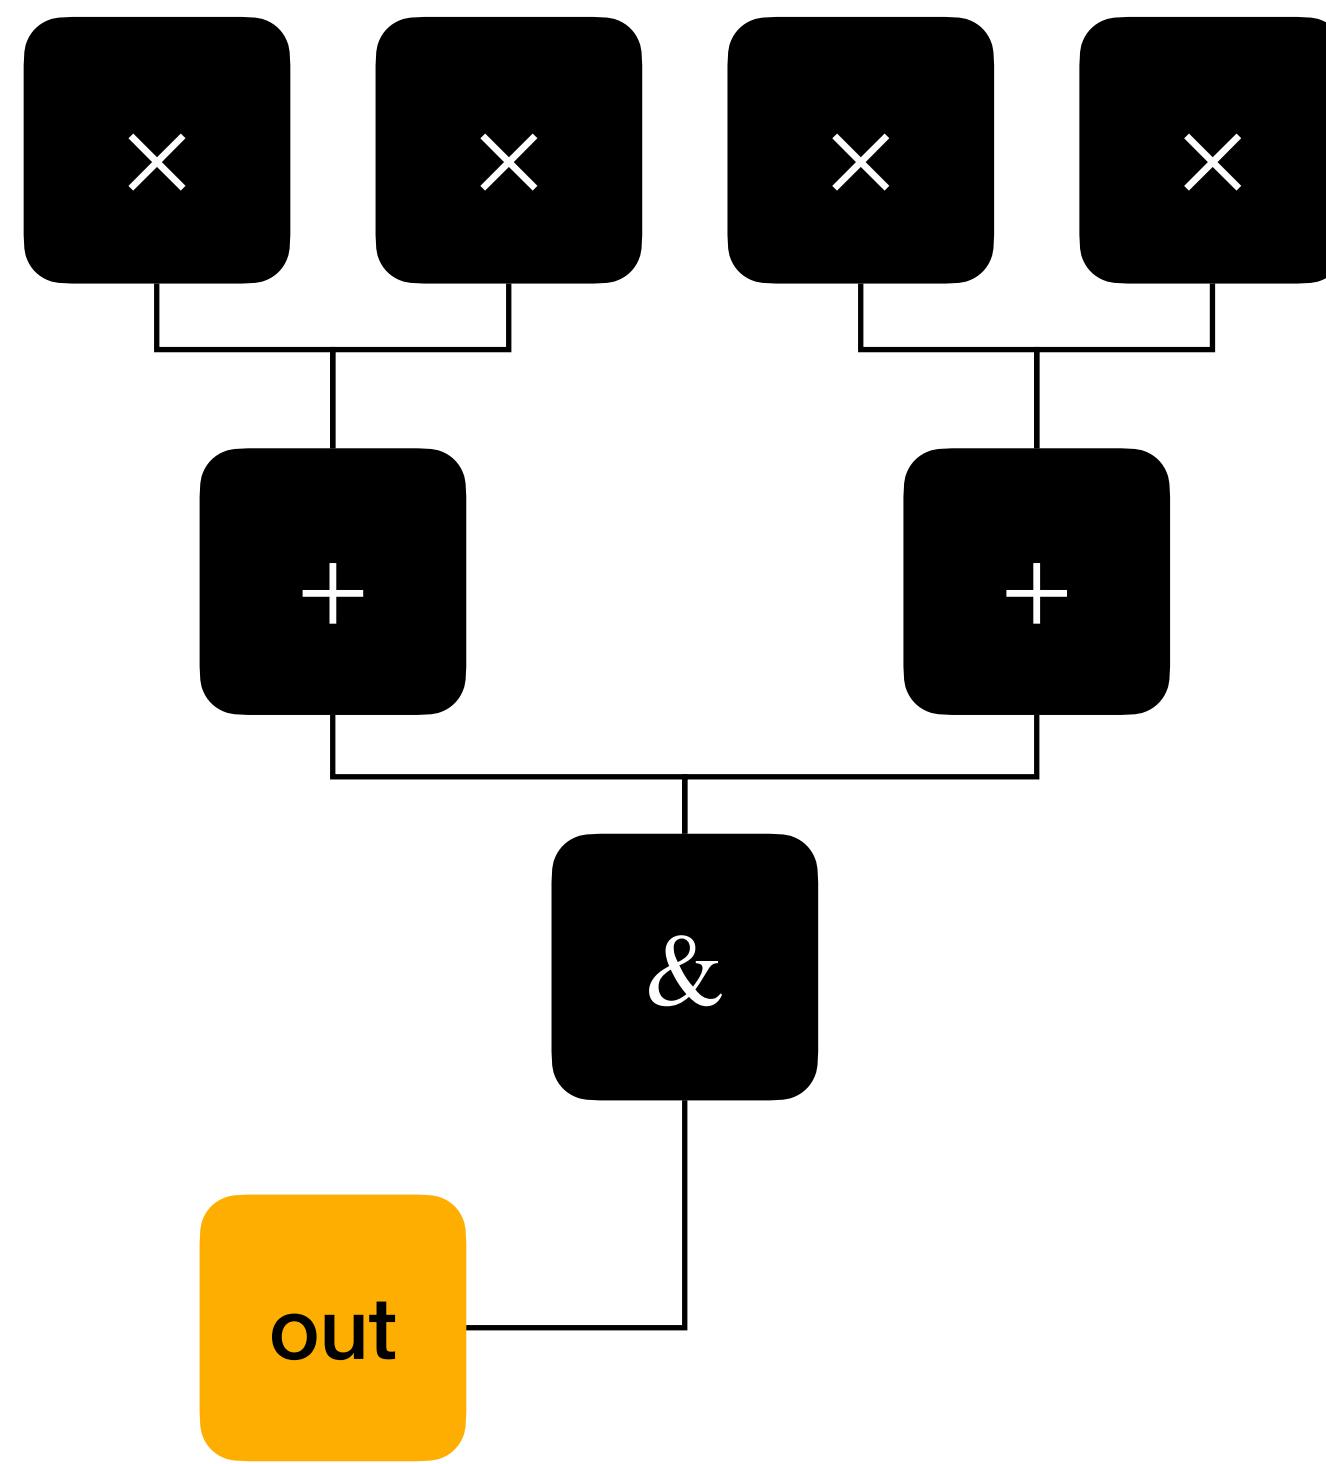
auto T = C.left_outer_join(O,
    {"[CustKey]"},
    {"Count", "Count", count});

T.convert_a2b("Count", "[Count]");

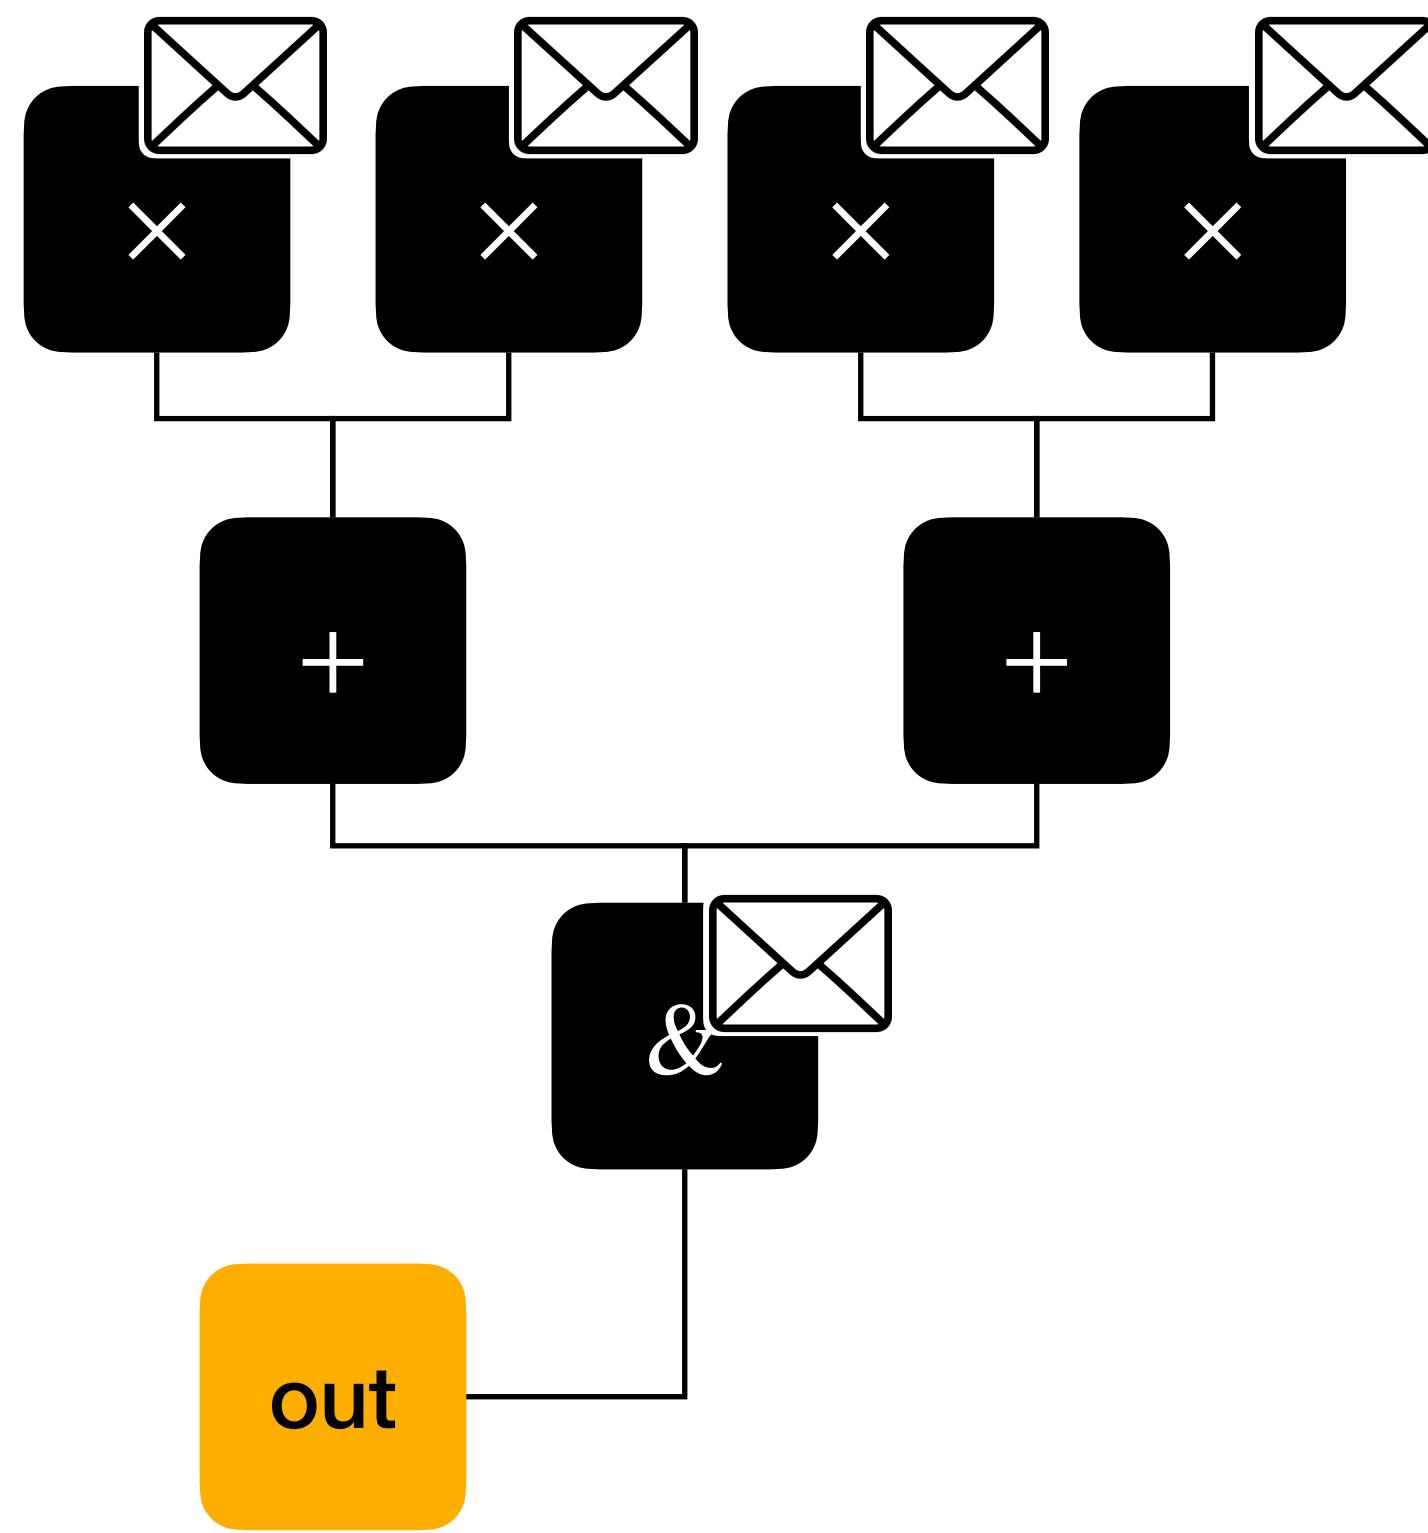
auto F = T.aggregate(
    {"[Count]"},
    {"CustDist", "CustDist", count});

F.convert_a2b("CustDist", "[CustDist]");
F.sort({"[CustDist]", "[Count]"}, DESC);
```

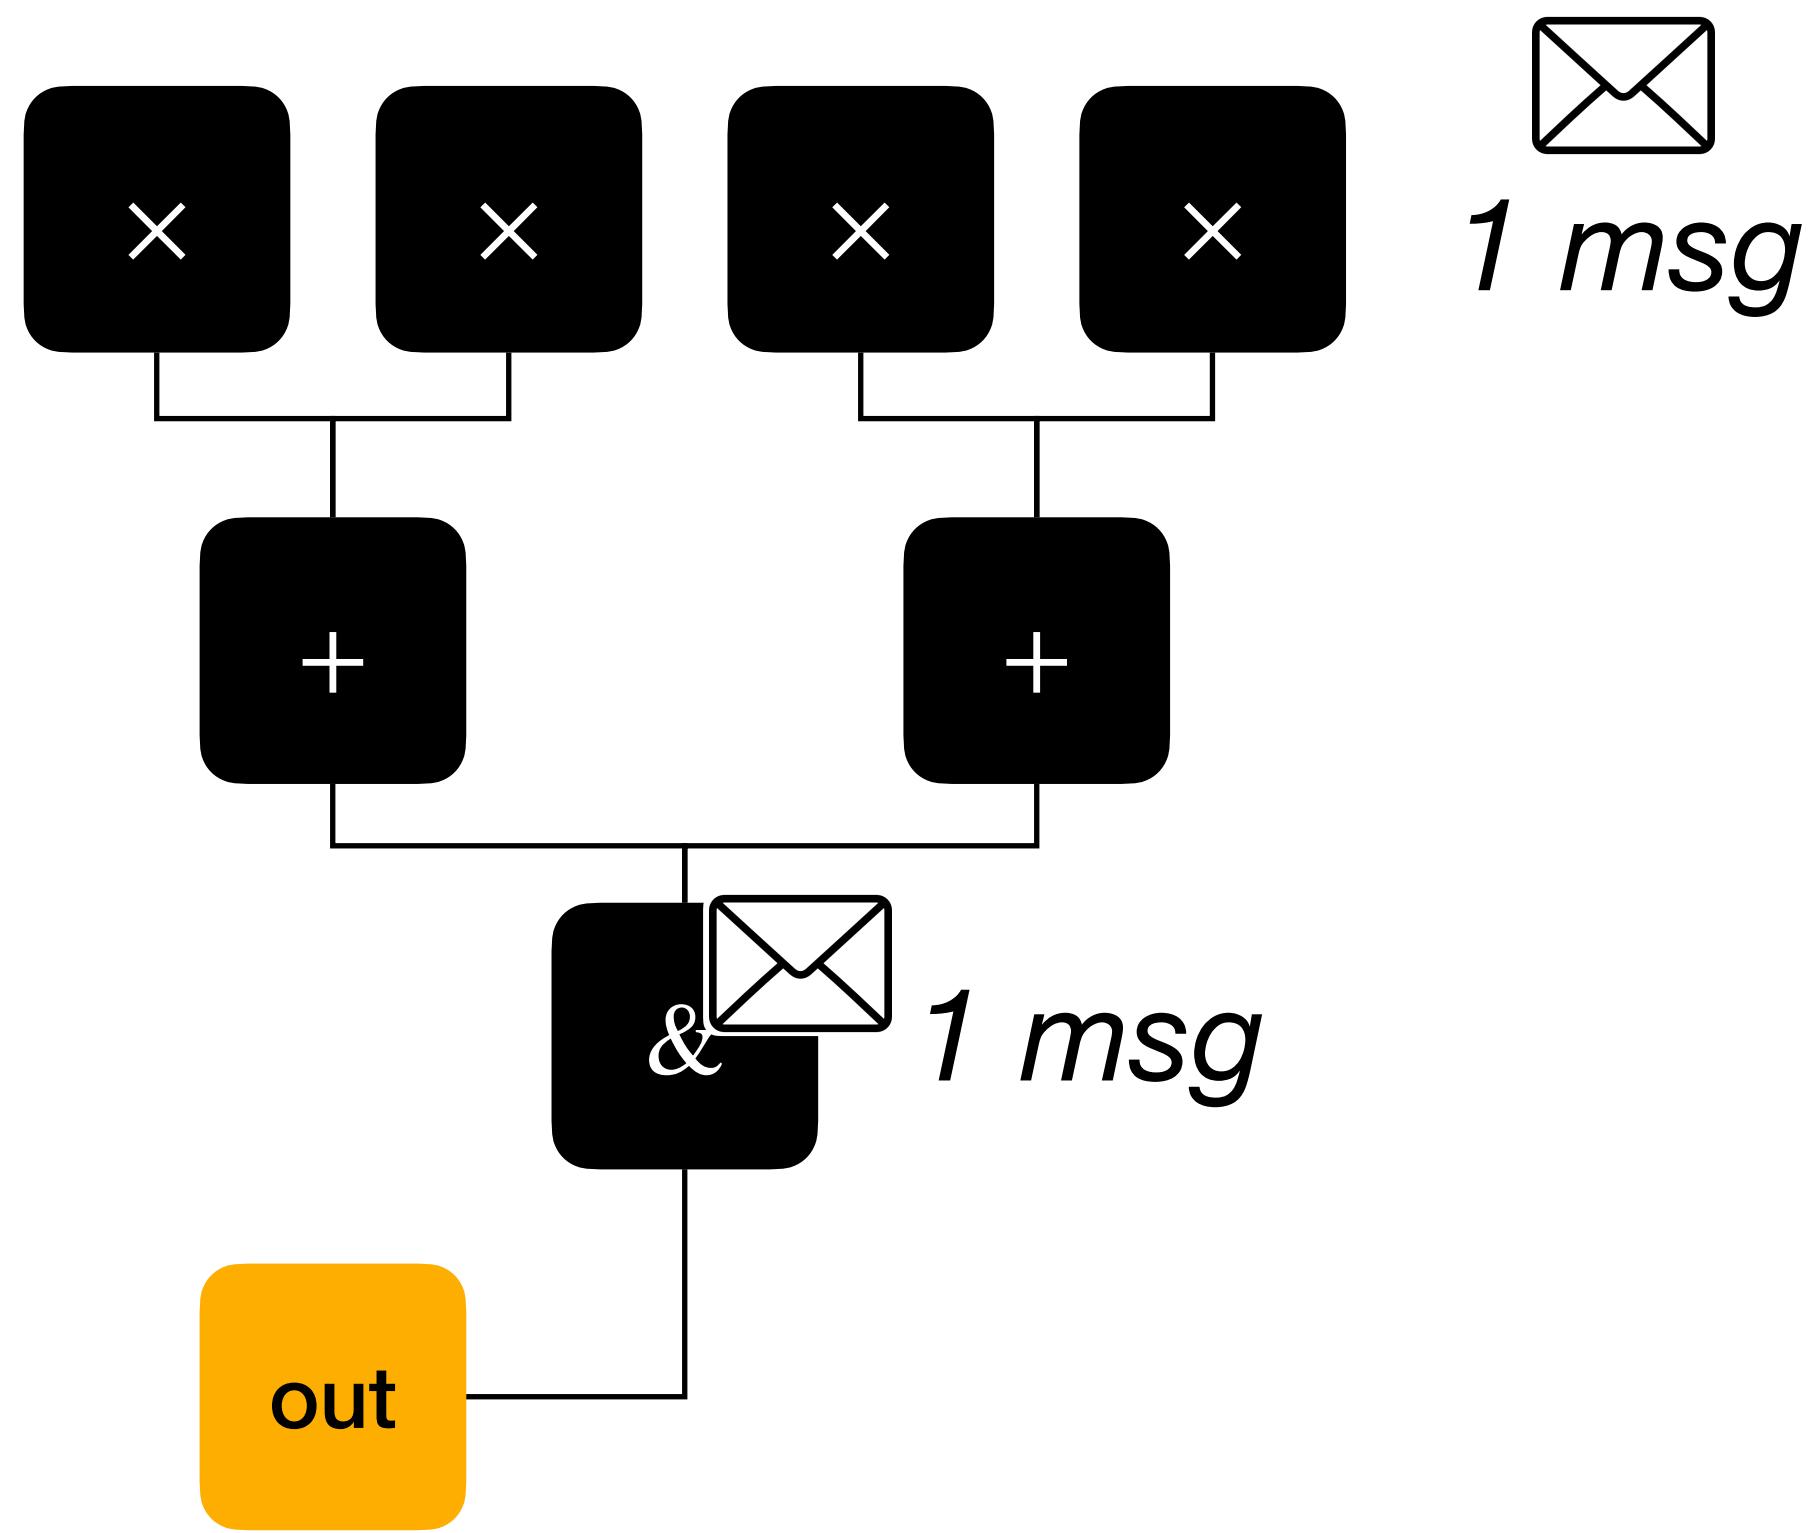
# Vectorization & Message Batching



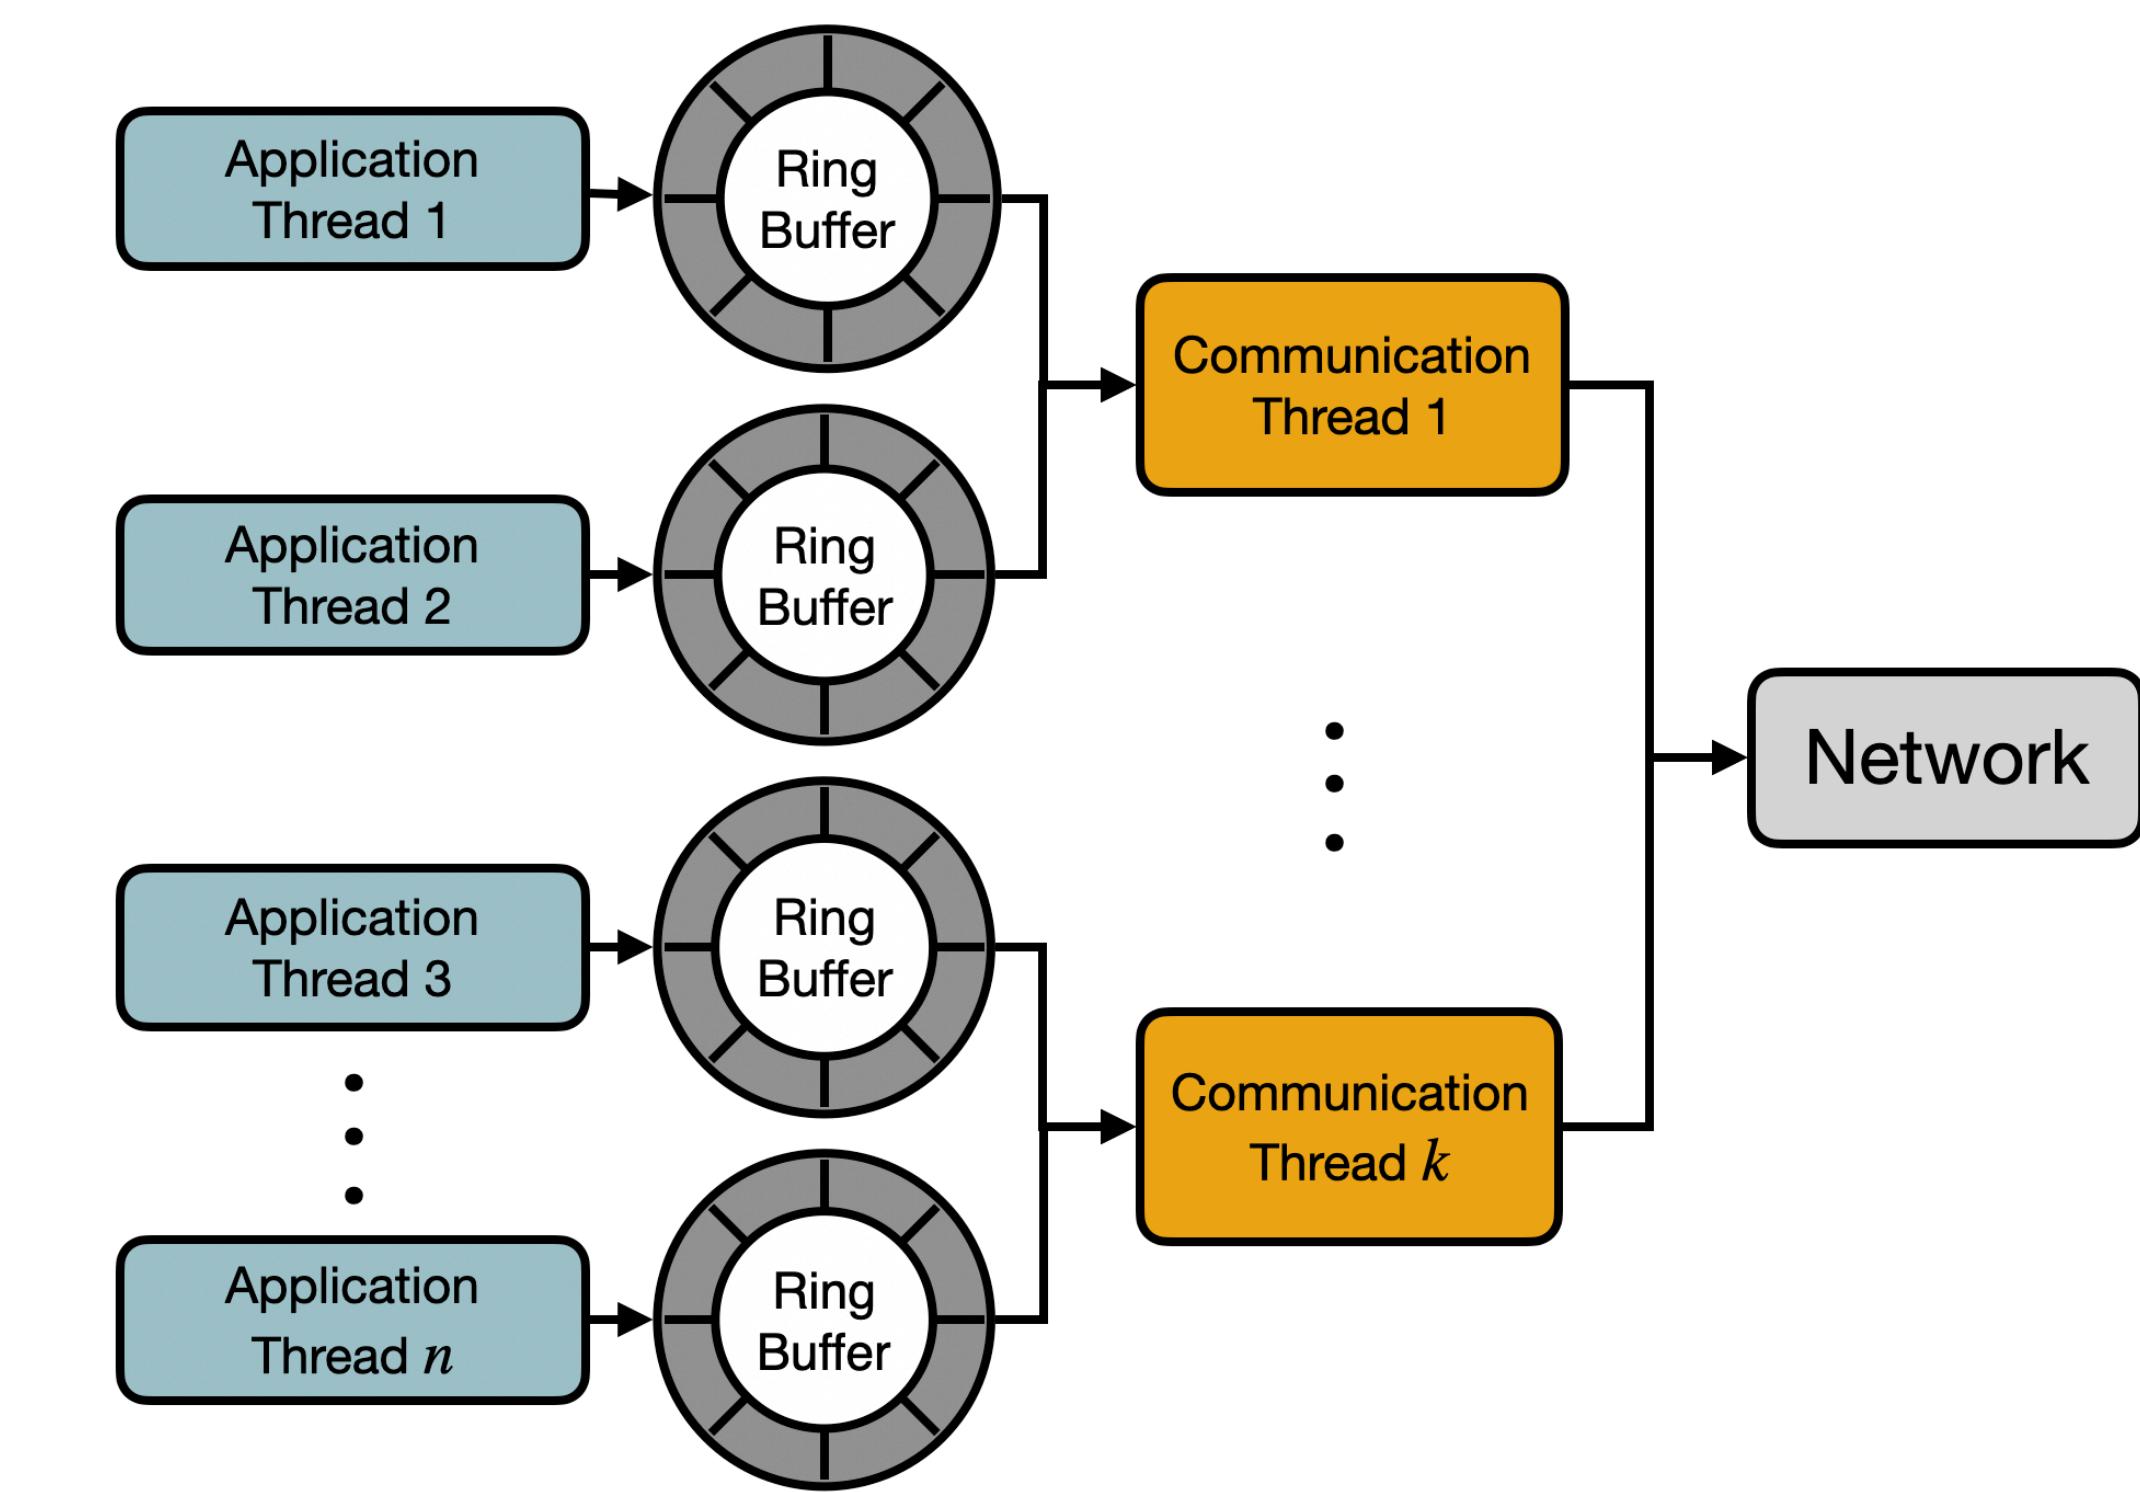
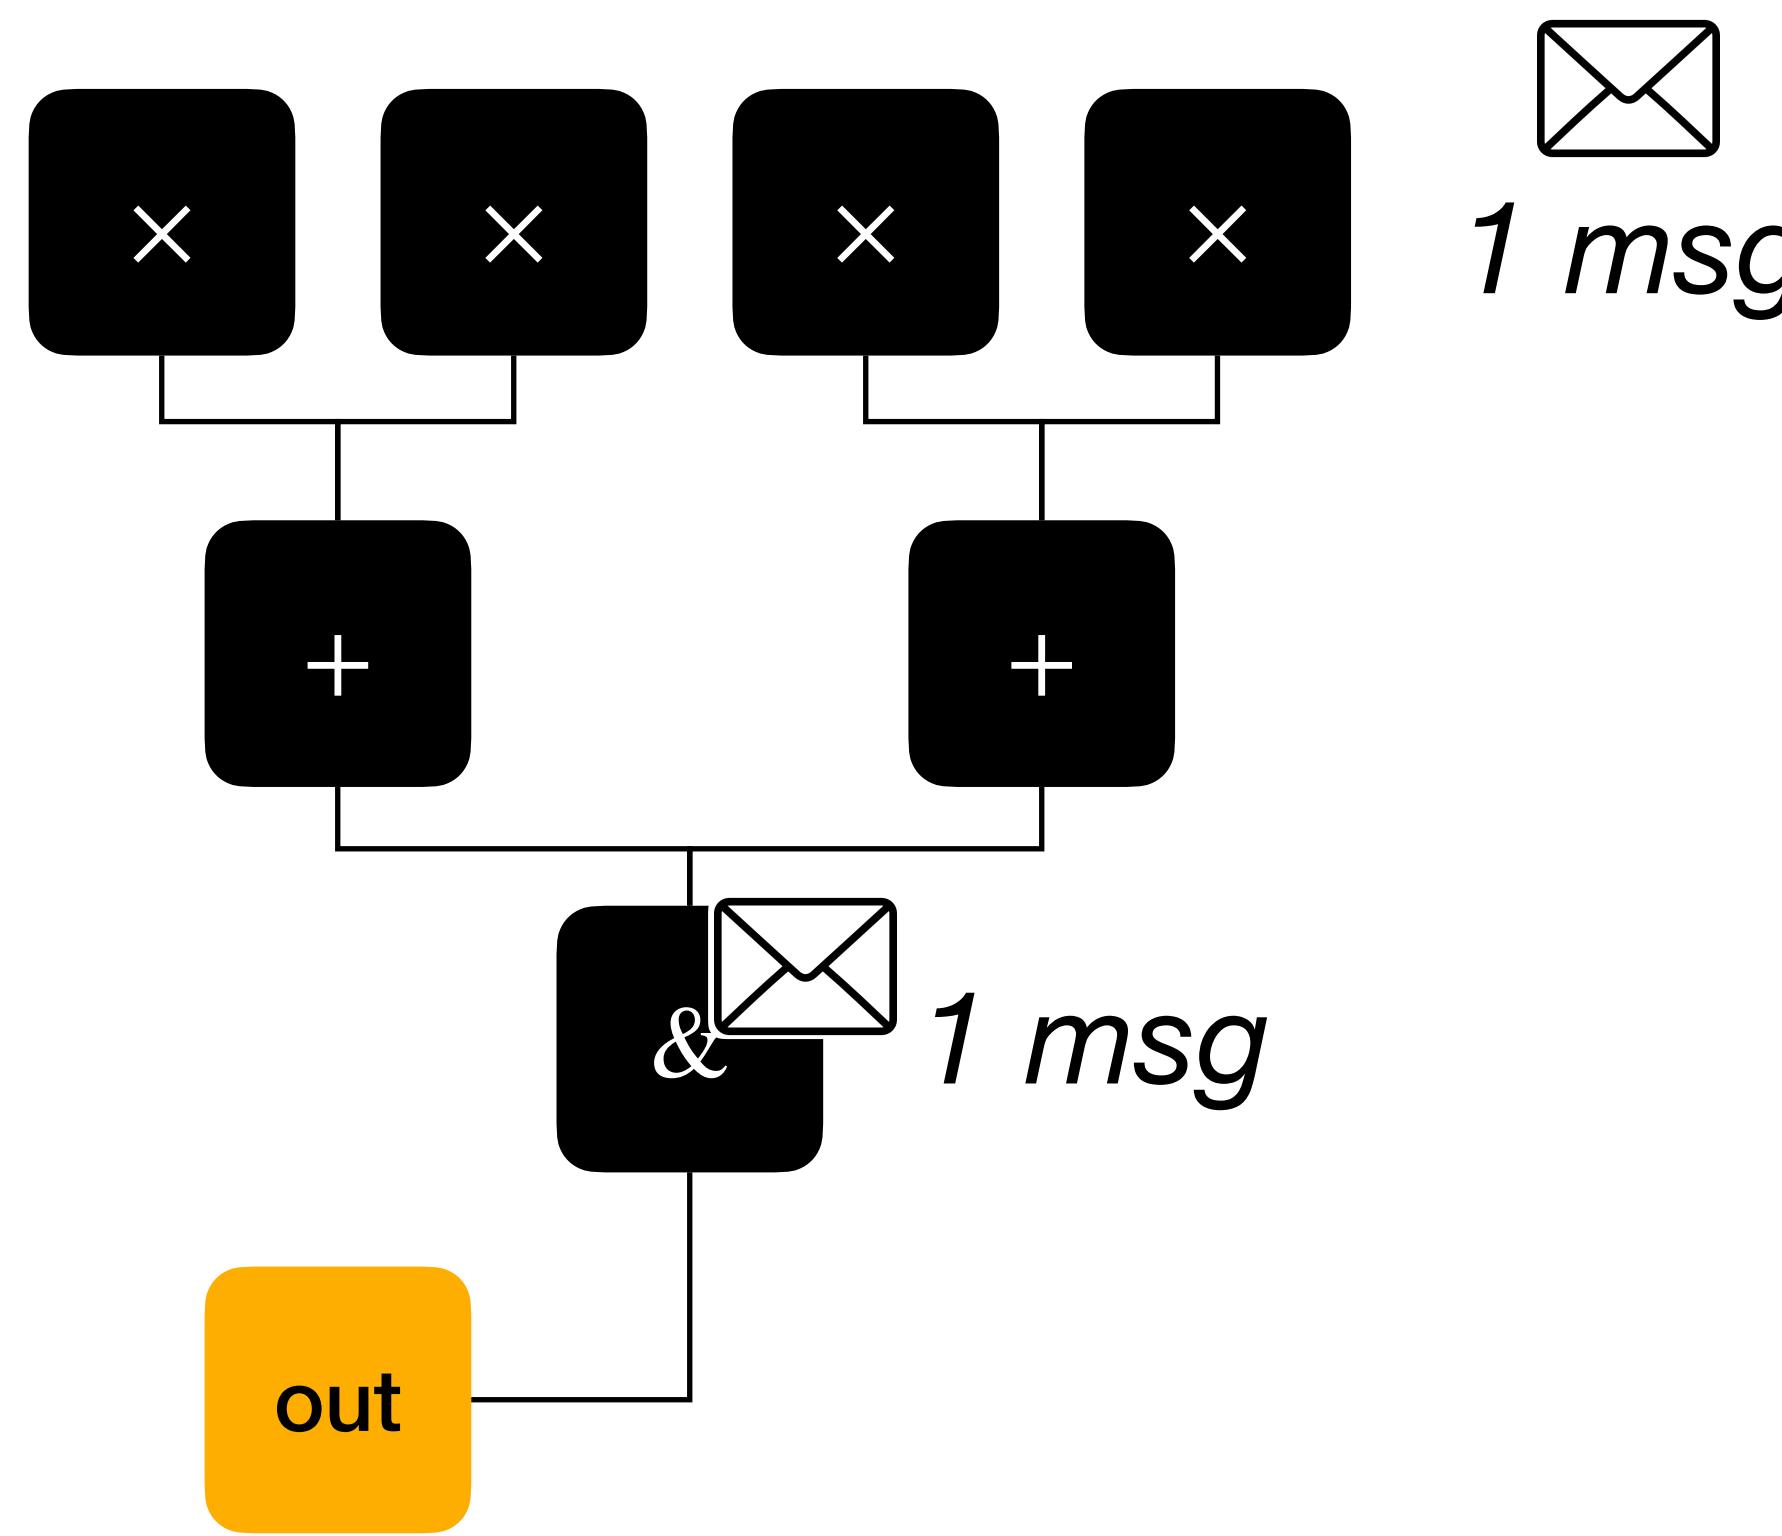
# Vectorization & Message Batching



# Vectorization & Message Batching



# Vectorization & Message Batching



Flexible TCP Communicator

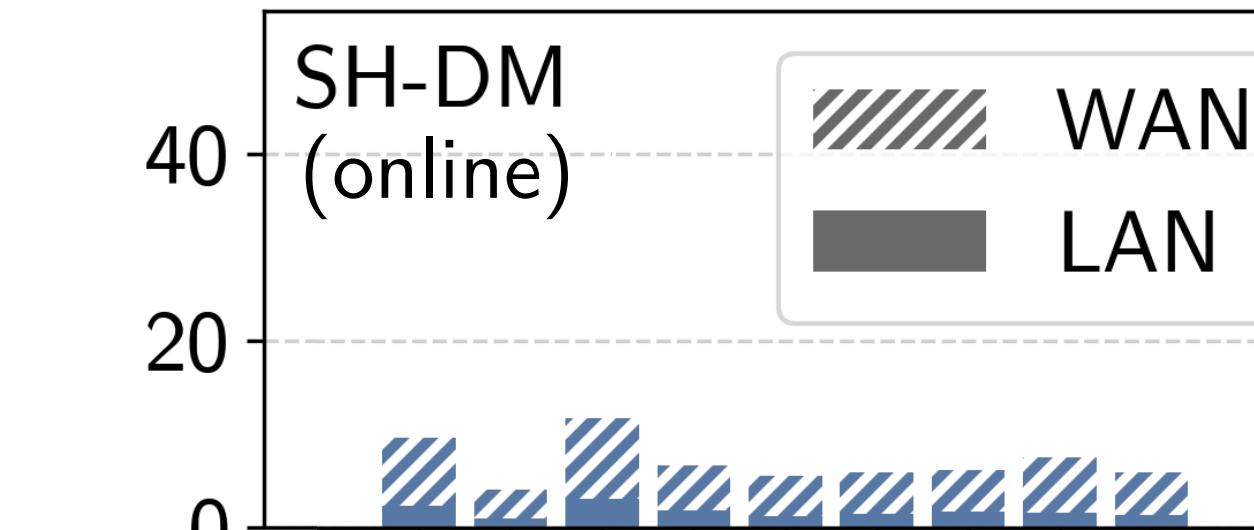
# Prior MPC Benchmarks

with ~5M input rows

# Prior MPC Benchmarks

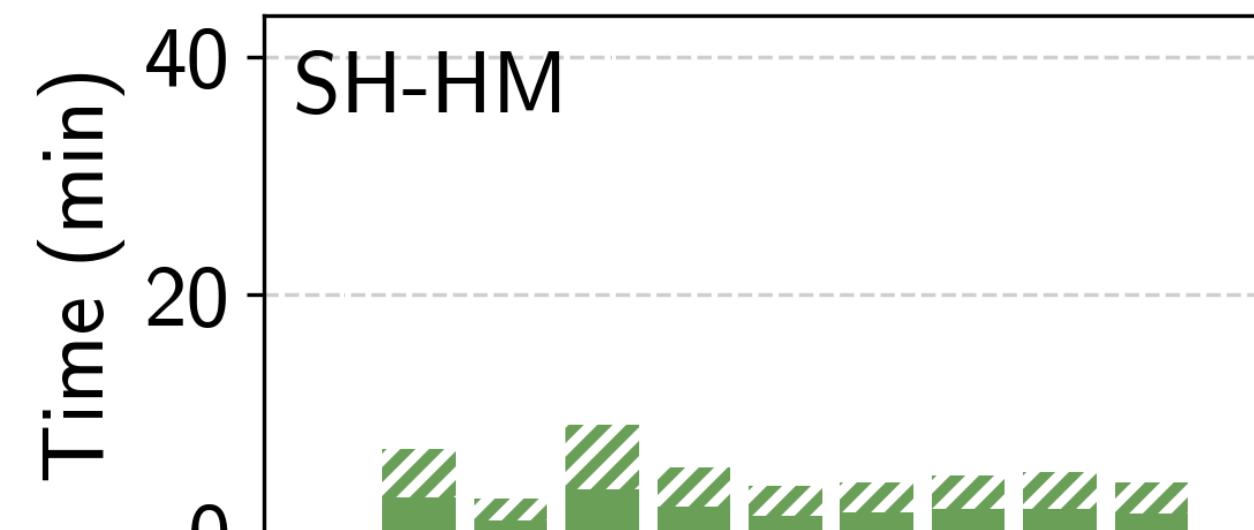
with ~5M input rows

**ABY 2PC**  
Semihonest / Dishonest Majority



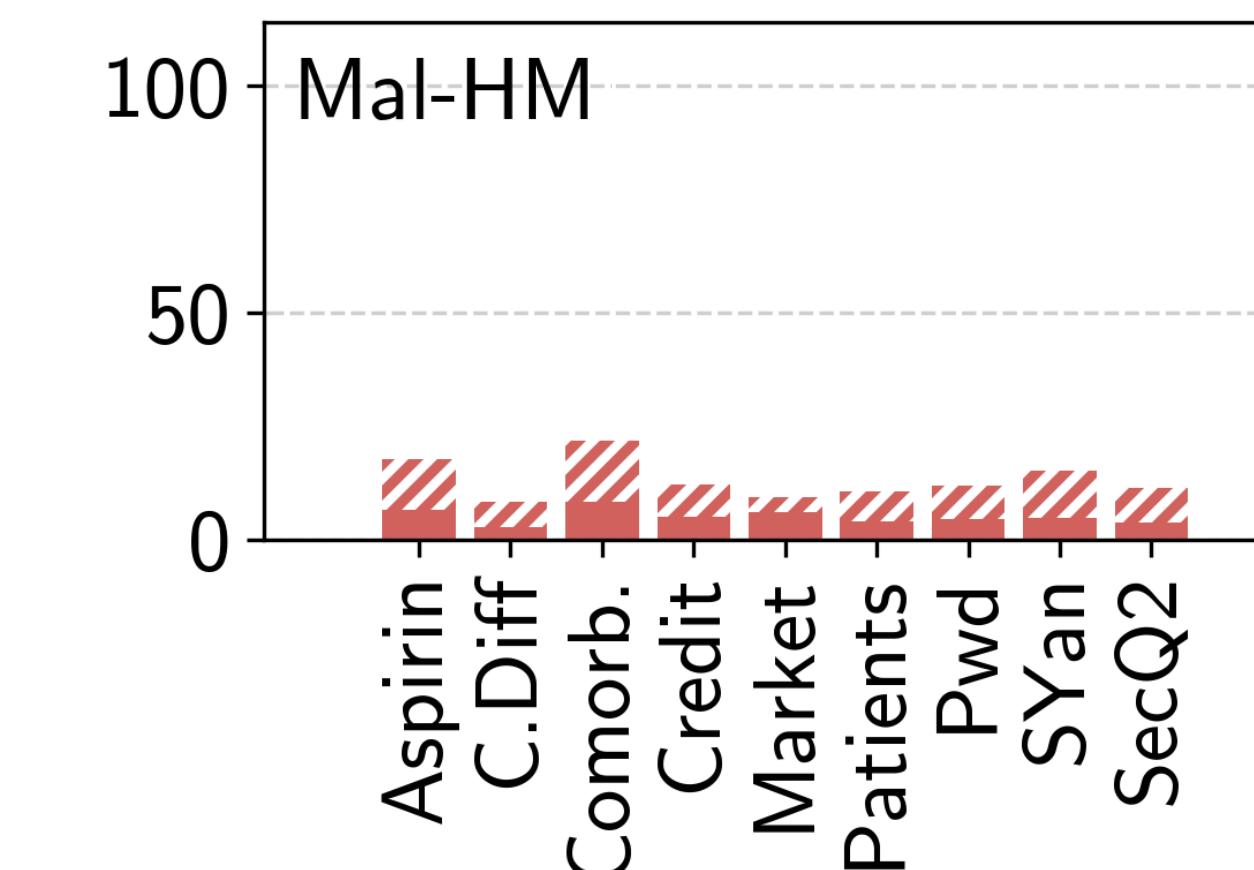
all  $\leq$  3 min

**Replicated 3PC**  
Semihonest / Honest Majority



all  $\leq$  4 min

**Fantastic 4PC**  
Malicious / Honest Majority

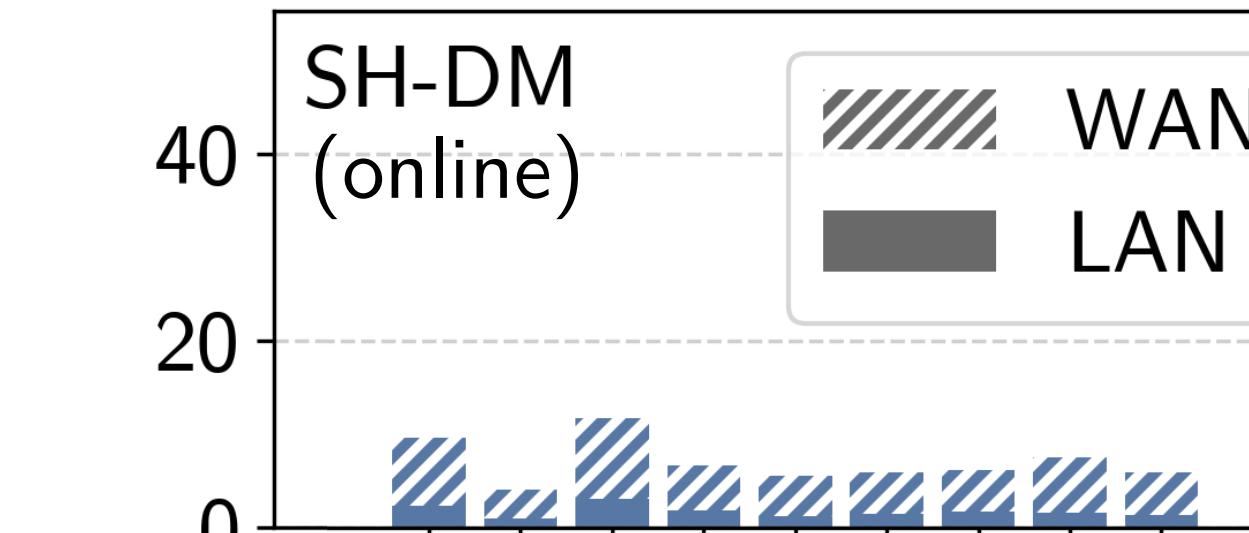


all  $\leq$  9 min

# Prior MPC Benchmarks

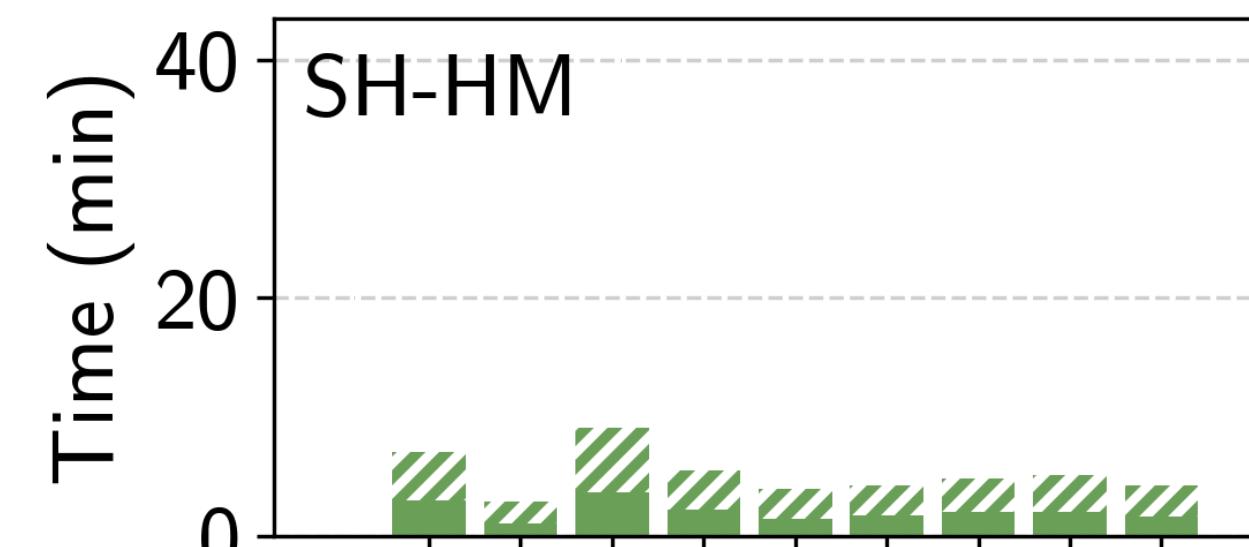
with ~5M input rows

**ABY 2PC**  
Semihonest / Dishonest Majority



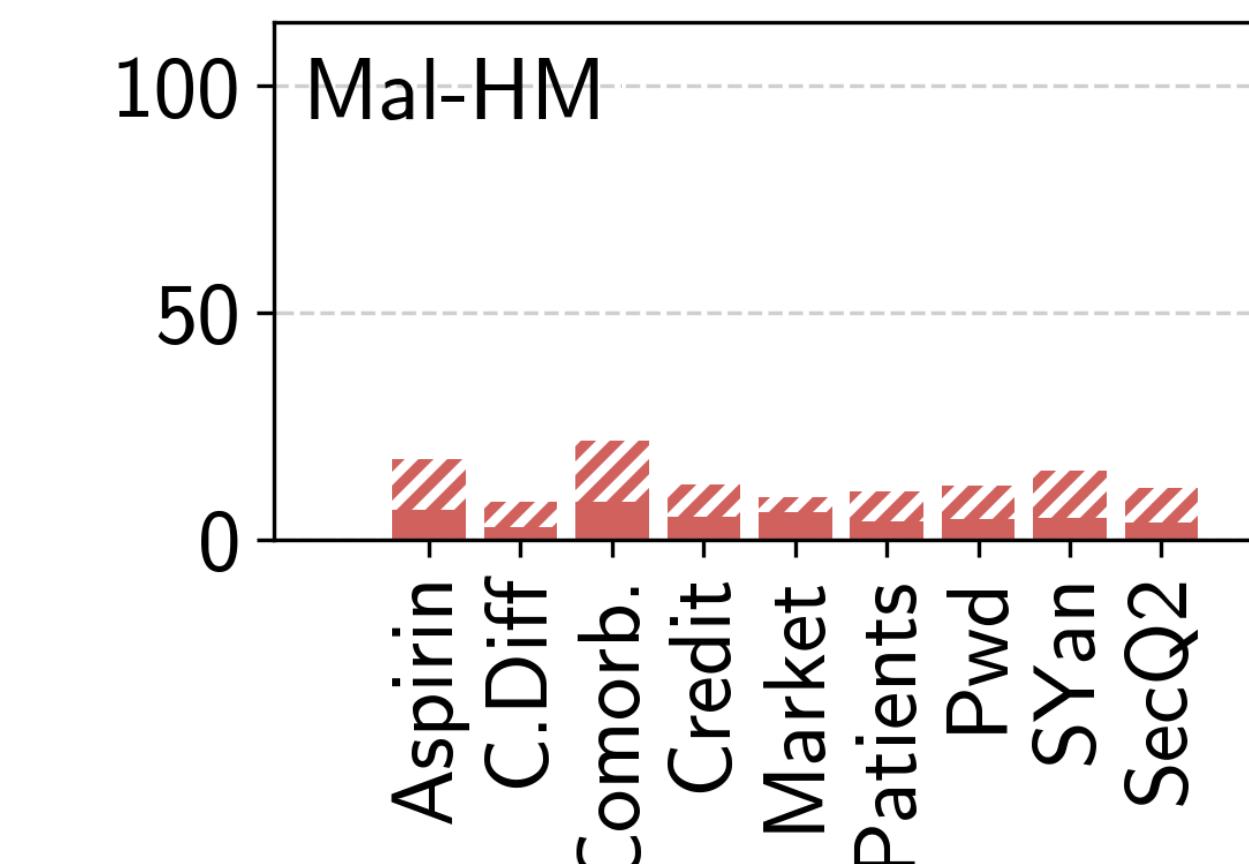
all  $\leq$  3 min

**Replicated 3PC**  
Semihonest / Honest Majority



all  $\leq$  4 min

**Fantastic 4PC**  
Malicious / Honest Majority



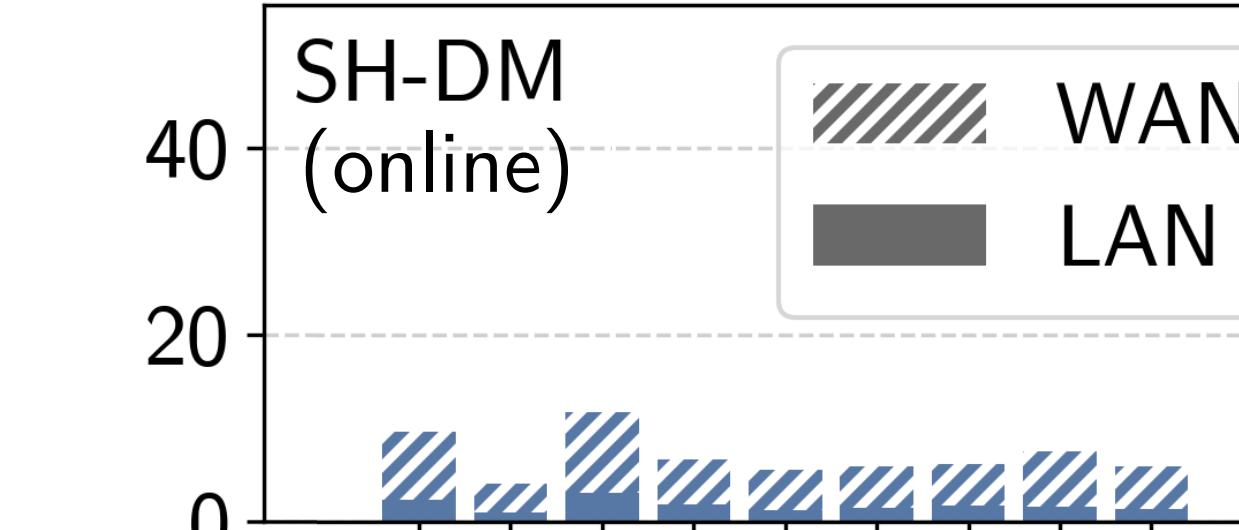
Motivation for leakage!

all  $\leq$  9 min

# Prior MPC Benchmarks

with ~5M input rows

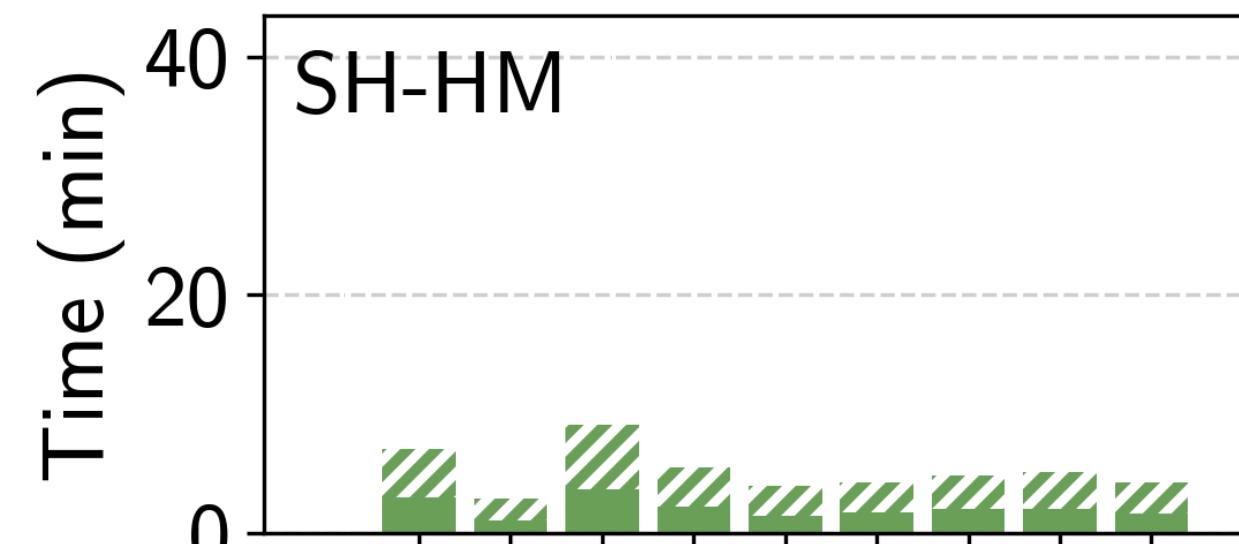
**ABY 2PC**  
Semihonest / Dishonest Majority



6 Gbps 20 ms  
25 Gbps 0.3 ms

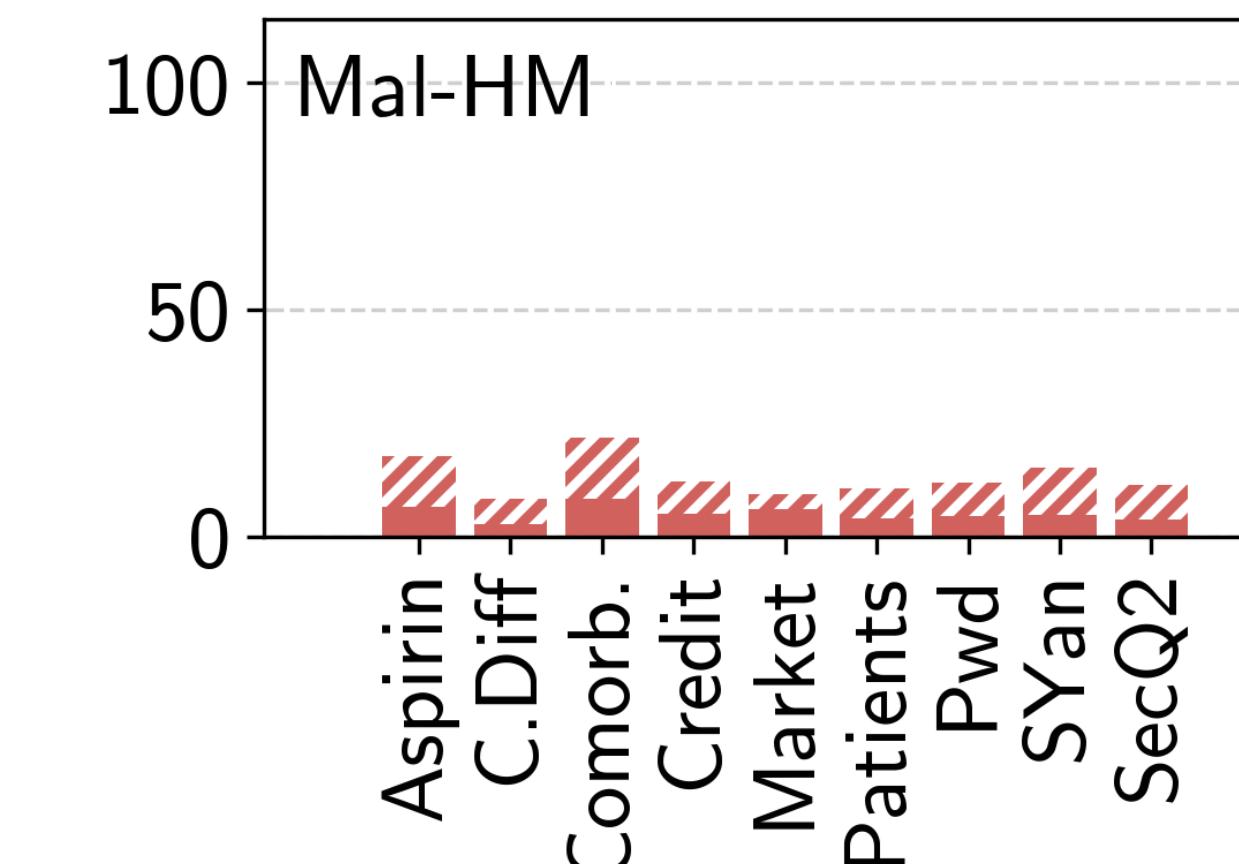
all  $\leq$  3 min

**Replicated 3PC**  
Semihonest / Honest Majority



all  $\leq$  4 min

**Fantastic 4PC**  
Malicious / Honest Majority

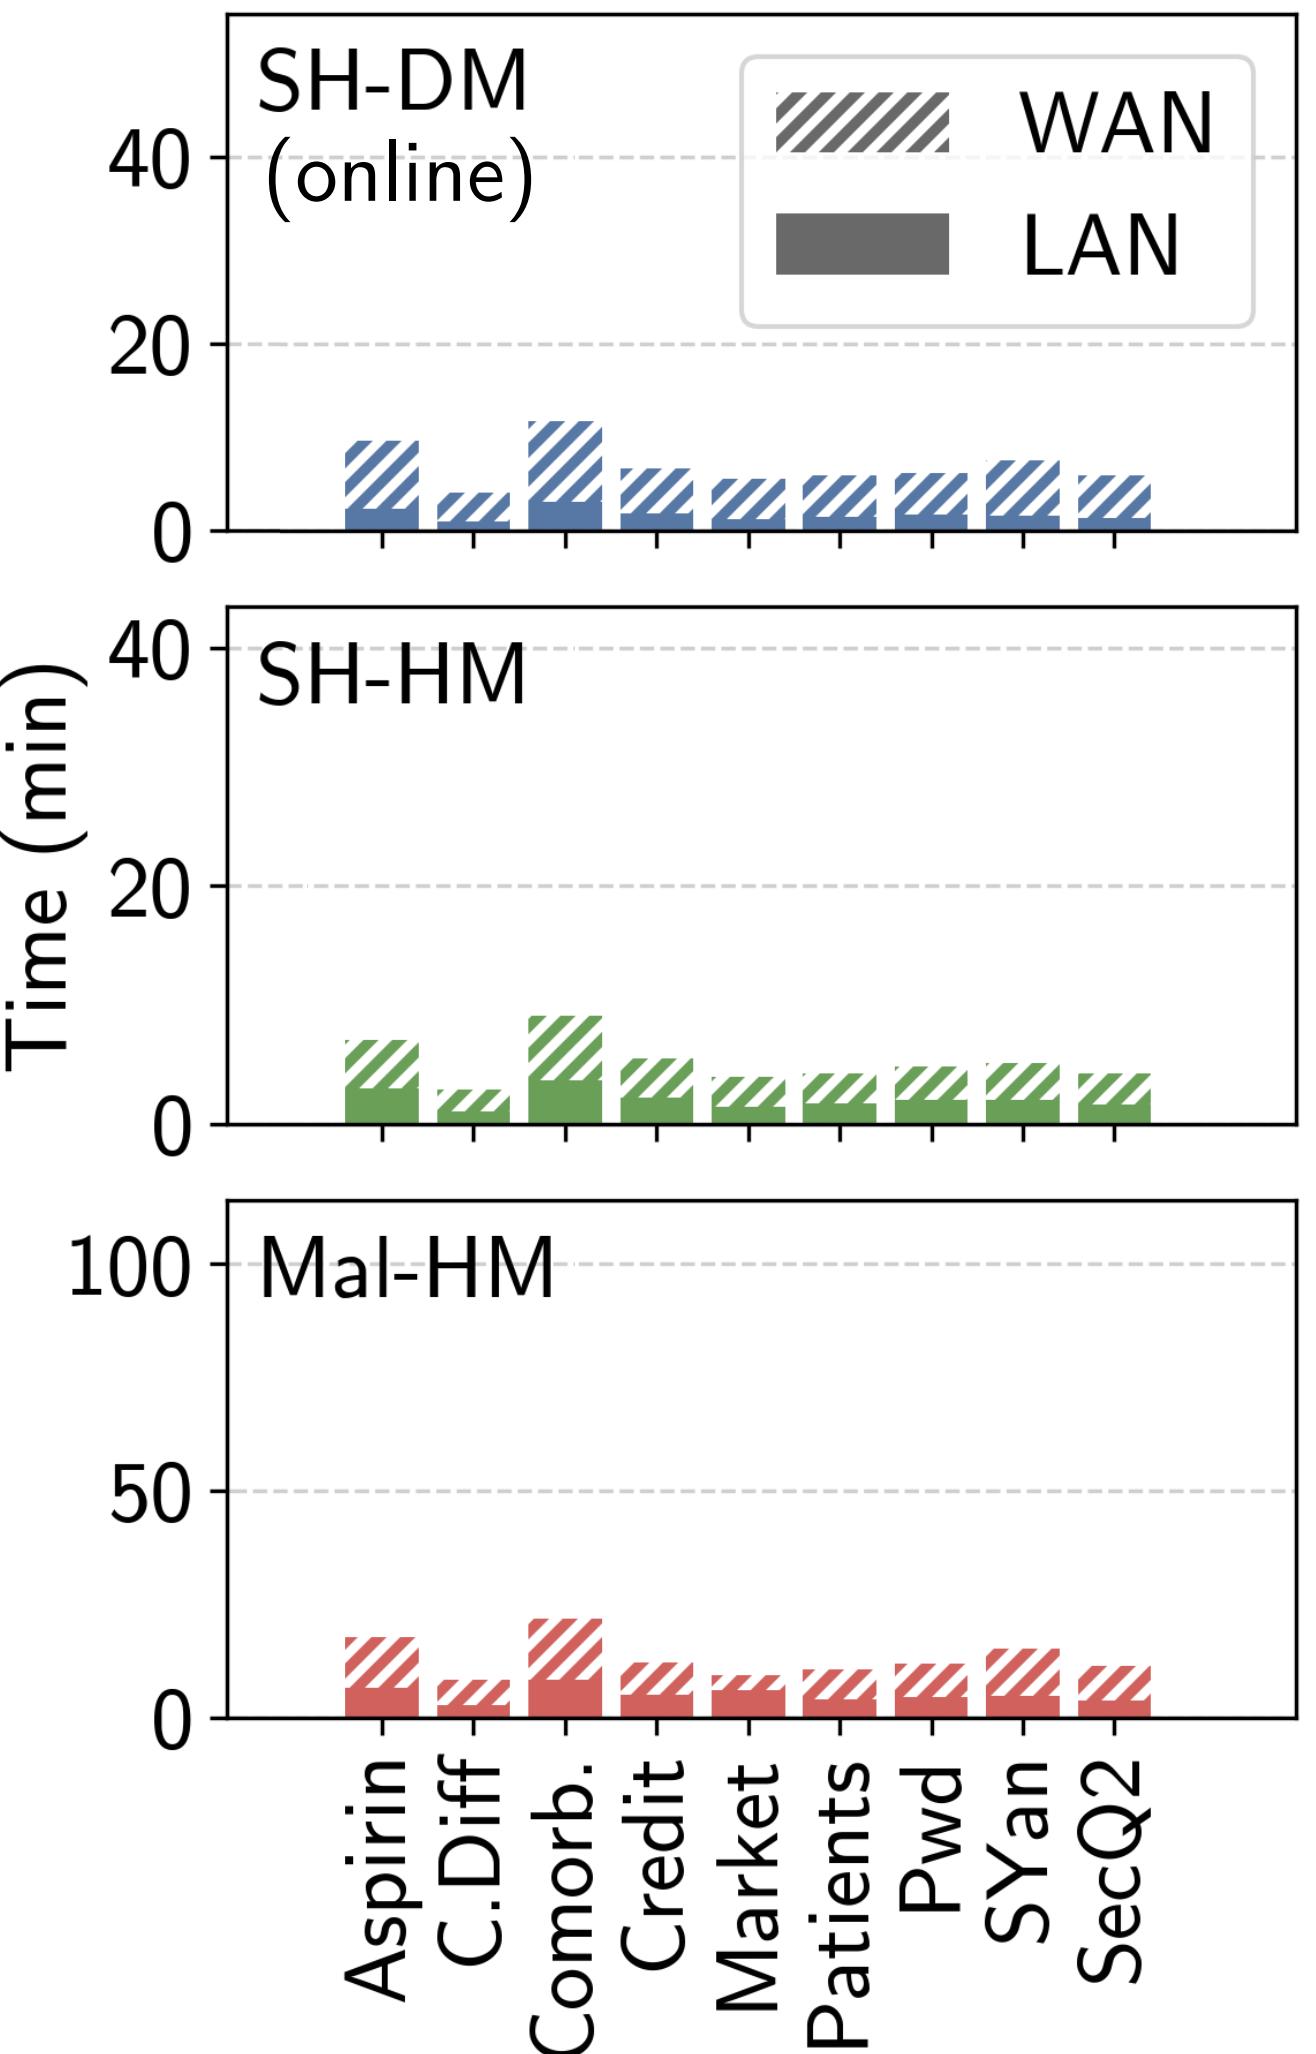


Motivation for leakage!

all  $\leq$  9 min

# Prior MPC Benchmarks

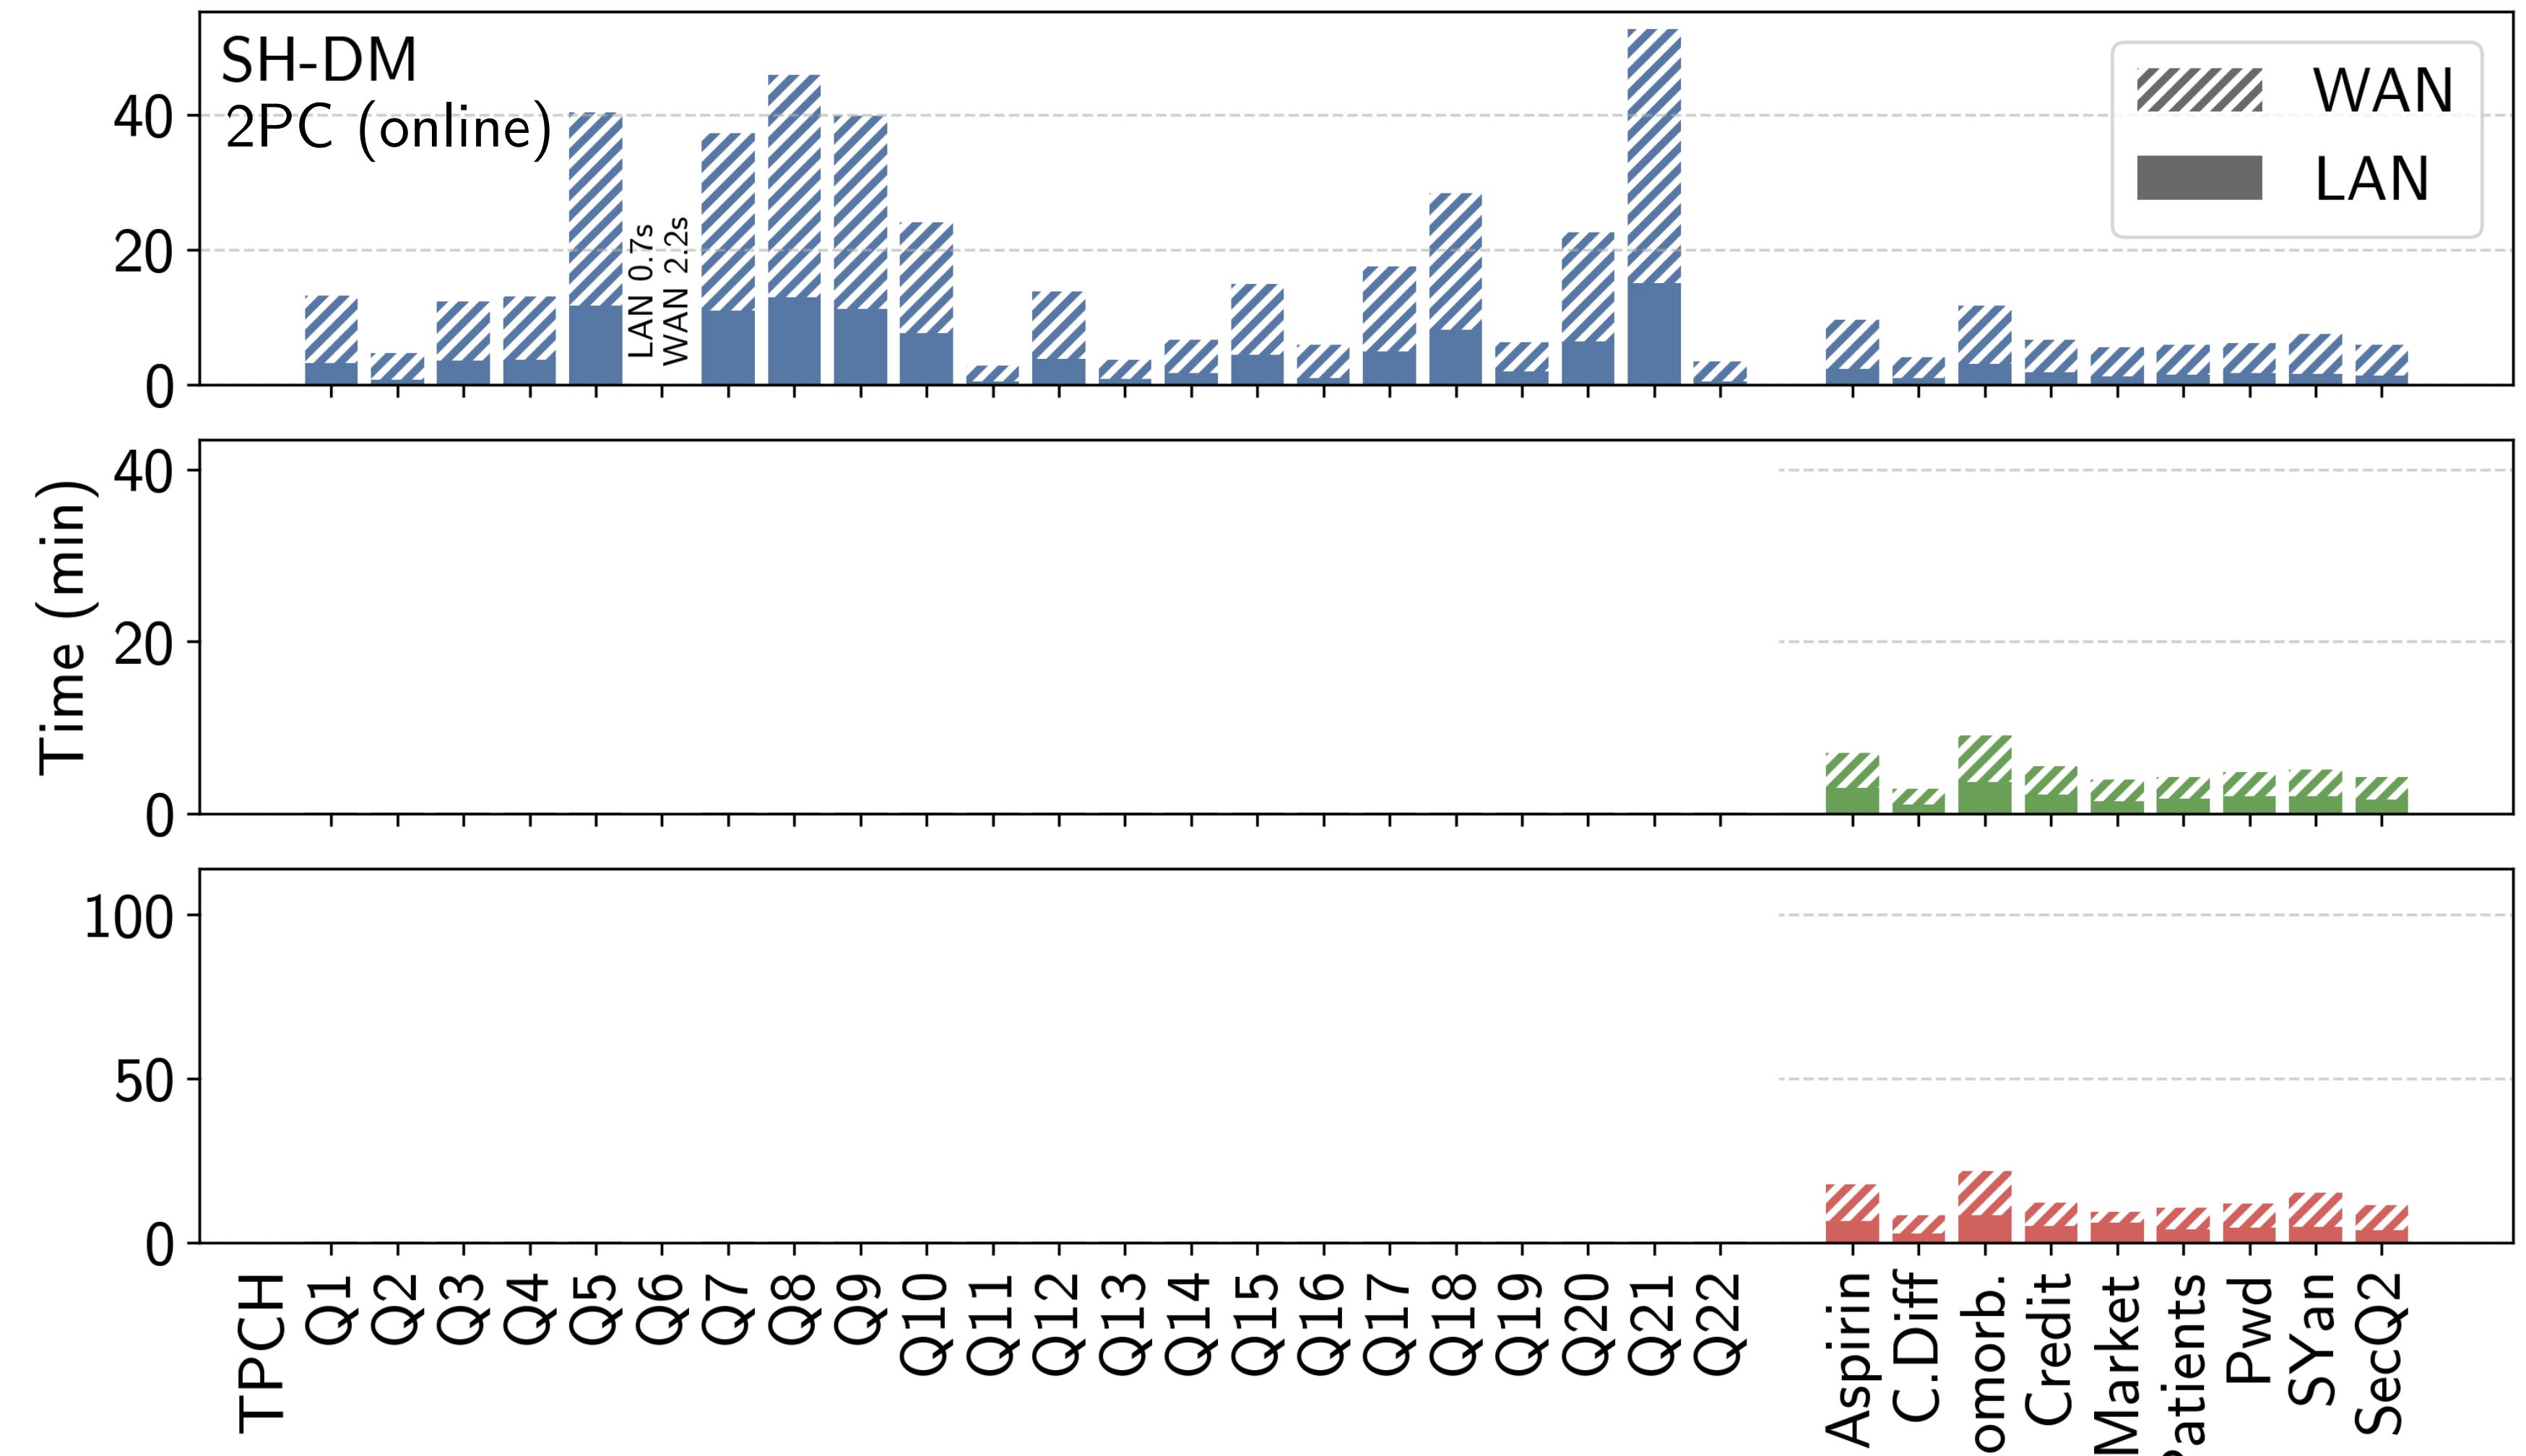
with ~5M input rows



# Full TPC-H Benchmark under MPC

@ Scale Factor 1 ( $\approx 1$  GB input)

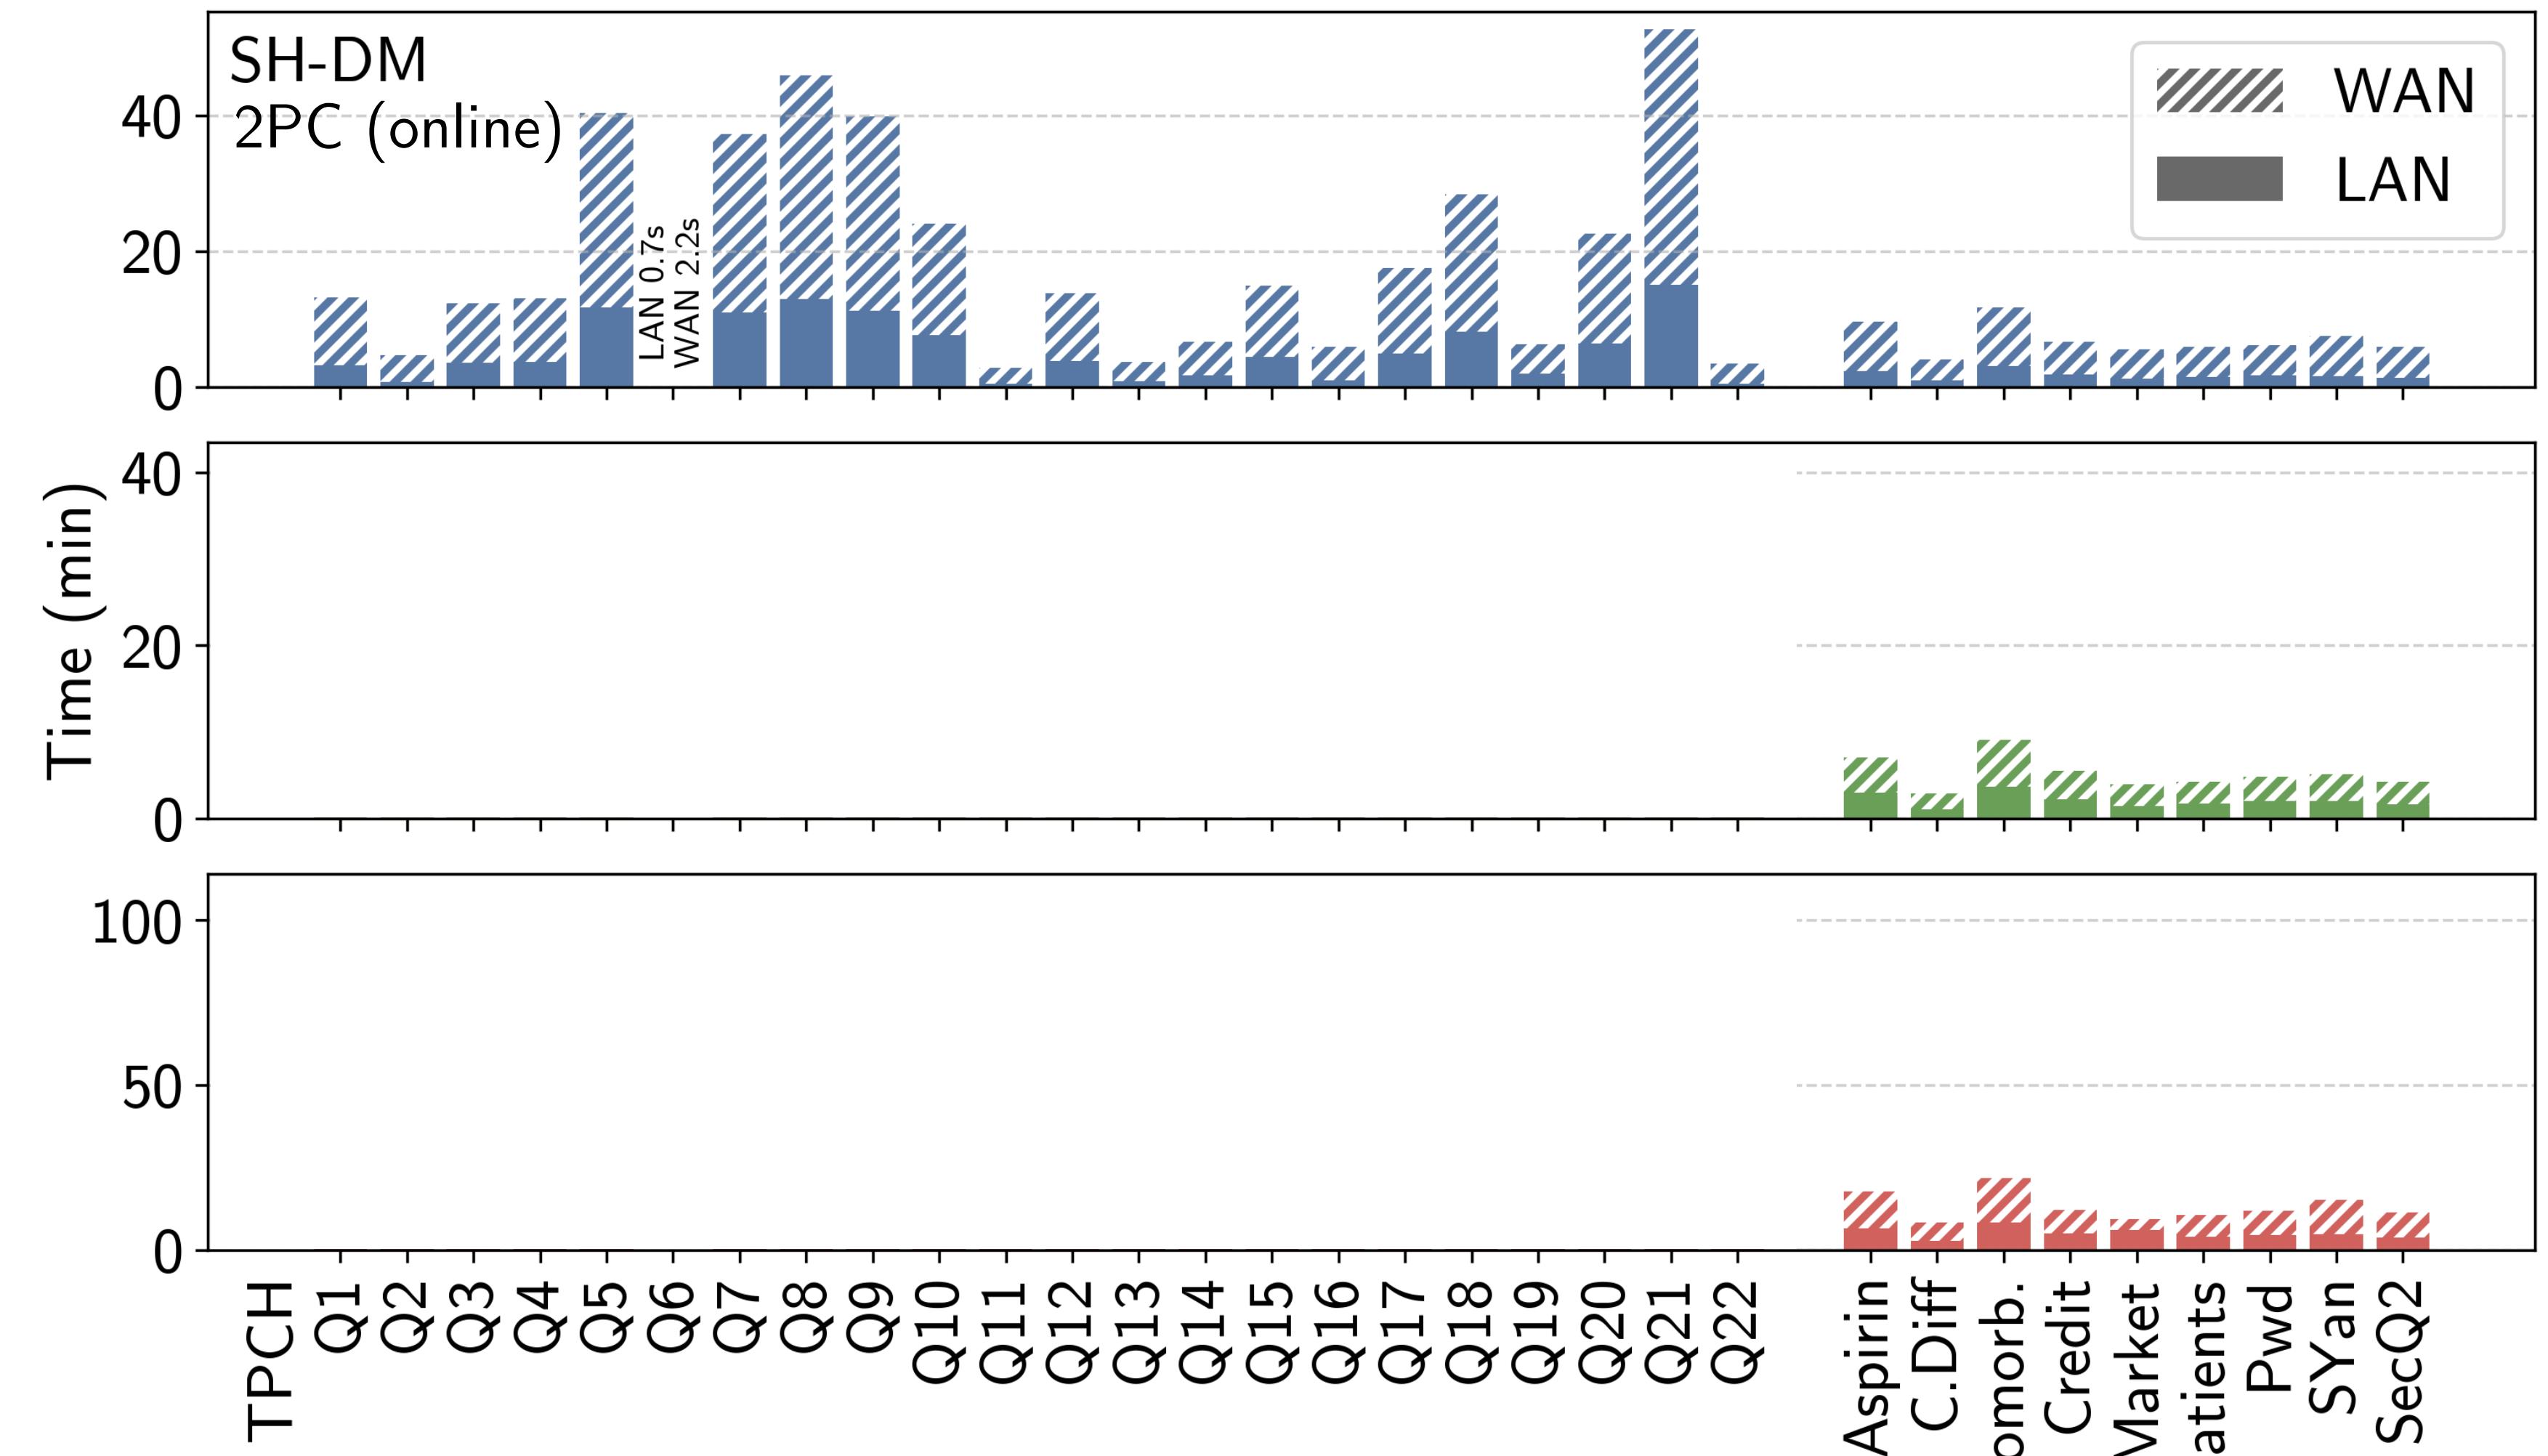
TPC-H		
	Median	Max
SH-DM LAN	3.8	15.0
WAN	13.5	52.7



# Full TPC-H Benchmark under MPC

@ Scale Factor 1 ( $\approx 1$  GB input)

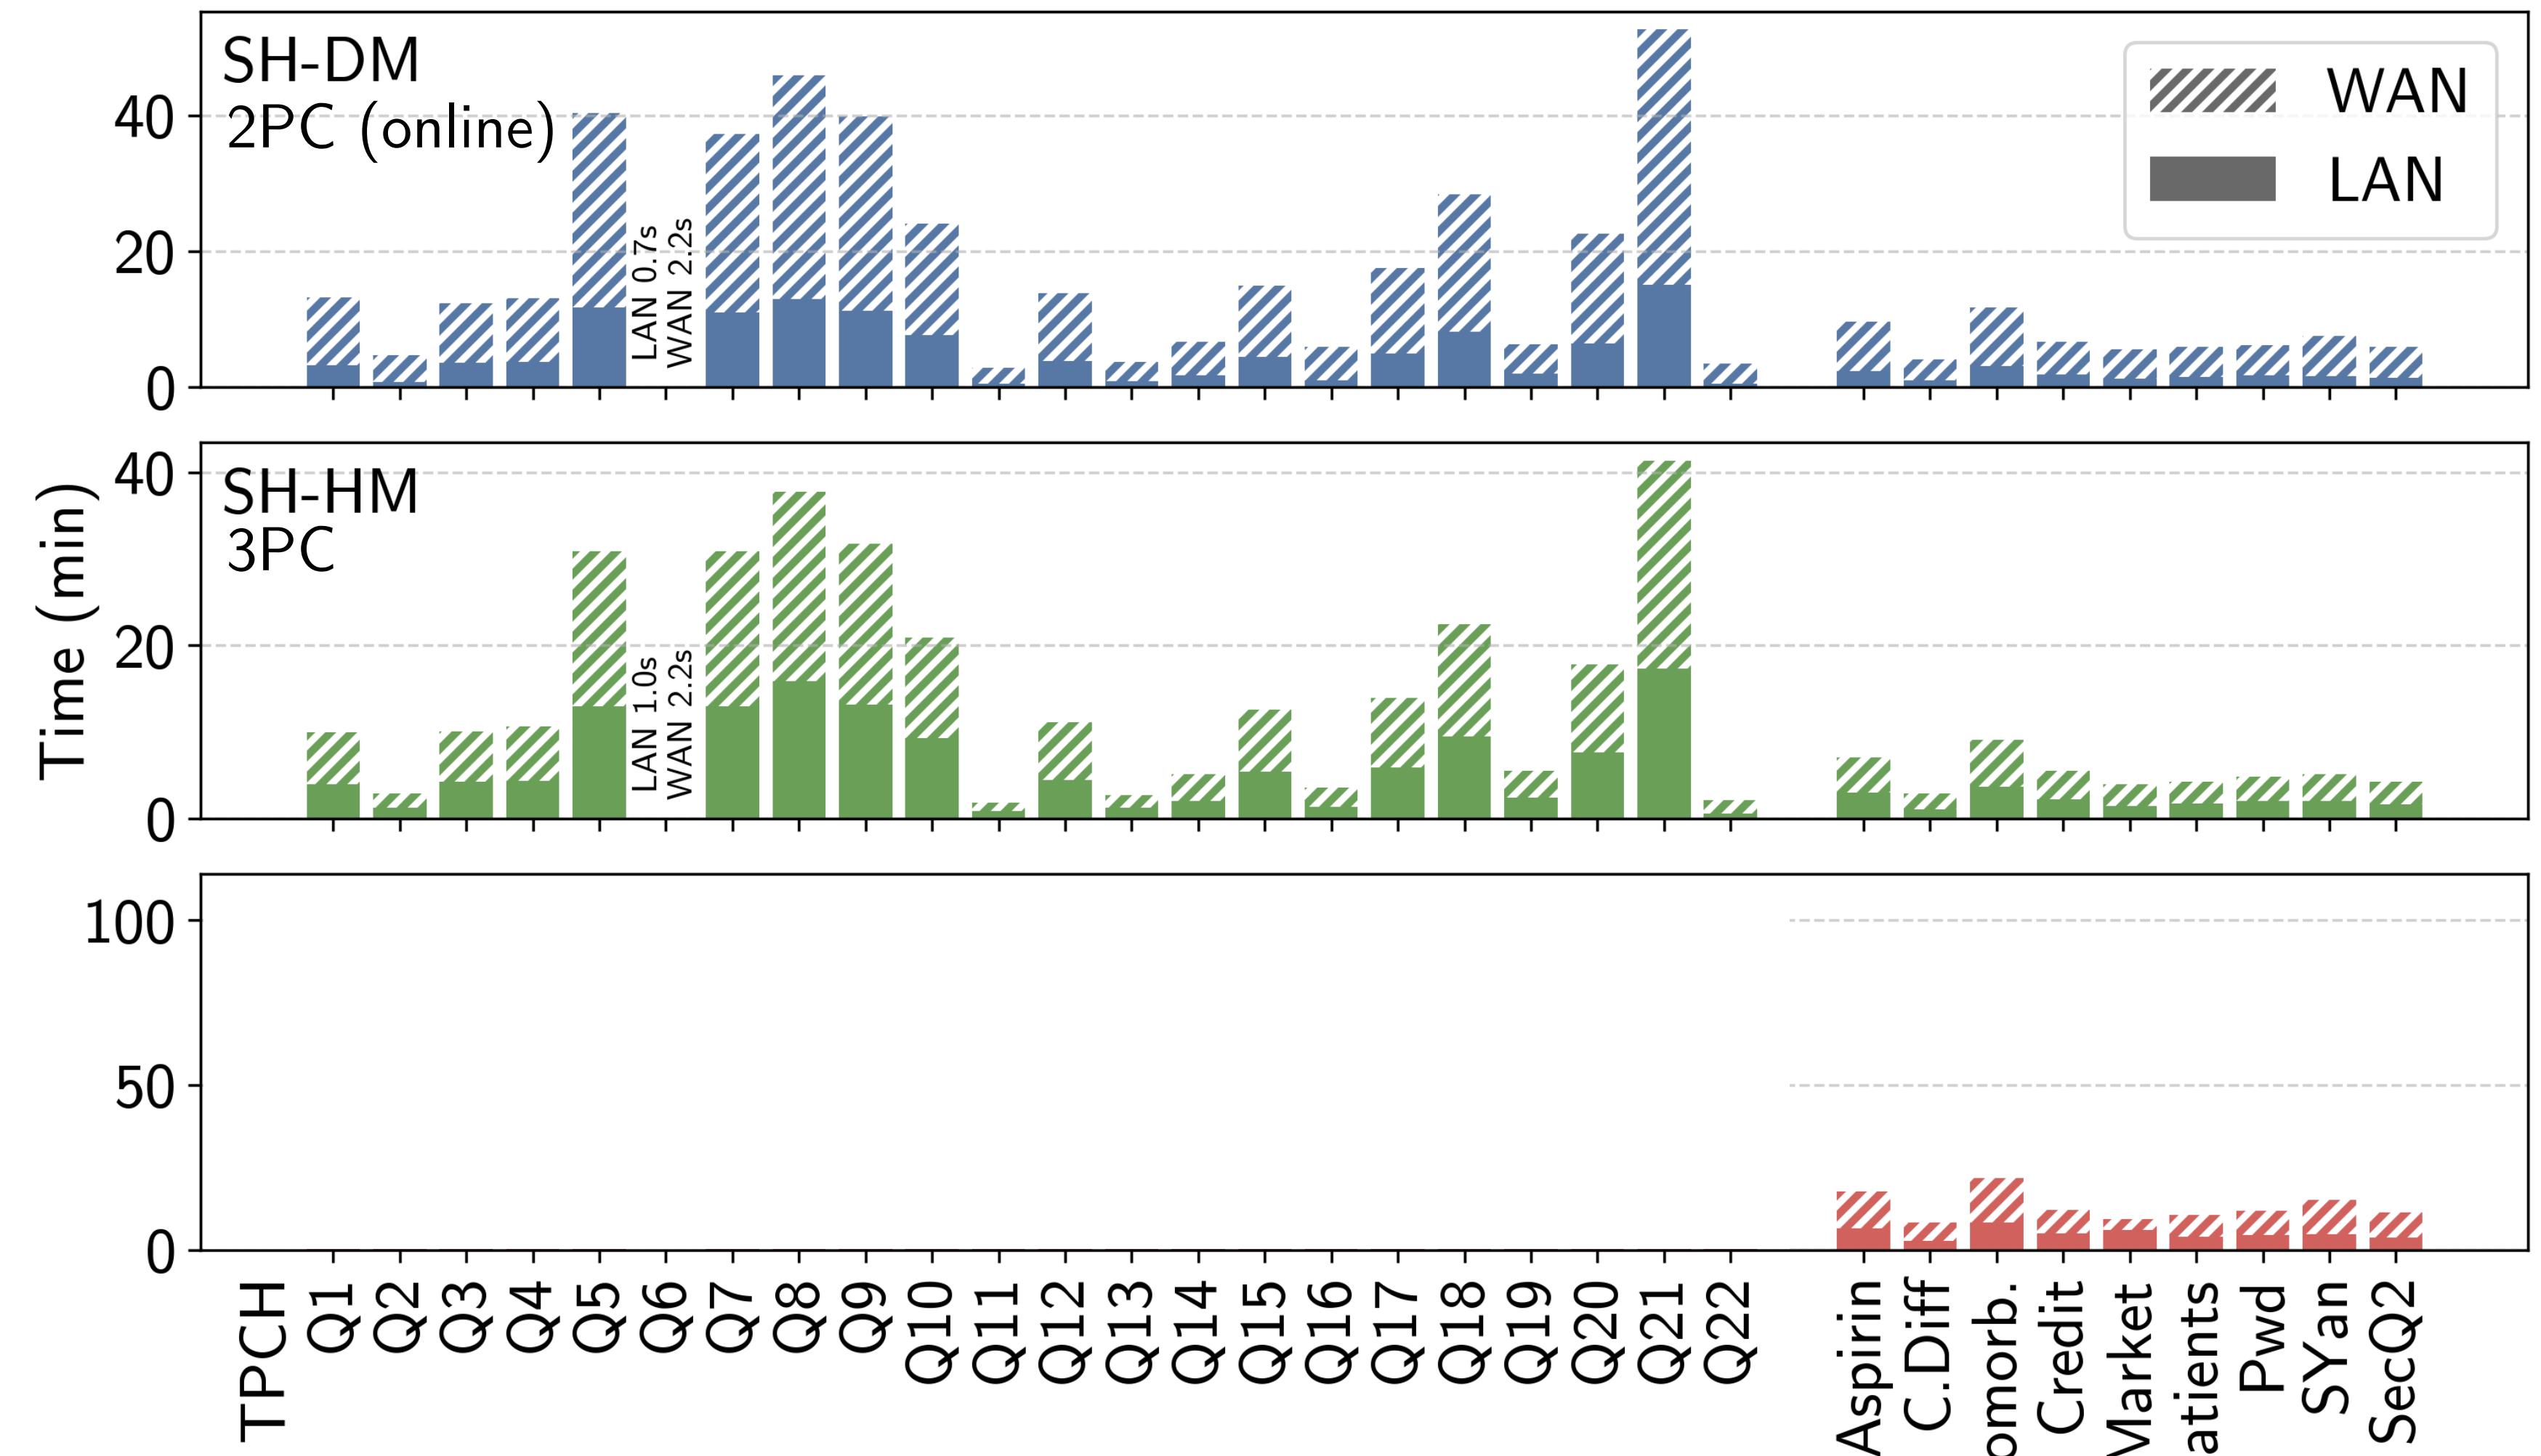
TPC-H	
	minutes
Median	3.8
Max	15.0



# Full TPC-H Benchmark under MPC

@ Scale Factor 1 ( $\approx 1$  GB input)

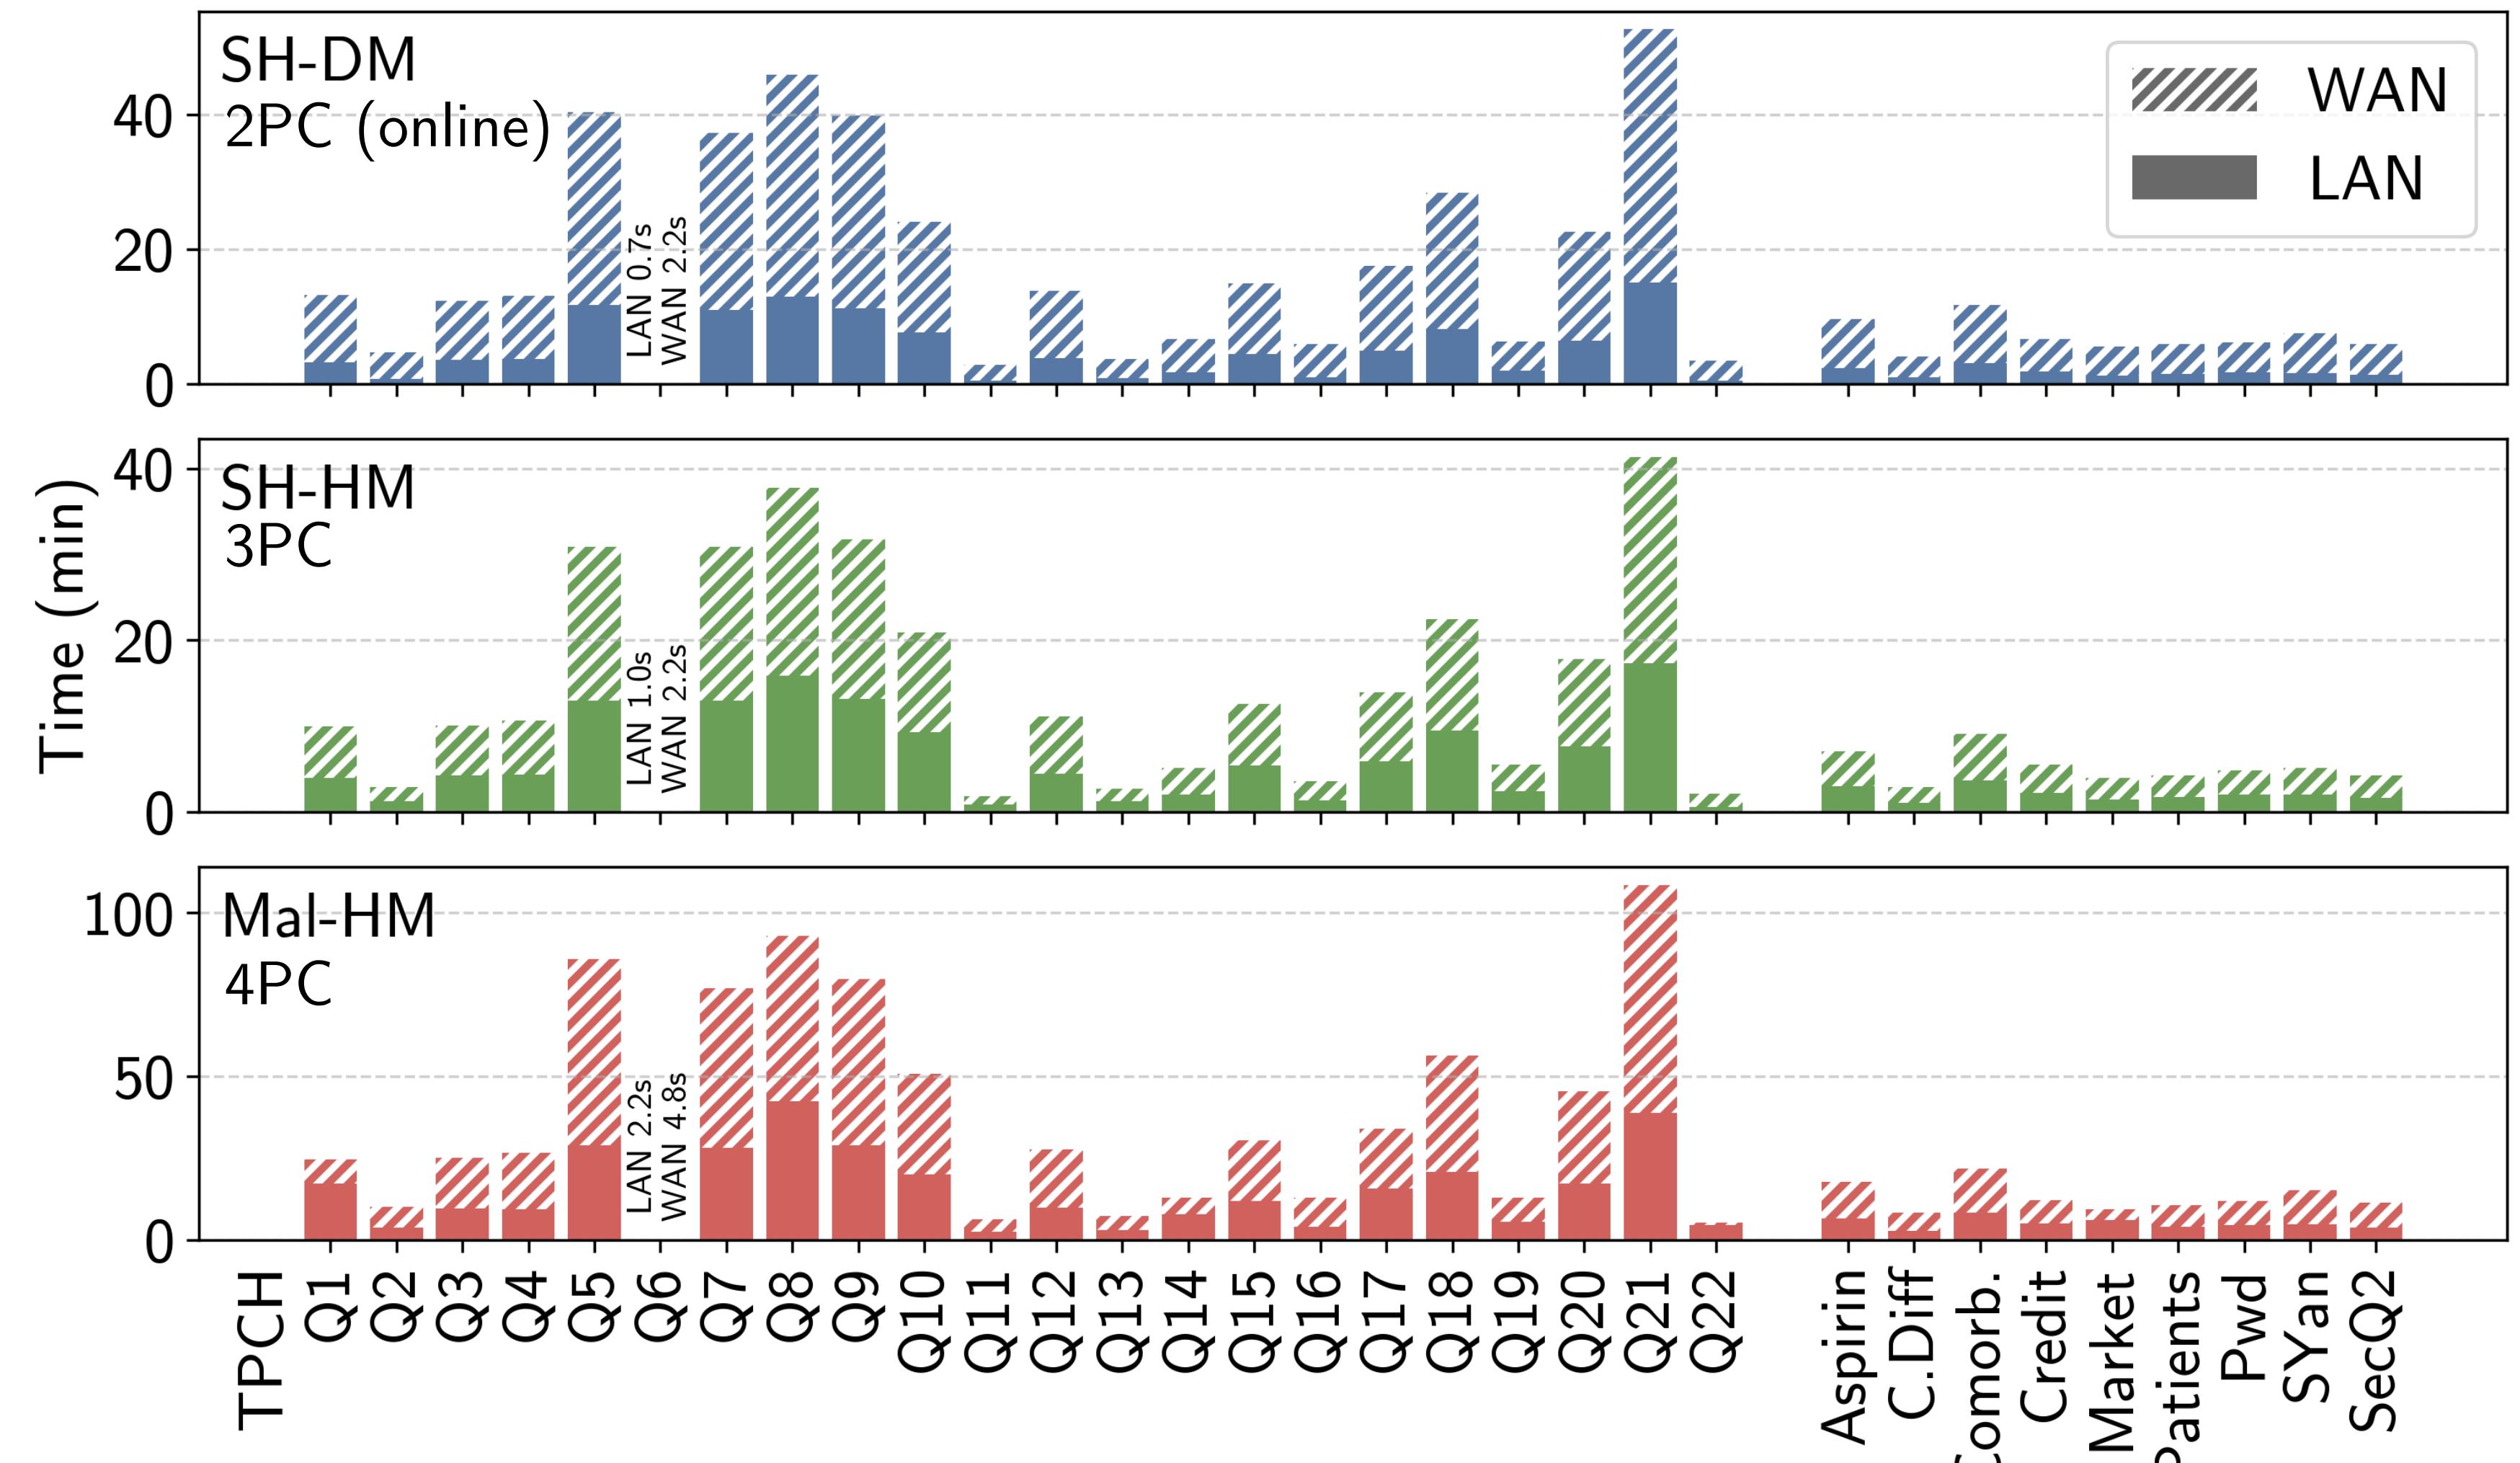
TPC-H		
	Median	Max
SH-DM LAN	3.8	15.0
WAN	13.5	52.7
SH-HM LAN	4.4	17.4
WAN	10.9	41.4



# Full TPC-H Benchmark under MPC

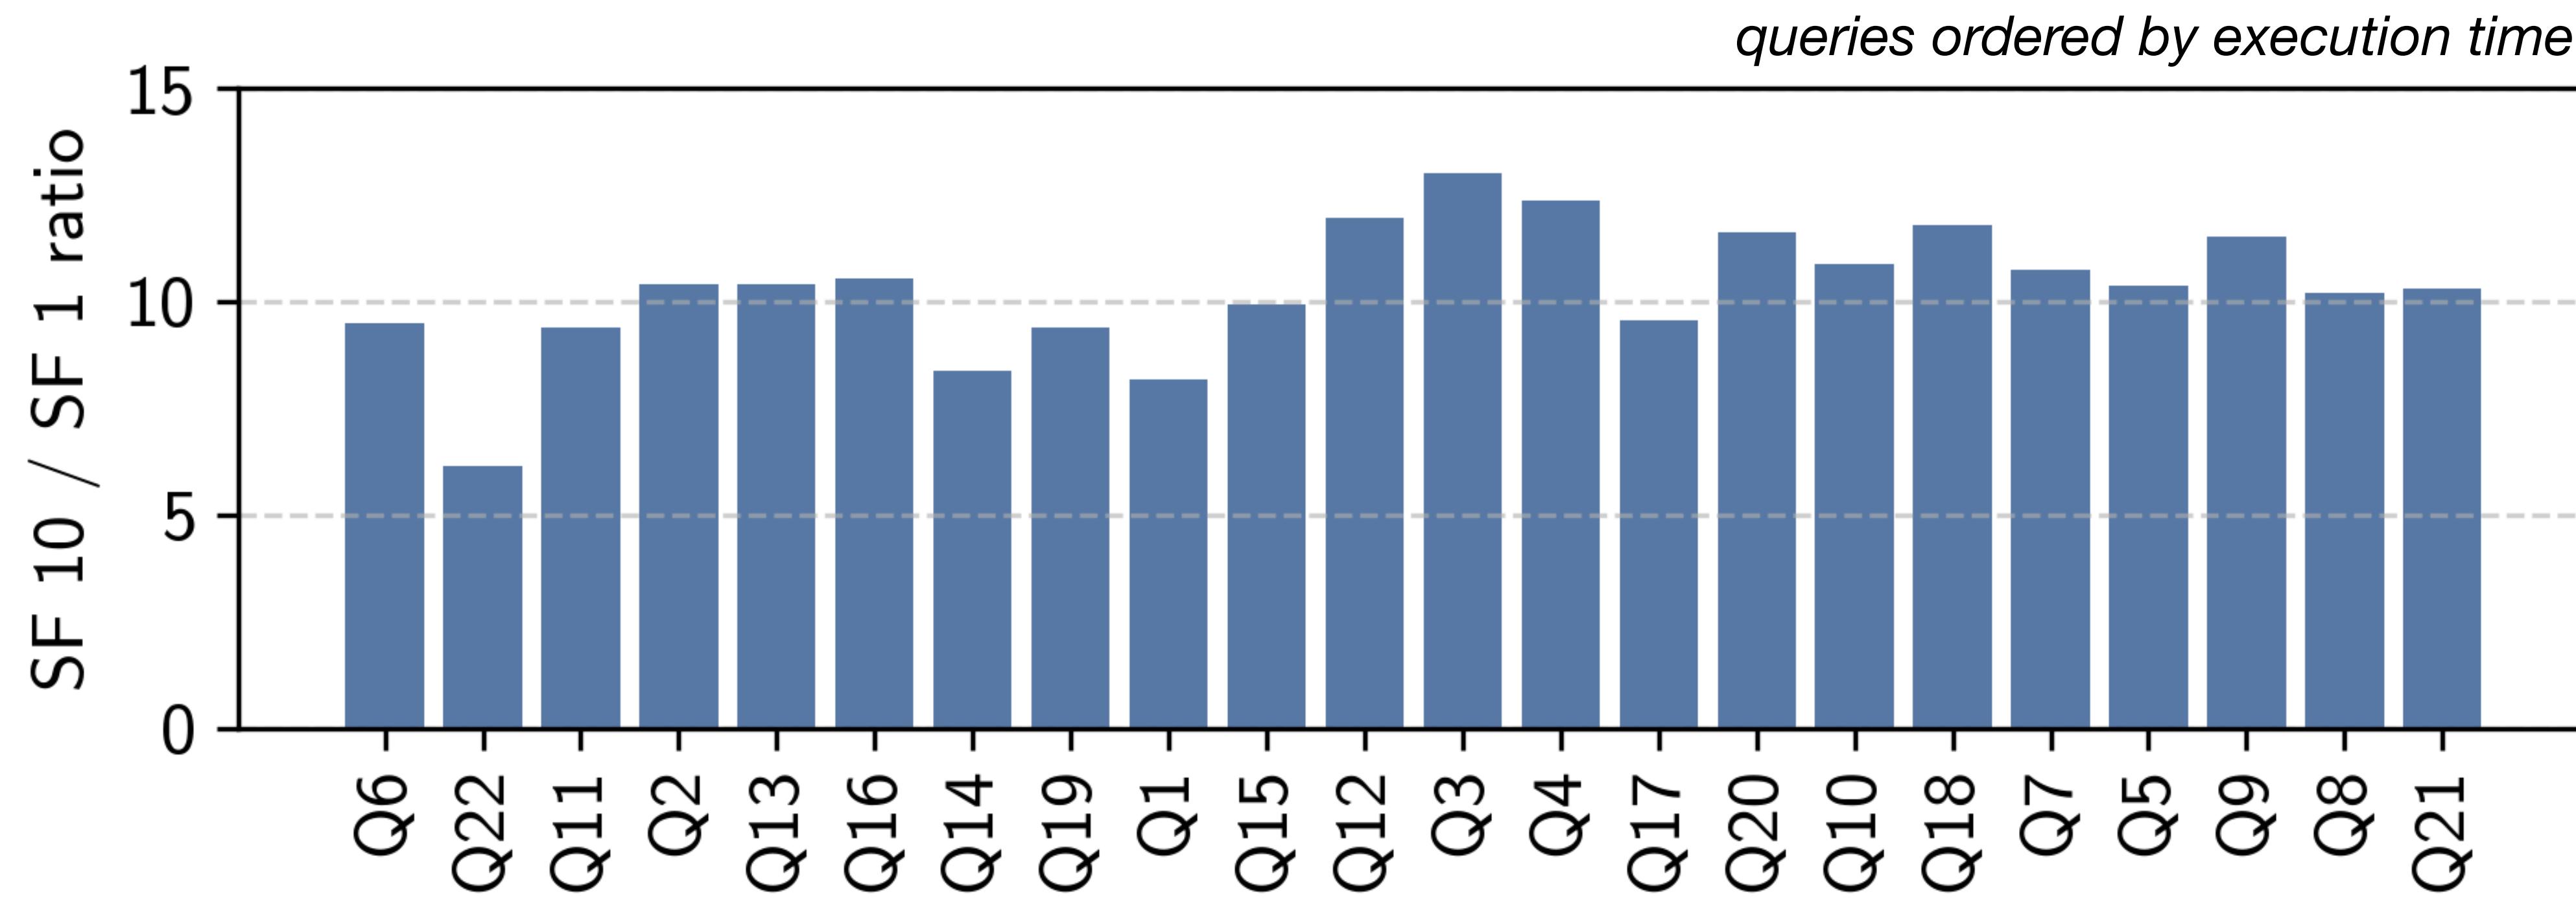
@ Scale Factor 1 ( $\approx 1$  GB input)

	TPC-H	
	Median	Max
SH-DM LAN	3.8	15.0
WAN	13.5	52.7
SH-HM LAN	4.4	17.4
WAN	10.9	41.4
Mal-HM LAN	10.9	42.3
WAN	27.1	108.4



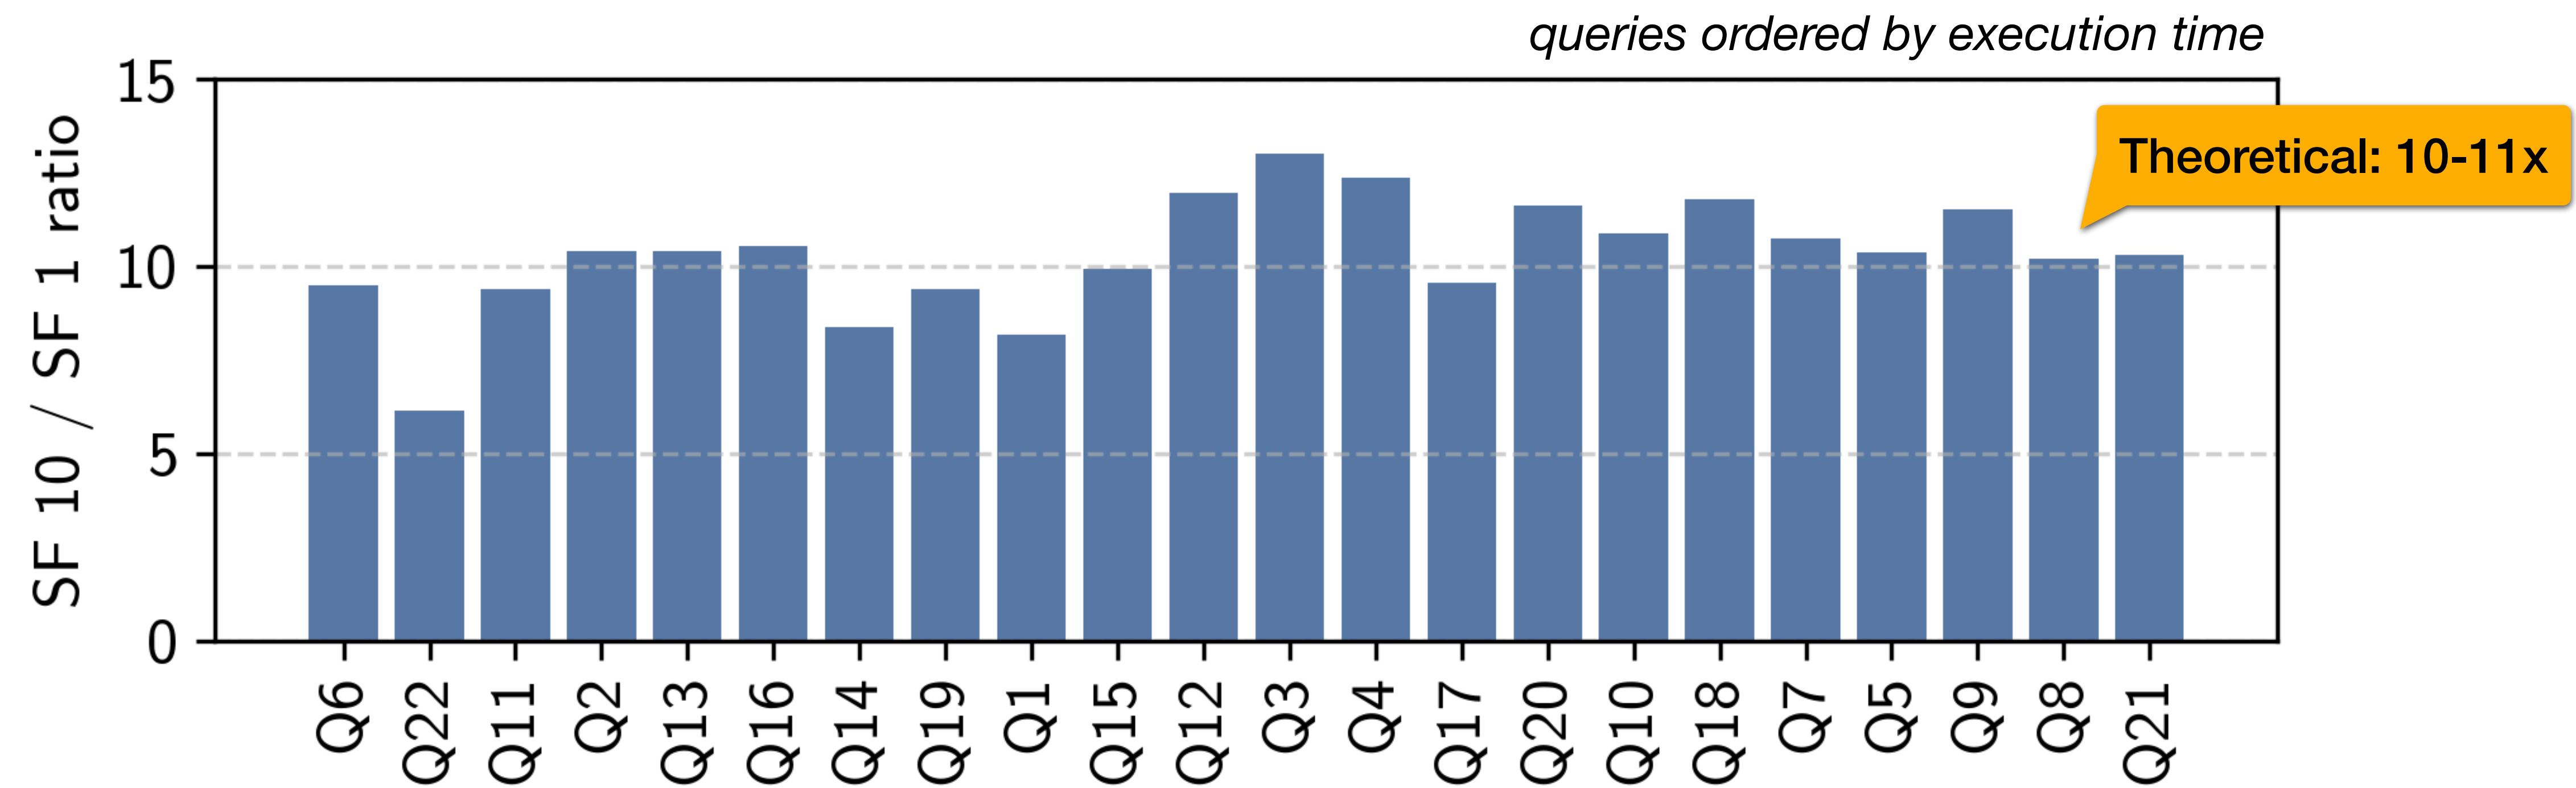
# Scalability

10x larger inputs (SF 10)



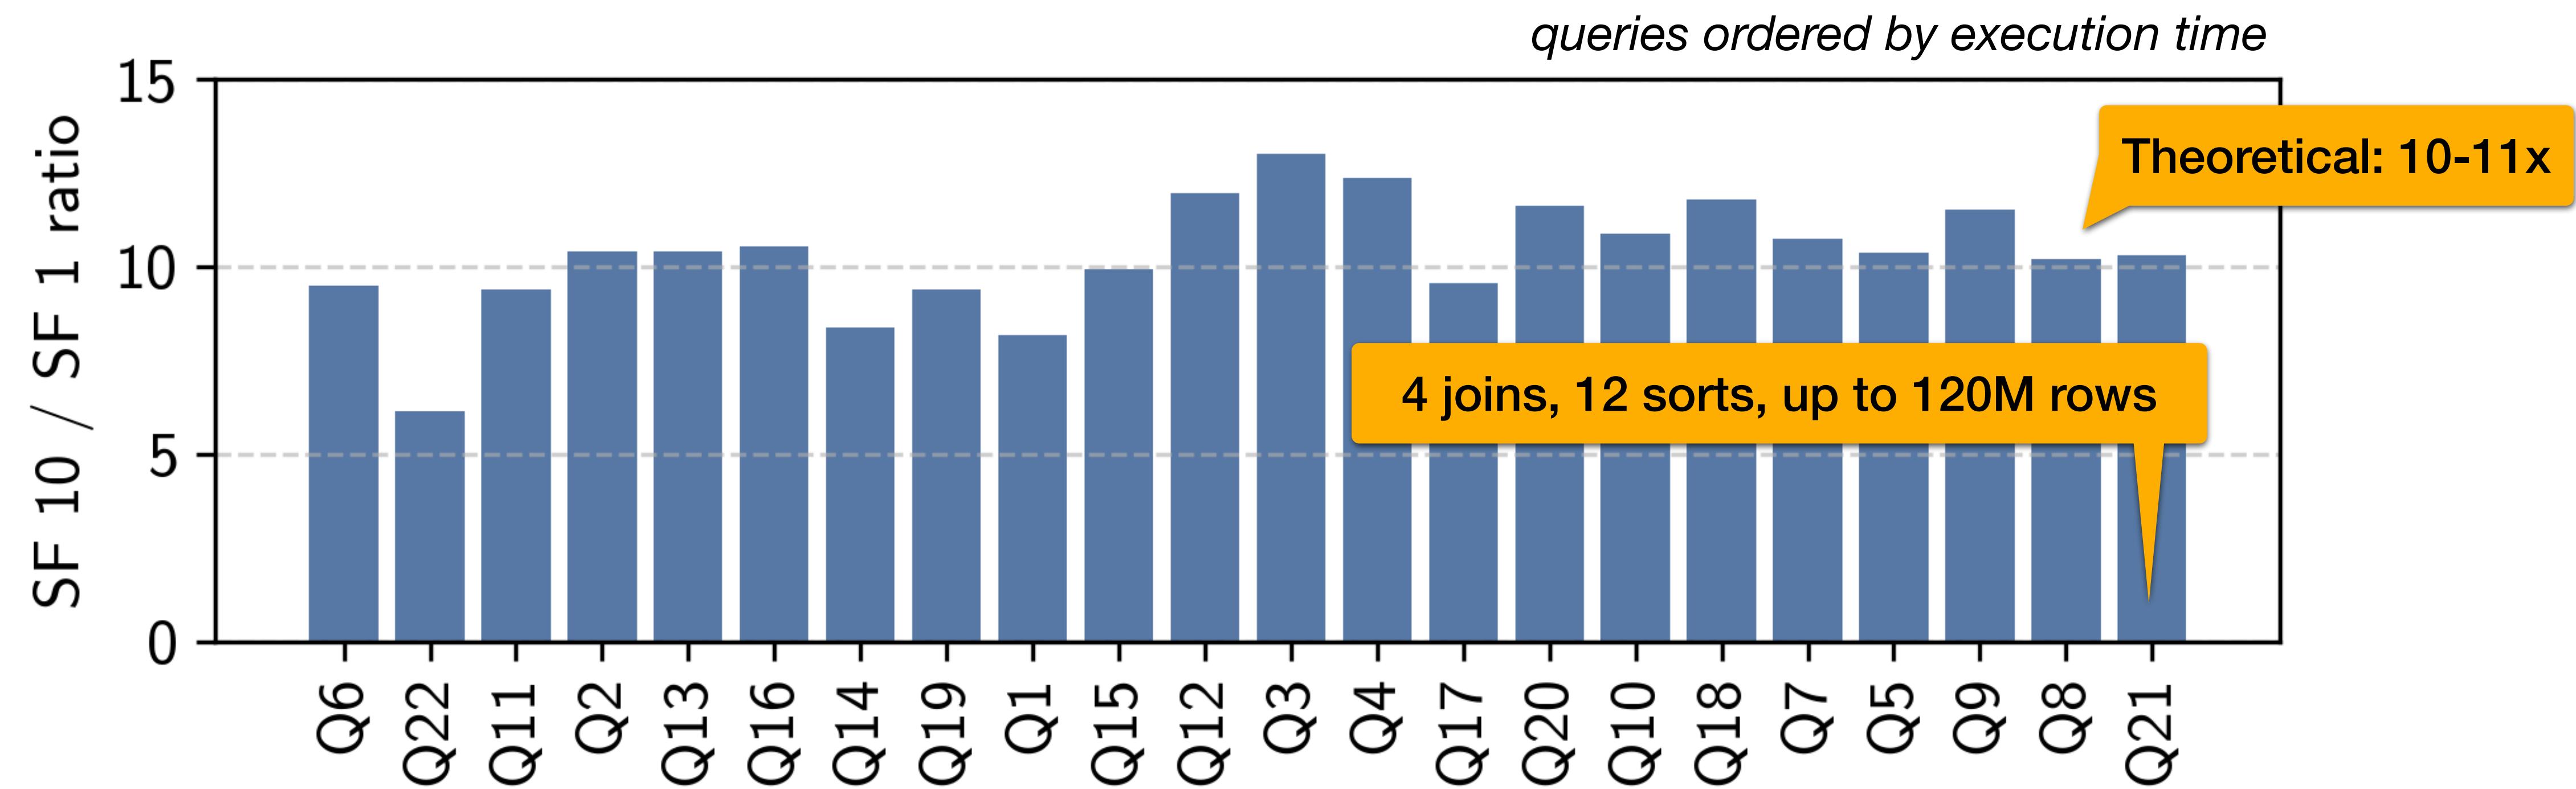
# Scalability

10x larger inputs (SF 10)



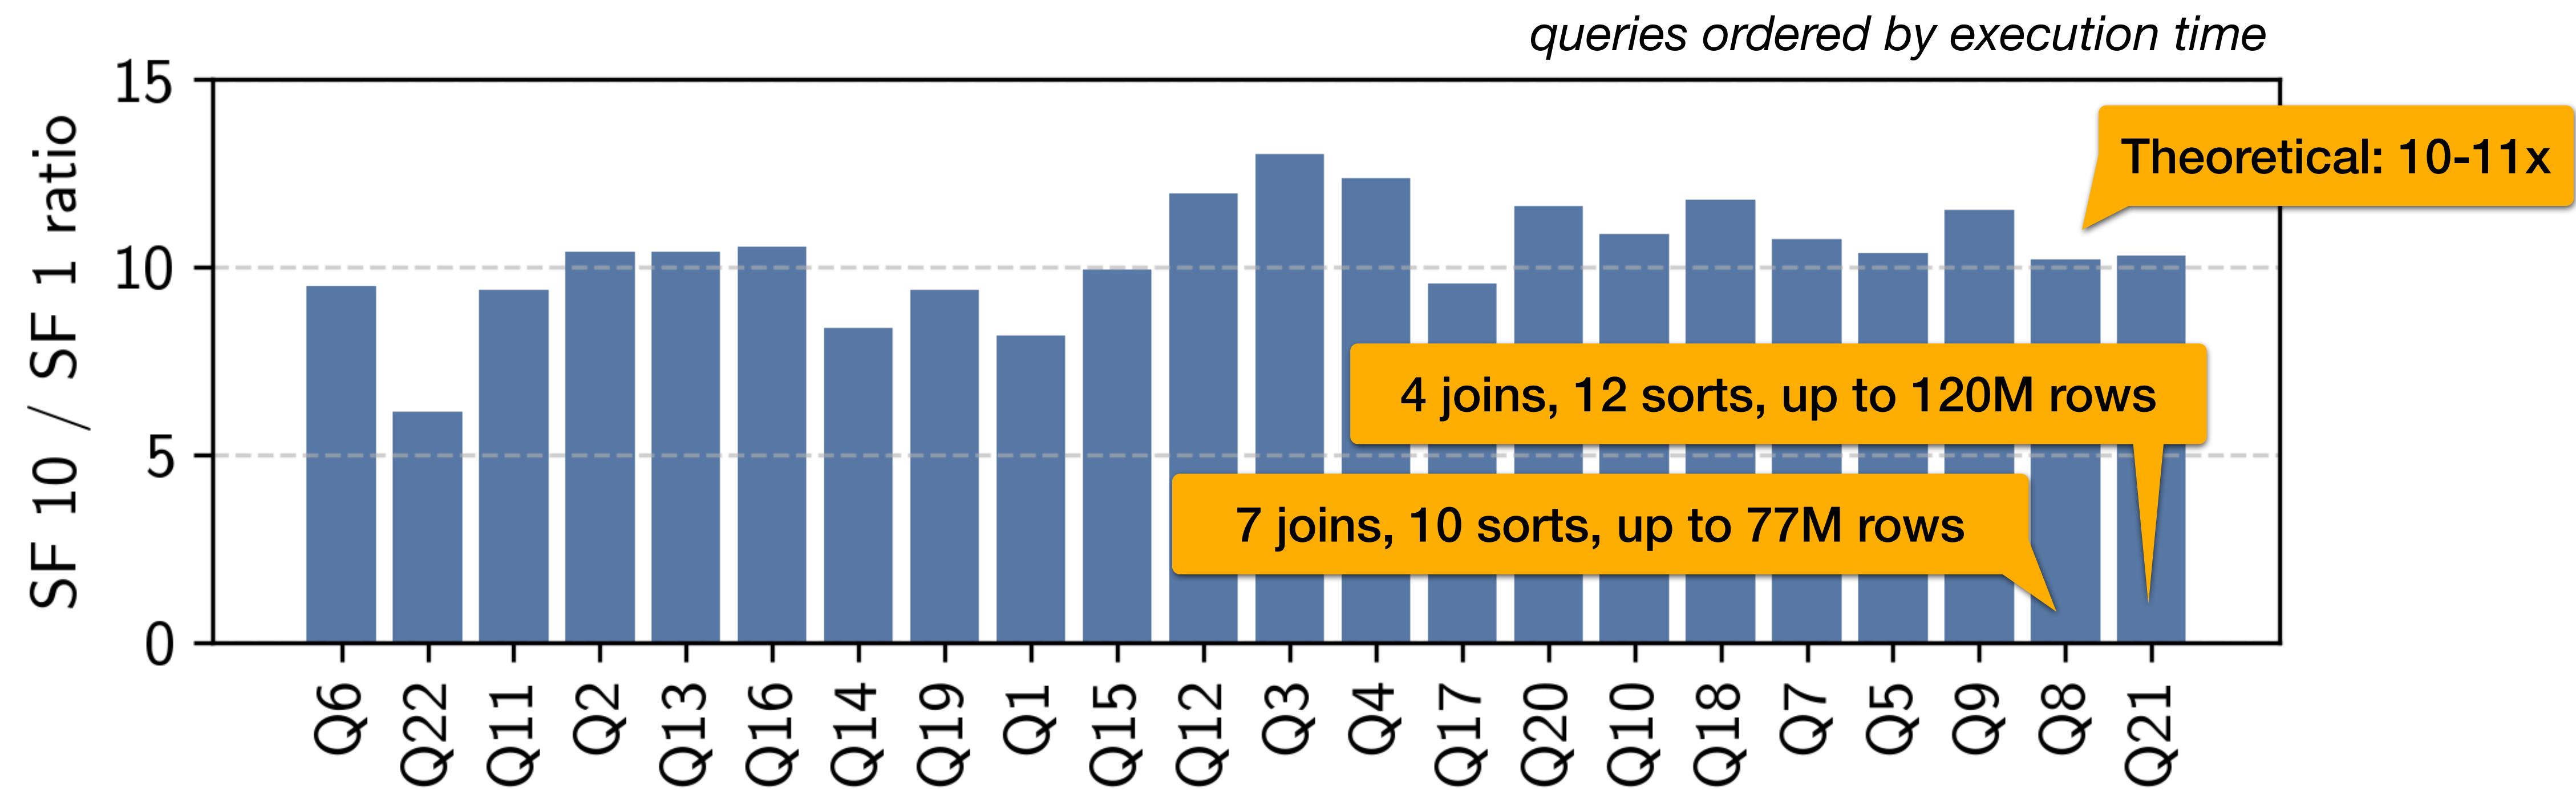
# Scalability

10x larger inputs (SF 10)



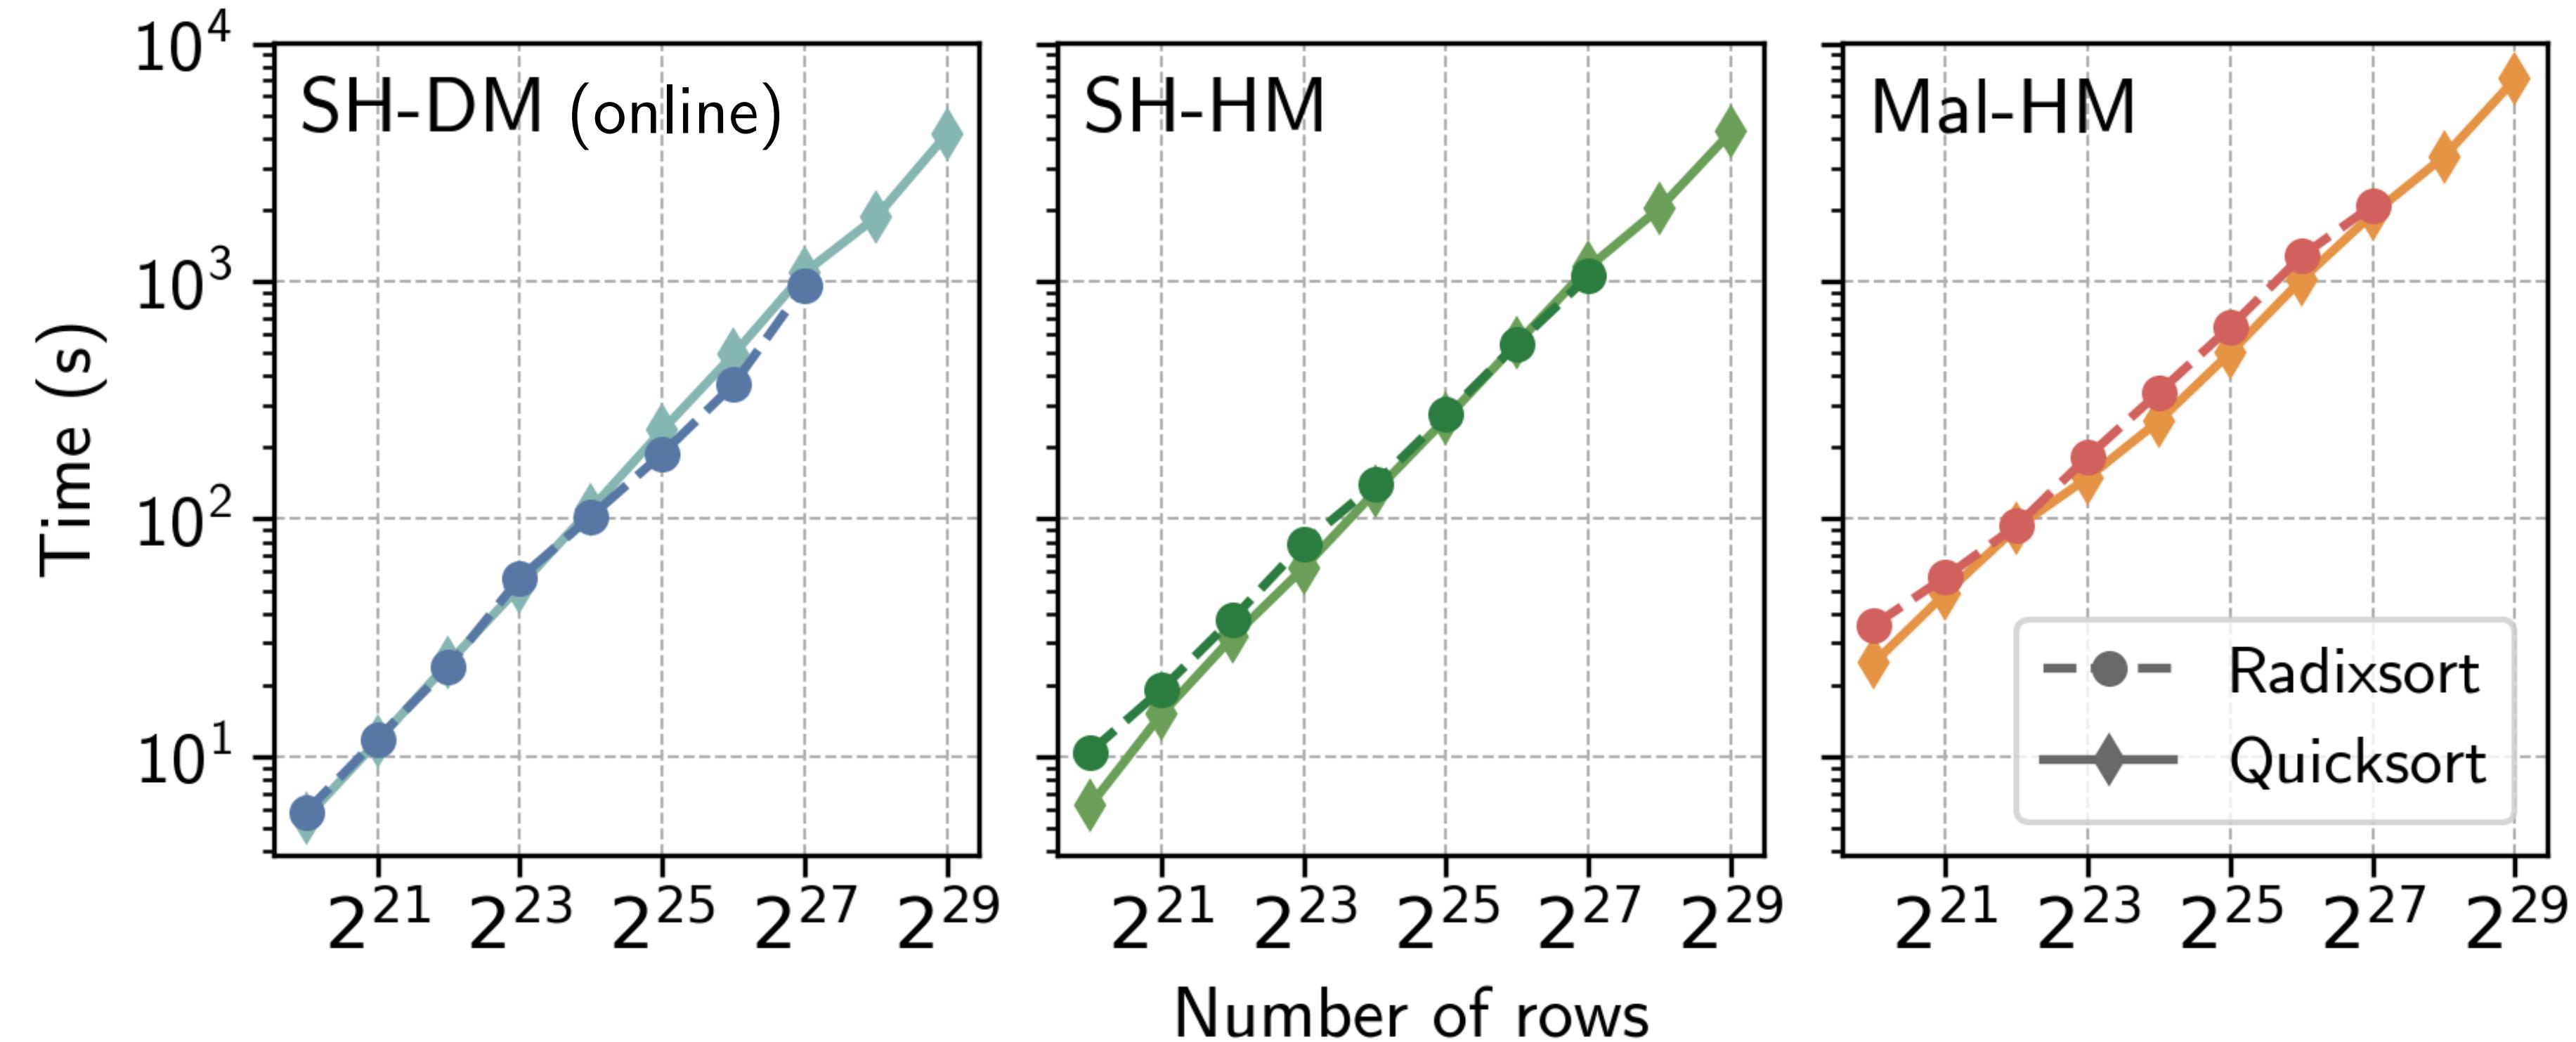
# Scalability

10x larger inputs (SF 10)



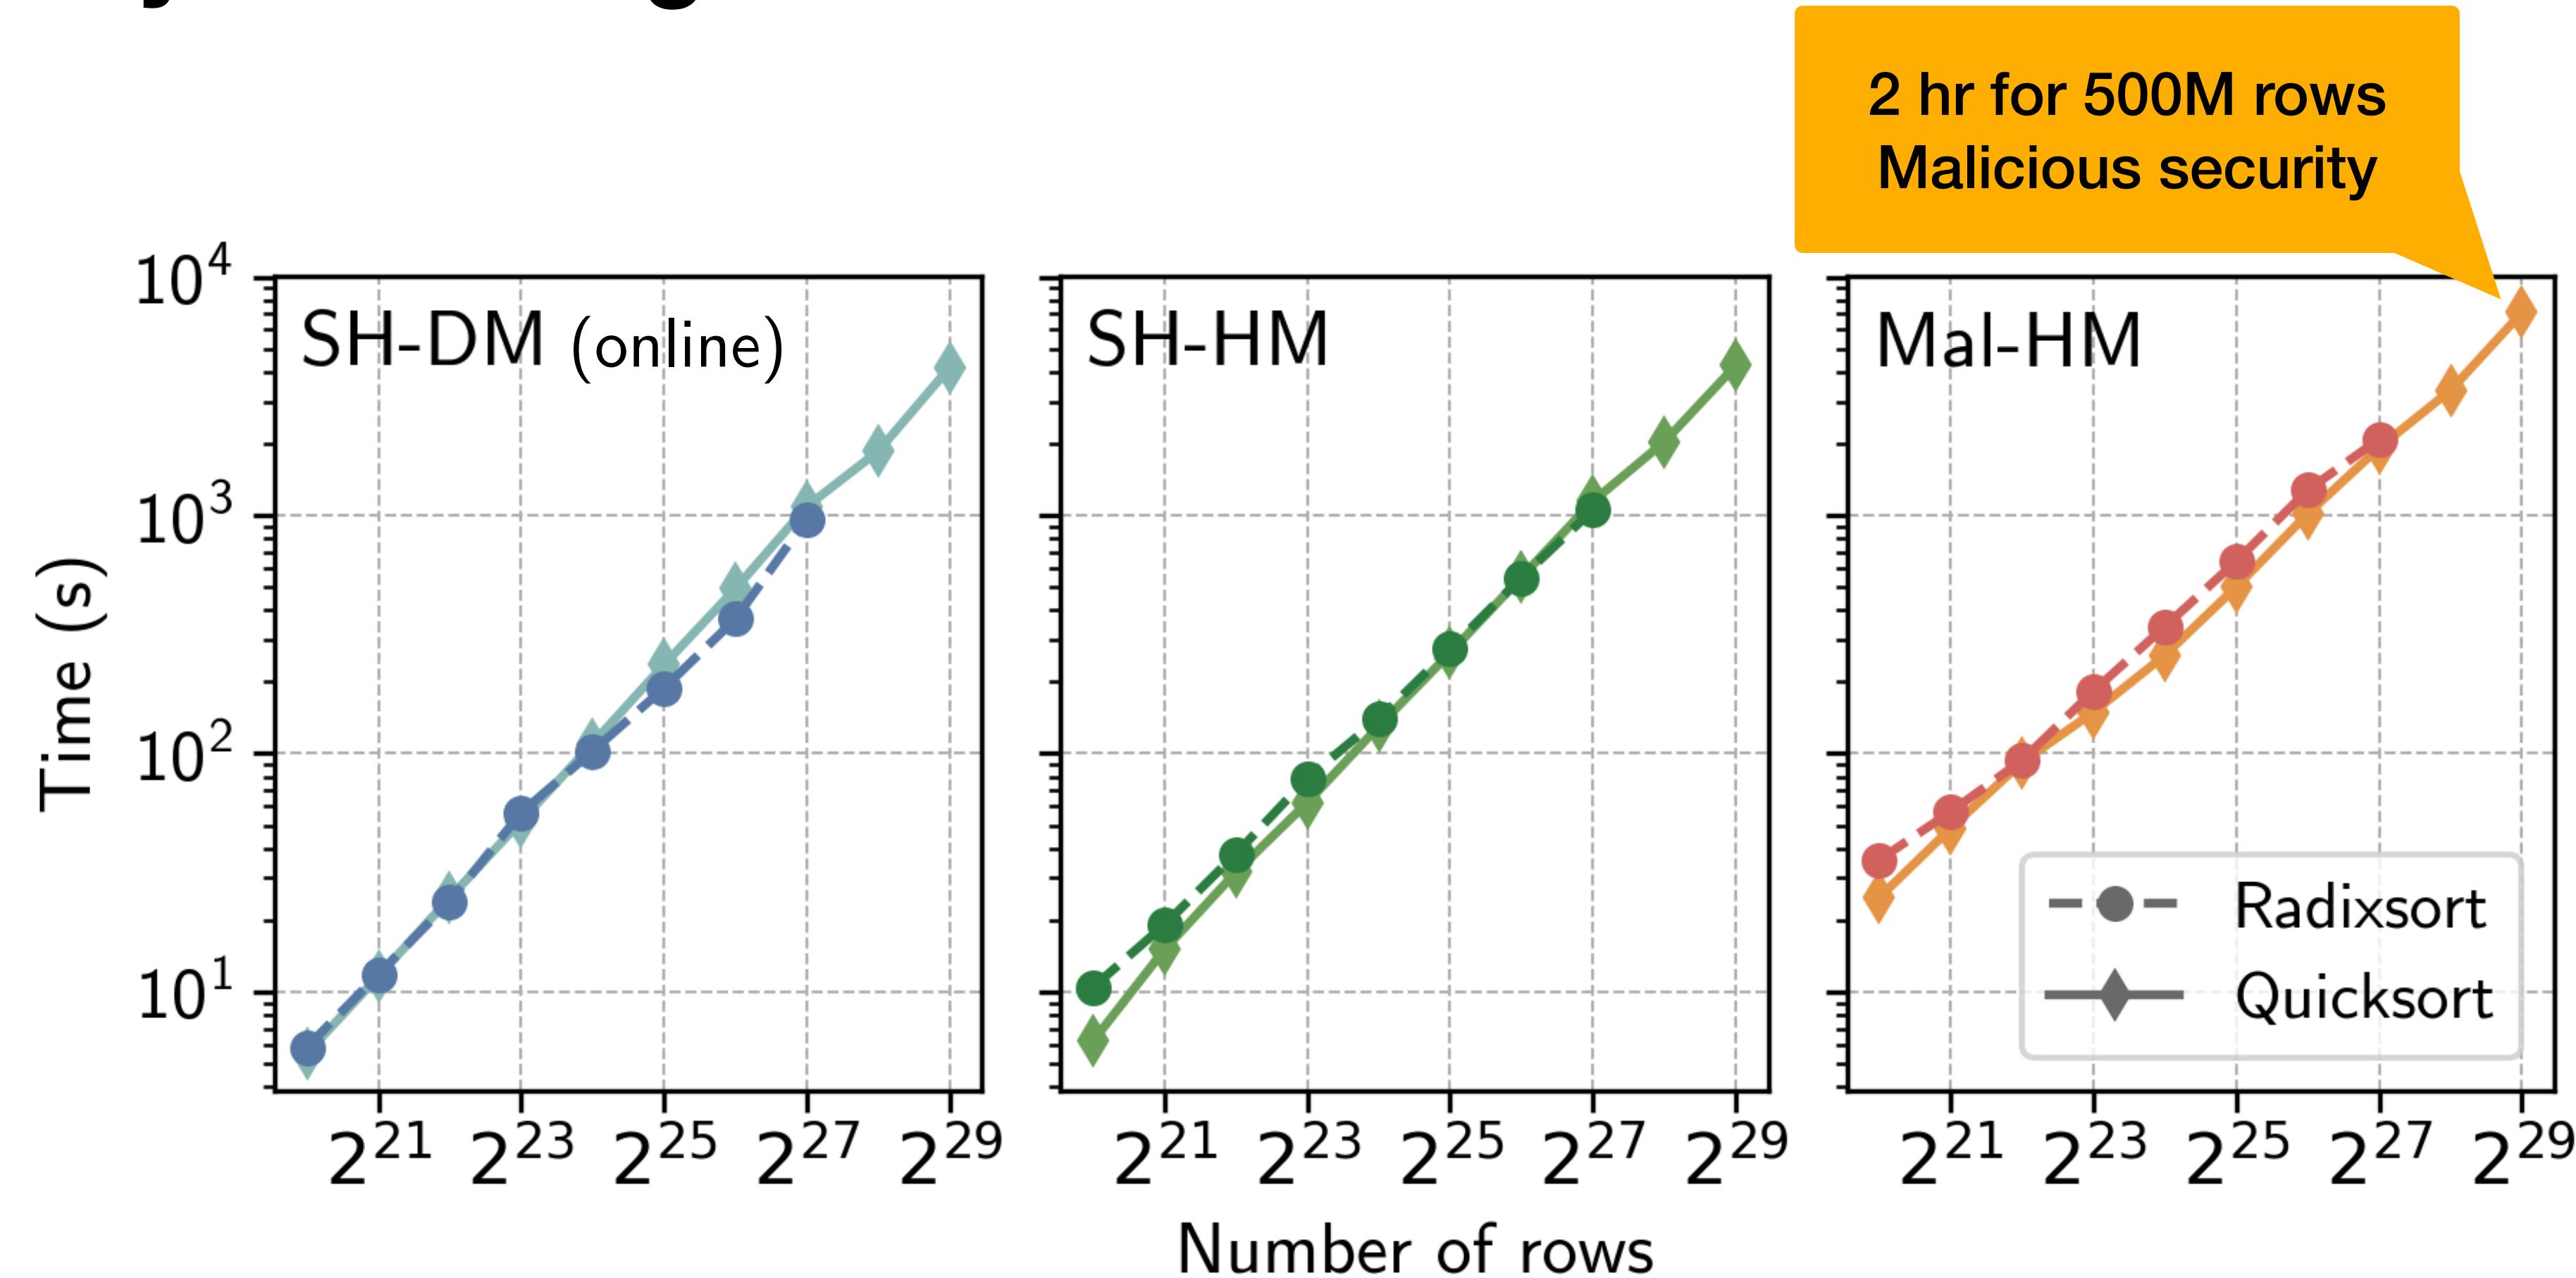
# Going Even Further?

## Scalability of Sorting



# Going Even Further?

## Scalability of Sorting



# ORQ: Complex Analytics on Private Data with Strong Security Guarantees

email me: [elibaum@bu.edu](mailto:elibaum@bu.edu)

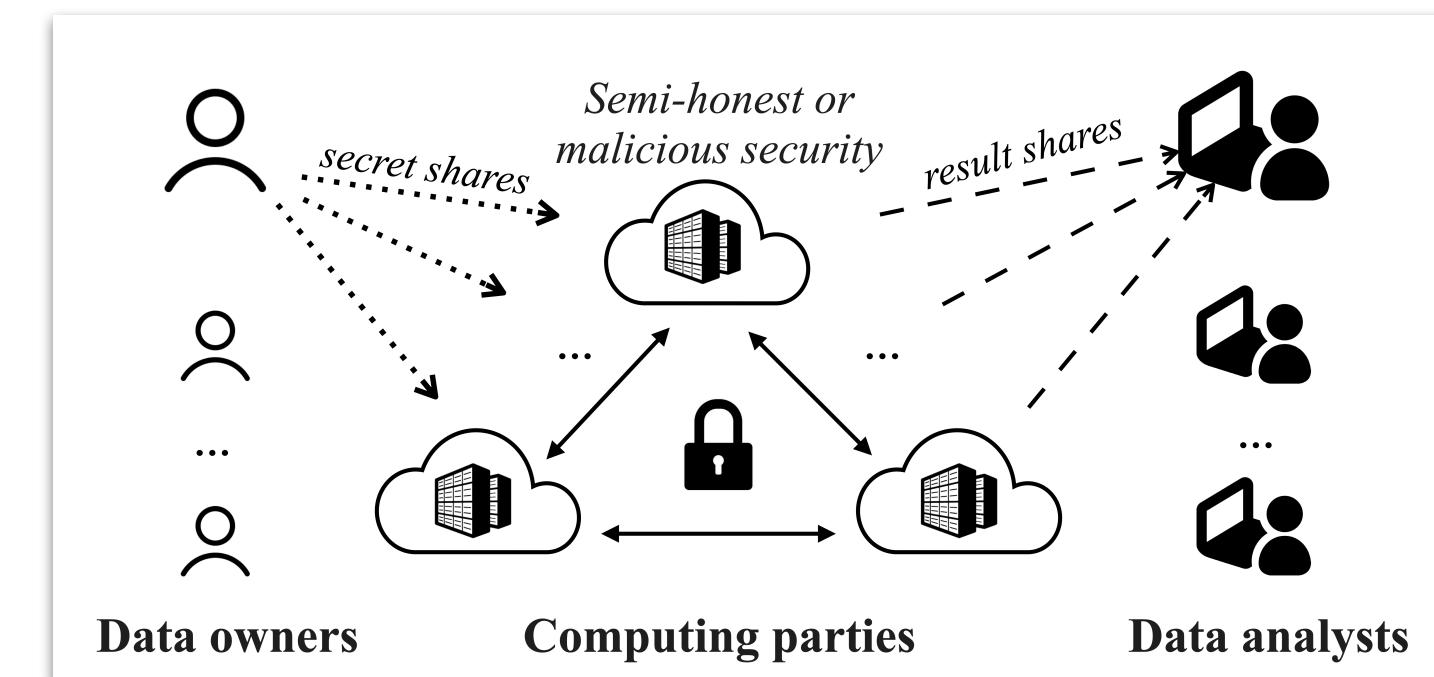


**Distinguished Artifact**



Paper/Code/Docs

[elibaum.com](http://elibaum.com)



## Privacy

Runs on untrusted compute with end-to-end oblivious execution

## Generality

Uses crypto protocols as a blackbox to meet different security requirements

## Expressivity

Efficiently computes all queries from prior work & TPC-H, with no leakage

## Scalability

Evaluates queries with millions of rows & efficiently uses available resources



CASPLab  
**We're hiring!**